

CSC112 Lab1: Printing Vowels and Consonants

You are given a C program “pthread vows cons.c” that creates two threads (`vow` and `cons`) that print words starting with vowels and consonants, respectively, while maintaining the input order of words. It uses the syscall `sched_yield()` to allow threads to take turns.

Compile the program using the command line:

```
gcc -pthread -o threadtest threadtest.c
```

Run the program by supplying a list of words as arguments, for example:

```
./threadtest apple banana orange grape kiwi umbrella
```

The output looks like this:

```
Vowel: apple
Consonant: banana
Vowel: orange
Consonant: grape
Consonant: kiwi
Vowel: umbrella
```

The program works as follows:

1. Input Handling: The program takes command-line arguments as input. Each argument is treated as a word.
2. Thread Creation: The `vow` thread prints words starting with vowels. The `cons` thread prints words starting with consonants.
3. Turn-Based Synchronization:
 - A shared variable `turn` determines which thread should process the current word:
 - `turn == 0`: Vowel thread's turn.
 - `turn == 1`: Consonant thread's turn.
 - Each thread checks whether it's its turn and processes the word accordingly. If it's not its turn, it calls `sched_yield()` to yield control to the other thread.
4. Word Processing:
 - Each thread checks whether the current word starts with a vowel or consonant and prints it accordingly.
 - The `current_index` ensures that both threads process words in sequence.

The system call `sched_yield()` lets the calling thread voluntarily give up control to the OS scheduler, in order to give other active threads a chance to run. The two threads use `sched_yield()` to take turns in printing out the vows and cons. The program preserves the original word sequence in the argument list because both threads operate on a shared index `current_index` and take turns based on the `turn` variable. No mutexes or explicit locking mechanisms are used.

Task 1: Use pthreads and mutex lock to print words in input order

Your first task is to modify the program to use a mutex lock instead of `sched_yield()` while keeping the same program behavior (of printing words starting with vowels and consonants, respectively while maintaining the input order of words). Remove the calls to `sched_yield()`.

Declare a global mutex lock:

```
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

Add calls to `pthread_mutex_lock(&mutex)` and `pthread_mutex_unlock(&mutex)` within the while loop in functions `print_vowels()` and `print_consonants()` to ensure mutual exclusion between them, since they are not yielding to each other voluntarily. Please refer to the following code snippet in `main()`:

```
#include <pthread.h>
...
pthread_t vow_thread, cons_thread;
pthread_create(&vow_thread, NULL, print_vowels, NULL);
pthread_create(&cons_thread, NULL, print_consonants, NULL);
pthread_join(vow_thread, NULL);
pthread_join(cons_thread, NULL);
pthread_mutex_destroy(&mutex);    // Cleanup mutex
```

Try the following code, what do you observe?

```
pthread_create(&vow_thread, NULL, print_vowels, NULL);
pthread_join(vow_thread, NULL);
pthread_create(&cons_thread, NULL, print_consonants, NULL);
pthread_join(cons_thread, NULL);
```

Side note: Mutex locks are essential for thread synchronization when accessing shared resources. After you get the program to work, if you remove the mutex lock/unlock instructions (and also `sched_yield()`), several problems can occur due to race conditions when multiple threads access shared variables (`current_index` and `turn`) simultaneously:

1. Data Corruption: Both threads might try to modify `current_index` at the same time, leading to: Words being skipped; Words being printed twice; Incorrect increment of `current_index`
2. Race Condition on `turn`: Both threads might read `turn` at the same time and: Both threads might think it's their turn; Both threads might print at the same time; The turn-taking mechanism might break down.
3. Output Interleaving: Without synchronization, you might see output like this, where both threads try to print simultaneously, mixing their output.

```
VowConsonant: el: apple
tree
```

However, you are not likely to see these problems since they may occur very rarely, e.g., once in a million runs. Most likely the program will seem to work fine, but it is not reliable and could fail in unpredictable ways.

Task 2: Use pthreads to print all vowels before all consonants

Your second task is to modify the program to use pthreads to achieve the program behavior of printing out all vowels before all consonants. First, you need to modify the functions `print_vowels()` and `print_consonants()` to each iterate through all the input arguments and print out all vowels and consonants, respectively. Second, you need to make sure that the thread that runs `print_vowels()` runs and finishes before the thread that runs `print_consonants()`. This can be done by calling `pthread_create()` to create the first thread that runs `print_vowels()`, and `pthread_join()` to wait for it to finish, before calling `pthread_create()` to create the second thread that runs `print_consonants()`, and `pthread_join()` to wait for it to finish. (You do not need mutex protection for each function, since

the two threads run sequentially, not concurrently.) Please refer to the following code snippet in `main()`:

```
#include <pthread.h>
...
pthread_t vow_thread, cons_thread;
pthread_create(&vow_thread, NULL, print_vowels, NULL);
pthread_join(vow_thread, NULL);
pthread_create(&cons_thread, NULL, print_consonants, NULL);
pthread_join(cons_thread, NULL);
```

Try the following code, what do you observe? Explain why.

```
pthread_create(&vow_thread, NULL, print_vowels, NULL);
pthread_create(&cons_thread, NULL, print_consonants, NULL);
pthread_join(vow_thread, NULL);
pthread_join(cons_thread, NULL);
```

References:

Thread functions in C/C++

<https://www.geeksforgeeks.org/thread-functions-in-c-c/>

How to create and join threads in C (pthreads).

<https://www.youtube.com/watch?v=uA8X5zNOGw8>

Task 3: Use child processes to print all vowels before all consonants (only works on Linux)

Your third task is to modify the program to use process `fork()` and `join()` to achieve the program behavior of printing out all vowels before all consonants. This can be done by calling `fork()` to create the first child process that runs `print_vowels()`, and `waitpid()` to wait for it to finish, before calling `fork()` to create the second child process that runs `print_consonants()`, and `waitpid()` to wait for it to finish. (You do not need mutex protection for each function, since the two processes do not share memory, and they run sequentially.) Please refer to the following code snippet in `main()`:

```
#include <sys/wait.h> //This does not work on Windows
...
pid_t vowel_pid = fork();
waitpid(vowel_pid, NULL, 0);
pid_t cons_pid = fork();
waitpid(cons_pid, NULL, 0);
```

Try the following code, what do you observe? Explain why.

```
pid_t vowel_pid = fork();
pid_t cons_pid = fork();
waitpid(vowel_pid, NULL, 0);
waitpid(cons_pid, NULL, 0);
```

References:

Slide 13 “wait()” in “Lecture 2 Processes and Threads”

fork() in C

<https://www.geeksforgeeks.org/fork-system-call/>

Note: Since Windows and MacOS do not support the process fork() syscall, please work on a Linux machine.

Option 1: SSH to a terminal in Adams Hall. Please refer to the document [COMPUTER SCIENCE LAB TECHNOLOGY GUIDE](#) for more details.

“From a Terminal (e.g. WSL or Cygwin on Windows), enter the following command:

```
ssh -p 5010 username@adams204xx.hofstra.edu
```

The username is your Hofstra Portal username (e.g. h702345678) and the xx is a number from 01 to 30. Type yes and press Enter. Type in your Linux account password and press Enter. If the machine number you try is unreachable, try a different number. If this is the first time you are logging in to your Linux account, you need to login with the default password, which is your 700 number (e.g. 702345678) and it will require you to set a new password. After entering your 700 number as the password, it will ask you to enter the current password again (your 700#). Then you will need to provide your new password twice.”

You can use simple text editors such as Nano, Emacs or VI on the SSH terminal, or if you want to use the VS Code IDE, refer to this link: <https://devblogs.microsoft.com/cppblog/integrated-terminal-in-visual-studio-new-ssh-integration-for-linux-targeting/>

Option 2: Install the Windows Subsystem for Linux (WSL): <https://learn.microsoft.com/en-us/windows/wsl/install>. If you use VS Code, please install VS Code WSL extension <https://code.visualstudio.com/docs/remote/wsl>.

What to submit:

Please submit the following on Canvas:

1. Task 1: A C program named Task1.c that uses the mutex API, and prints words starting with vowels and consonants, respectively while maintaining the original order of input words , and screenshots for running it against some inputs.
2. Task 2: A C program named Task2.c that uses pthreads to print all vowels before all consonants.
3. Task 3: A C program named Task3.c that uses child processes to print all vowels before all consonants.
4. A PDF report explaining the code you have written, the execution results, and your observations of trying some alternatives mentioned in this document.