

CSC 112: Computer Operating Systems

Lecture 2

Processes

Department of Computer Science,
Hofstra University

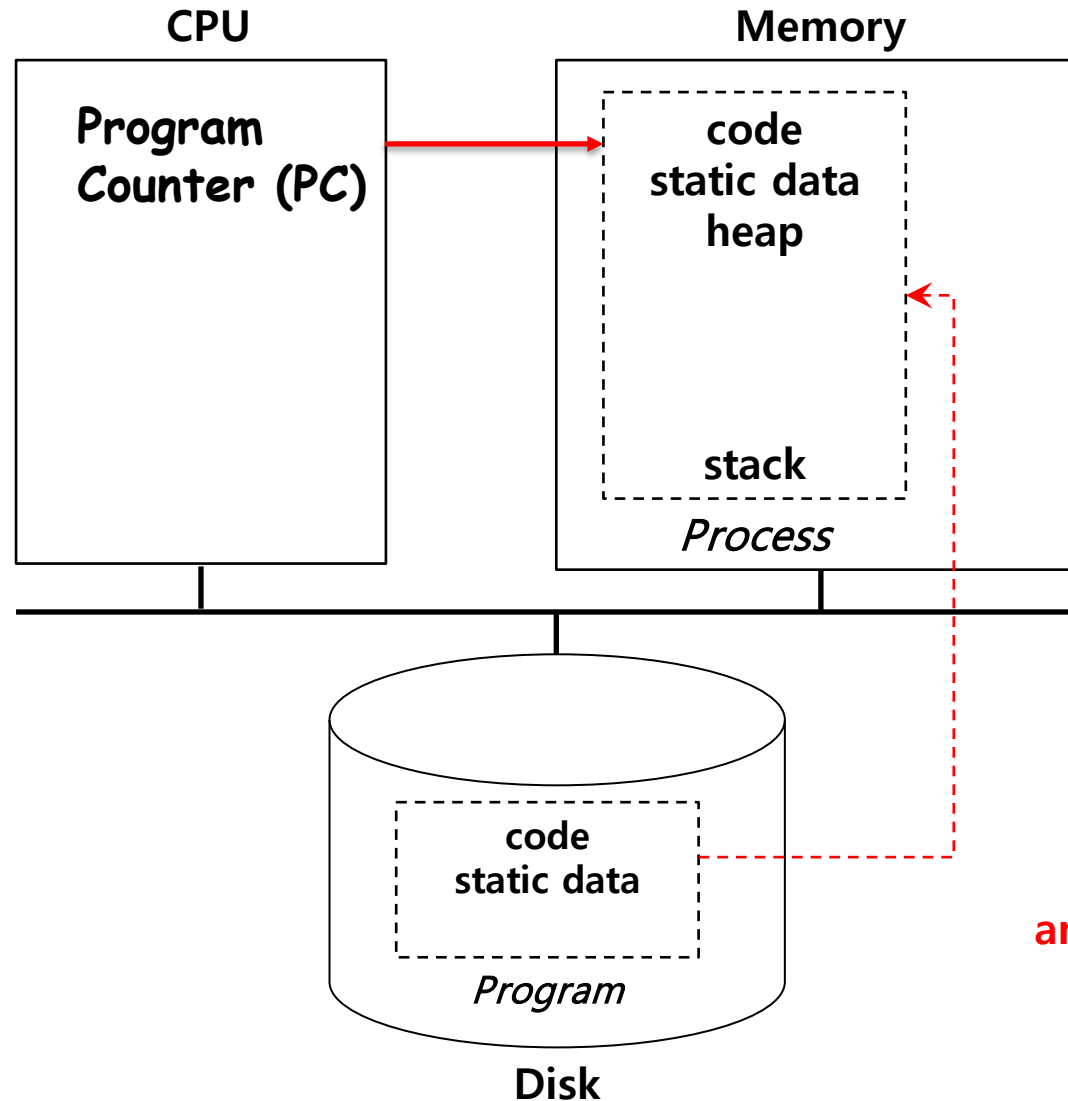
Overview

- Process concept
- Process state
- Process API (creation, wait)
- Process tree

Process

- Program is a *static* entity stored on disk (**executable file**), process is *active*
 - Program becomes process when executable file loaded into memory
 - Process is an abstraction of CPU
- Execution of program started via Graphic User Interface (GUI) mouse clicks, command line entry of its name, etc
- A physical CPU is shared by many processes
 - Time sharing: run one process for a little while, then run another one, and so forth.
 - Processes believe they are using CPU alone

Process



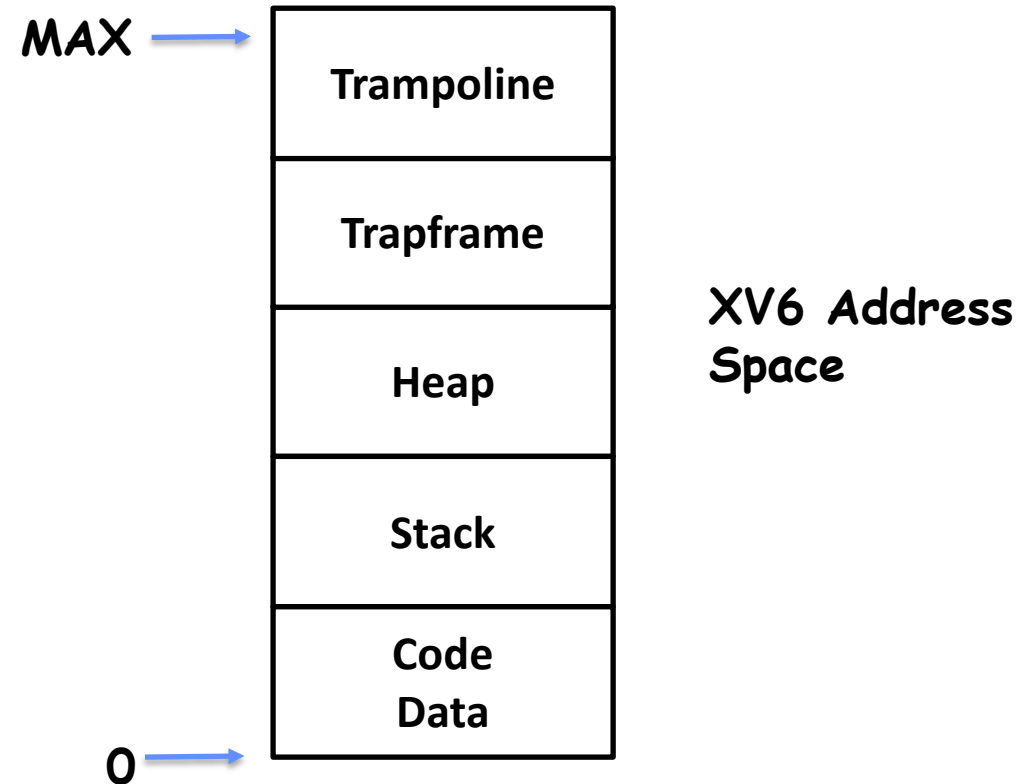
- A program becomes a process when it is selected to execute and loaded into memory.
- A process has an **address space**

Loading:
Takes on-disk
program
and reads it into the
address space of
process

Process

Process: a running program

- **Consists of:**
 - **Code:** Instructions
 - **Stack:** Temporary data, e.g., function parameters, returned addresses, local variables
 - **Registers:** Program counter (PC), general purpose, stack pointer
 - **Data:** Global variables
 - **Heap:** Dynamically allocated



Process

```
struct proc {
    struct spinlock lock; // p->lock must be held when using these:
    enum procstate state; // Process state
    void *chan; // If non-zero, sleeping on chan
    int killed; // If non-zero, have been killed
    int xstate; // Exit status to be returned to parent's wait
    int pid; // Process ID
    // wait_lock must be held when using this:
    struct proc *parent; // Parent process
    // these are private to the process, so p->lock need not be
    held.

    uint64 kstack; // Virtual address of kernel stack
    uint64 sz; // Size of process memory (bytes)
    pagetable_t pagetable; // User page table
    struct trapframe *trapframe; // data page for trampoline.S
    struct context context; // swtch() here to run process
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd; // Current directory
    char name[16]; // Process name (debugging)
};
```

XV6 (proc.h)

- A process is represented by a **process control block (PCB)**
 - Process ID (PID, unique)
 - State
 - Parent process
 - Opened files
 - etc.

Process State

- Process has different states

- **READY**

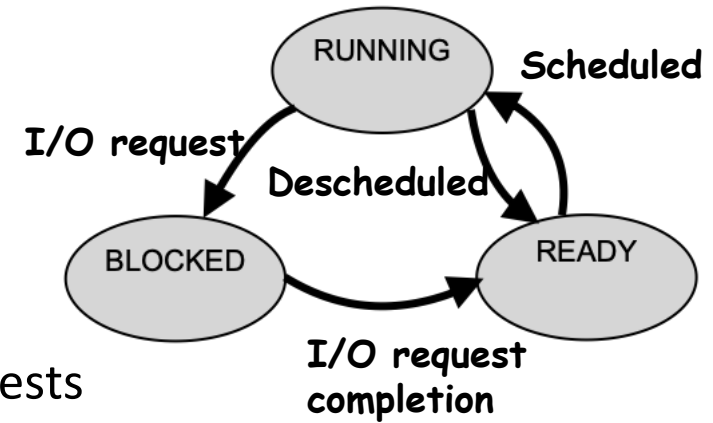
- » Ready to run and pending for running

- **RUNNING**

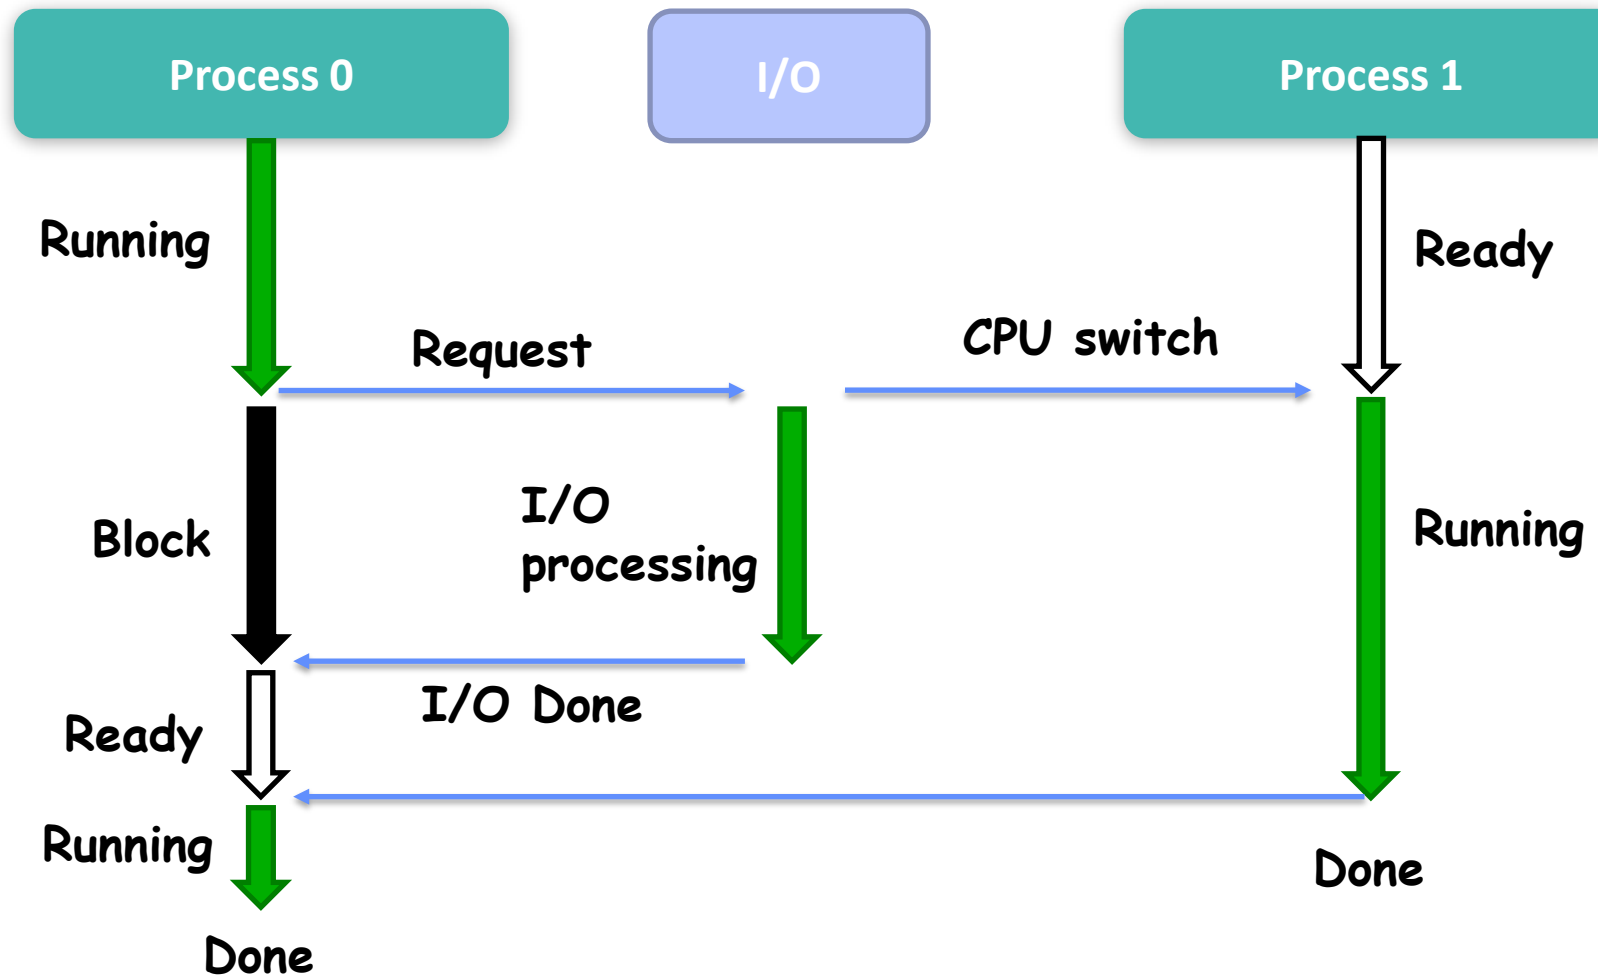
- » Being executed by OS

- **BLOCKED**

- » Suspended due to some other events, e.g., I/O requests



Process State



Process API

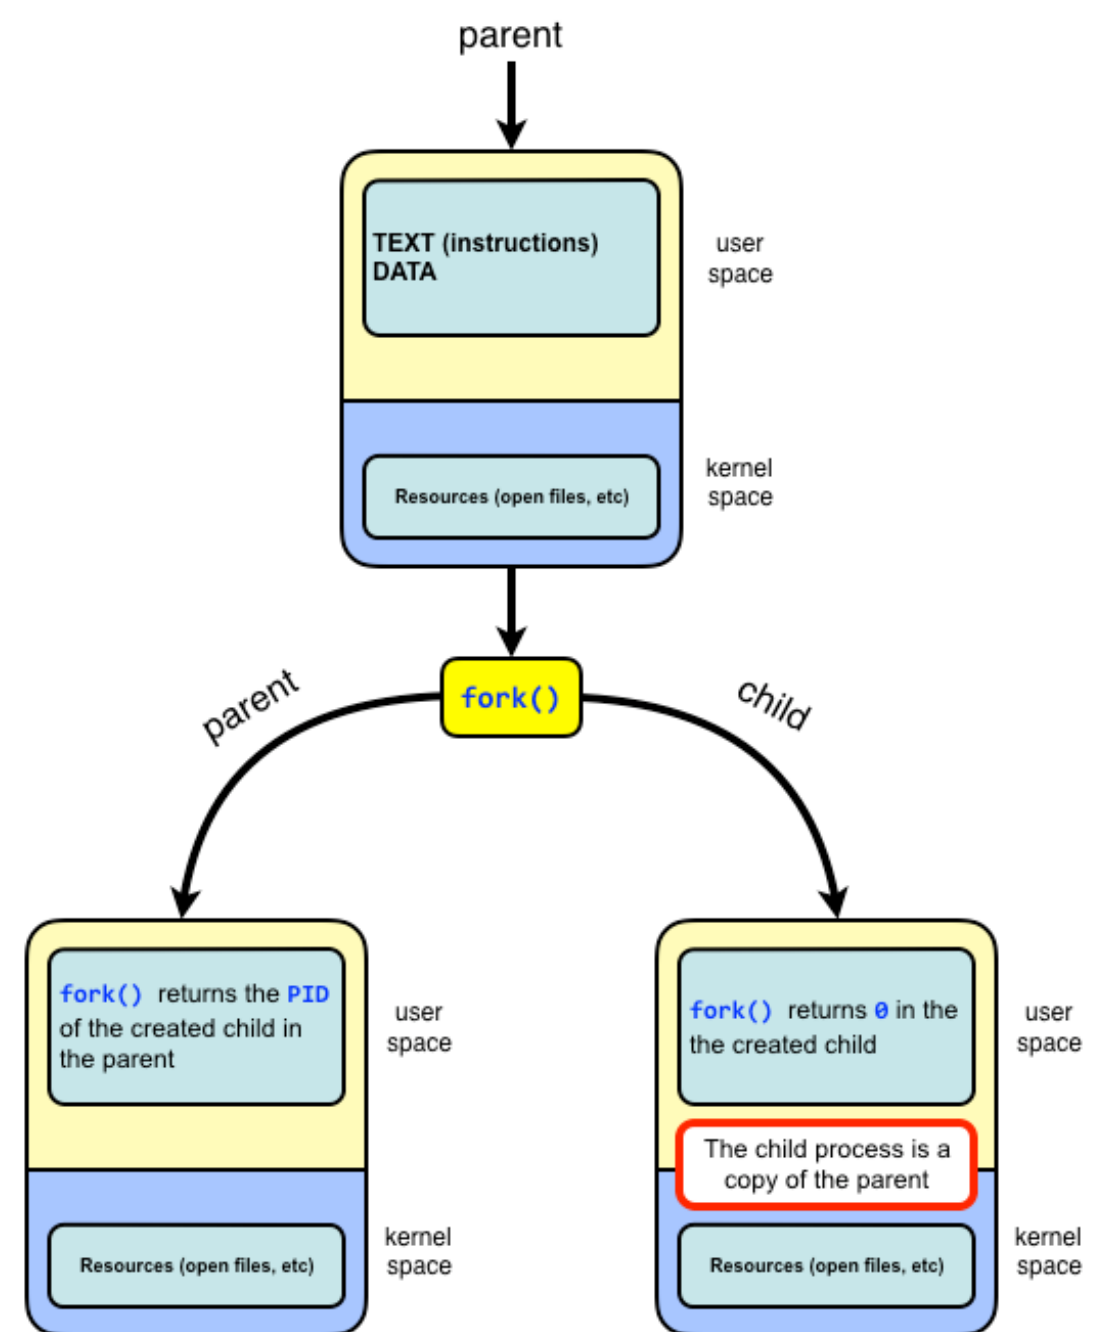
- Process API to manipulate processes
 - **CREATE**
 - » Create a new process, e.g., double click, a command in terminal
 - **WAIT**
 - » Wait for a process to stop
 - » Like I/O request
 - **DESTROY**
 - » Kill the processes
 - **STATUS**
 - » Obtain the information of a process
 - **OTHERS**
 - » Suspend or resume a process

Process Creation

- A process is created by another process, **parent process** or **calling process**
- Process creation relies on two system calls
 - **fork()**
 - » Create a new process and **clone** its parent process
 - **exec()**
 - » Overwrite the created process with a new program

fork()

- A function without any arguments
 - `pid = fork()`
- Both **parent process** and **child process** continue to execute **the instruction following the fork()**
- The return value indicates which process it is (**parent** or **child**)
 - **Non-0 pid** (pid of child process) : return value of the **parent** process,
 - **0** : return value of the new **child** process
 - **-1** : an error or failure occurs when creating new process
- Child process is a **duplicate** of its parent process and has same
 - **instructions, data, stack**
- Child and parents have **different**
 - **PIDs, memory spaces**



fork()

```
int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n"); exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {
        // parent goes down this path (original process)
        printf("hello, I am parent of %d (pid:%d)\n", rc, (int) getpid());
    }
    return 0;
}
```



Output


```
hello world (pid:96744)
hello, I am parent of 96745 (pid:96744)
hello, I am child (pid:96745)
```

Child Process

Parent Process

wait()

- Let the parent process wait for the completion of the child process
 - `pid = wait()`
- `waitpid()` is an alternative of `wait()`



wait()

```
int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        sleep(1);
    } else {
        // parent goes down this path (original process)
        int wc = wait(NULL);
        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n", rc, wc, (int) getpid());
    }
    return 0;
}
```

Child process sleeps for 1 second
Parent process waits for the child process
to finish sleeping

Child Process

Parent Process

wait()

- Without wait(): it is nondeterministic which process (parent or child) runs first

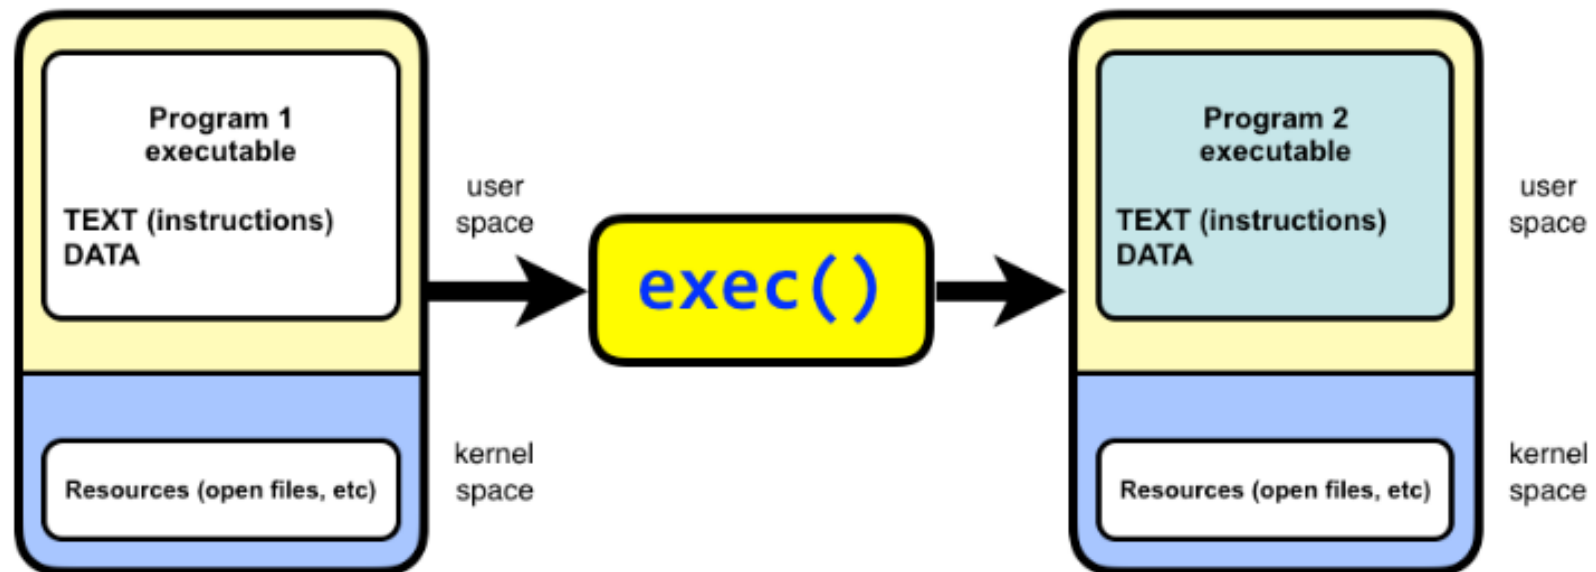
```
hello world (pid:96744)
hello, I am parent of 96745 (pid:96744)
hello, I am child (pid:96745)
```

- With wait(): child runs first, and parents waits for child to finish

```
hello world (pid:96848)
hello, I am child (pid:96849)
hello, I am parent of 96849 (wc:96849) (pid:96848)
```

exec()

- **exec(cmd, argv)** replaces the current process image with a new process image specified by the path to an executable file.
 - **exec()** does not return. It starts to execute the new program.
- There is a family of **exec()**, e.g., **execl()**, **execvp()**
 - **execl()** takes a variable number of arguments that represent the program name and its arguments.
 - **execvp()** takes an array of arguments instead of a variable-length argument list



exec() Example

```
int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n"); exit(1);
    } else if (rc == 0) { // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc"); // program: "wc" (word count)
        myargs[1] = strdup("p3.c"); // argument: file to count
        myargs[2] = NULL; // marks end of array
        execvp(myargs[0], myargs); // run word count
        printf("this will be replaced, so not printed out");
    } else { // parent
        int rc_wait = wait(NULL);
        printf("hello, I am parent of %d (rc_wait:%d) (pid:%d)\n", rc, rc_wait, (int) getpid());
    }
    return 0;
}
```

Output

```
hello world (pid:97511)
hello, I am child (pid:97512)
      32      123      966 p3.c
hello, I am parent of 97512 (wc:97512) (pid:97511)
```


Why fork() + exec()

- Why don't we directly create a new process instead of “**fork+exec**”?
 - **Simple** and **powerful**
 - **Important** to build Unix Shell (an interface to the Unix system)
- By separating **fork()** and **exec()**, we can manipulate various settings just before executing a new program and **make the IO redirection and pipe possible**.
 - IO redirection

– pipe


```
% cat w3.c > newfile.txt
```

```
% echo hello world | wc
```



pipe

- A communication method between two processes




```
(base) dliu@dhcp-10-24-18-121 my_code % cat hello.c
#include <stdio.h>
```

```
int main()
{
    printf("Hello World!\n");
    return 0;
}
```

```
(base) dliu@dhcp-10-24-18-121 my_code % cat hello.c |grep printf
printf("Hello World!\n");
(base) dliu@dhcp-10-24-18-121 my_code % █
```

Process Tree



Process Tree

- % pstree (to show the process tree)

```
(base) dliu@dhcp-10-24-17-236 ~ % pstree
-+-= 00001 root /sbin/launchd
    |--= 00322 root /usr/libexec/logd
    |--= 00323 root /usr/libexec/smd
    |--= 00324 root /usr/libexec/UserEventAgent (System)
```

- % ps (to show all processes)

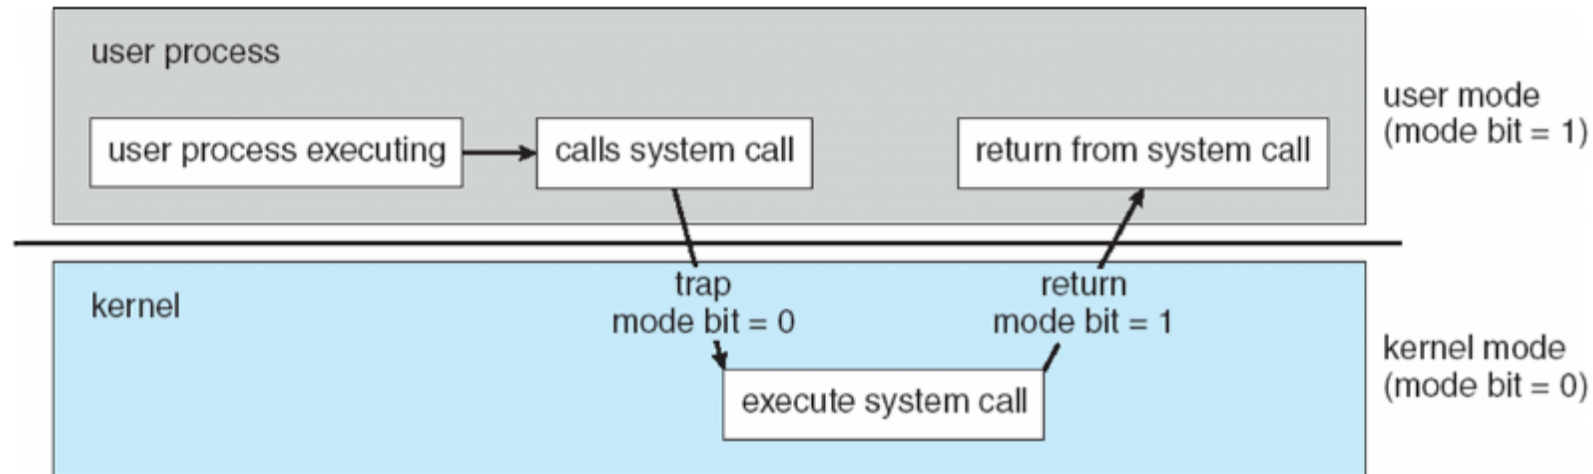
PID	TT	STAT	TIME	COMMAND
1	??	Ss	17:57.36	/sbin/launchd
322	??	Rs	6:29.86	/usr/libexec/logd
323	??	Ss	0:00.19	/usr/libexec/smd
324	??	Ss	0:19.58	/usr/libexec/UserEventAgent (System)

User/Kernel Mode Separation

- **User mode**: restricted, limited operations
 - Processes start in user mode
- **Kernel mode**: privileged, not restricted
 - OS starts in kernel mode
- What if a process wants to perform some restricted operations?
 - **System calls**: Allow the kernel services to provide some functionalities to user programs

User/Kernel Mode Separation

- A process starts in **user mode**
- If it needs to perform a restricted operation, it calls a system call by executing a **trap instruction**.
- The state and registers of the calling process are stored, the system enters **kernel mode**, OS completes the syscall work.
- **Return from syscall**, restore the states and registers of the process, and resume the execution of the process



Process Scheduling

- **Switching Between Processes**
 - Cooperative approach
 - Non-cooperative approach
- **Cooperative approach**
 - Trust process to relinquish CPU to OS through traps
 - » System calls
 - » Illegal operations, e.g., divided by zero
 - **Issue: if no system call**
- **Non-cooperative approach**
 - The OS takes control
 - OS obtains control periodically, e.g., timer interrupter

Summary

- In OS, process is a running program and has an address space
- We use process API to create and manage processes
- Fork() to duplicate a process, exec() to replace the command
- Process scheduling