

CSC 112: Computer Operating Systems

Lecture 2

Processes and Threads

Department of Computer Science,
Hofstra University

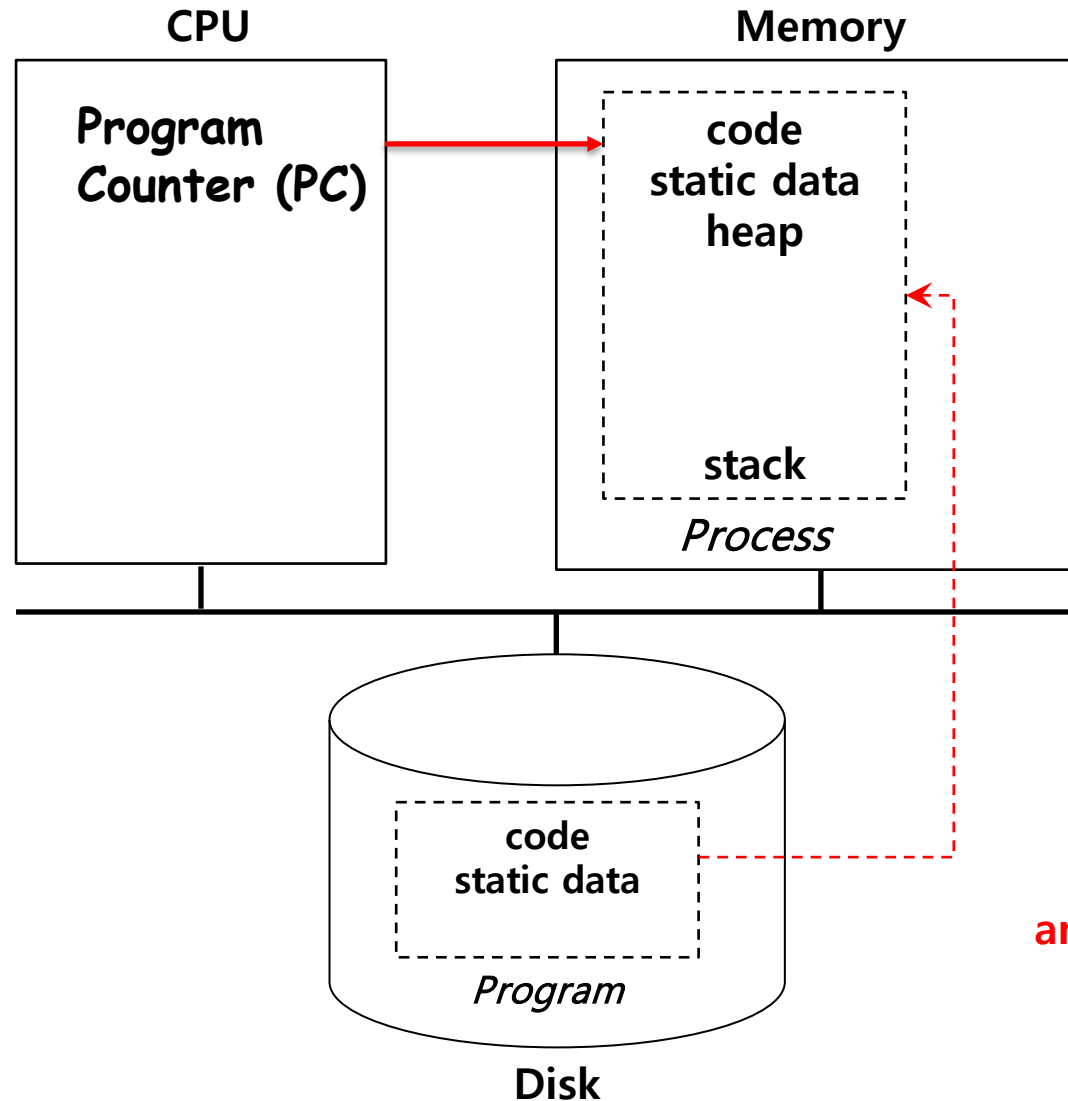
Overview

- Process concept
- Process state
- Process API (creation, wait)
- Process tree

Process

- Program is a *static* entity stored on disk (**executable file**), process is *active*
 - Program becomes process when executable file loaded into memory
 - Process is an abstraction of CPU
- Execution of program started via Graphic User Interface (GUI) mouse clicks, command line entry of its name, etc
- A physical CPU is shared by many processes
 - Time sharing: run one process for a little while, then run another one, and so forth.
 - Processes believe they are using CPU alone

Process



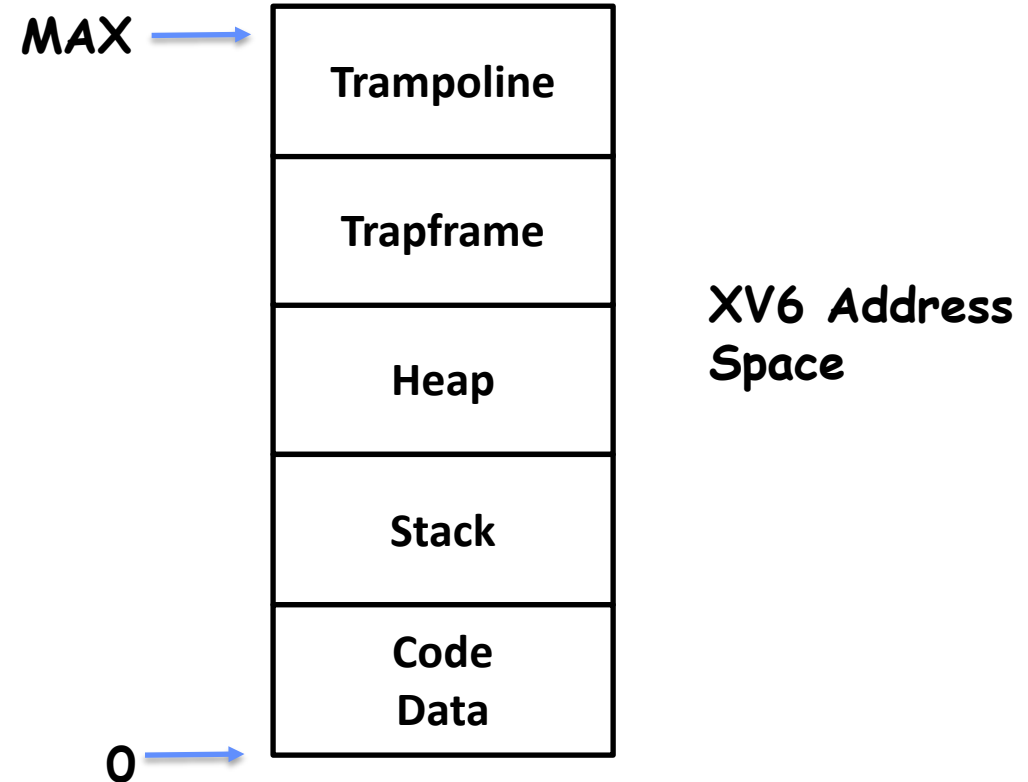
- A program becomes a process when it is selected to execute and loaded into memory.
- A process has an **address space**

Loading:
Takes on-disk
program
and reads it into the
address space of
process

Process

Process: a running program

- **Consists of:**
 - **Code:** Instructions
 - **Stack:** Temporary data, e.g., function parameters, returned addresses, local variables
 - **Registers:** Program counter (PC), general purpose, stack pointer
 - **Data:** Global variables
 - **Heap:** Dynamically allocated



Process

```
struct proc {
    struct spinlock lock; // p->lock must be held when using these:
    enum procstate state; // Process state
    void *chan; // If non-zero, sleeping on chan
    int killed; // If non-zero, have been killed
    int xstate; // Exit status to be returned to parent's wait
    int pid; // Process ID
    // wait_lock must be held when using this:
    struct proc *parent; // Parent process
    // these are private to the process, so p->lock need not be
    held.

    uint64 kstack; // Virtual address of kernel stack
    uint64 sz; // Size of process memory (bytes)
    pagetable_t pagetable; // User page table
    struct trapframe *trapframe; // data page for trampoline.S
    struct context context; // swtch() here to run process
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd; // Current directory
    char name[16]; // Process name (debugging)
};
```

XV6 (proc.h)

- A process is represented by a **process control block (PCB)**
 - Process ID (PID, unique)
 - State
 - Parent process pointer
 - Opened files
 - Many other fields
 - PCB in XV6 does not include pointers to child processes for simplicity, but PCB in Linux include them for convenient references to its child processes

Process State

- Process has different states

- **READY**

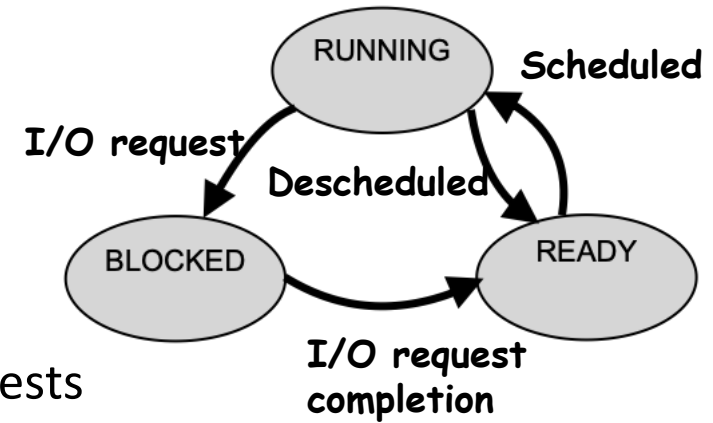
- » Ready to run and pending for running

- **RUNNING**

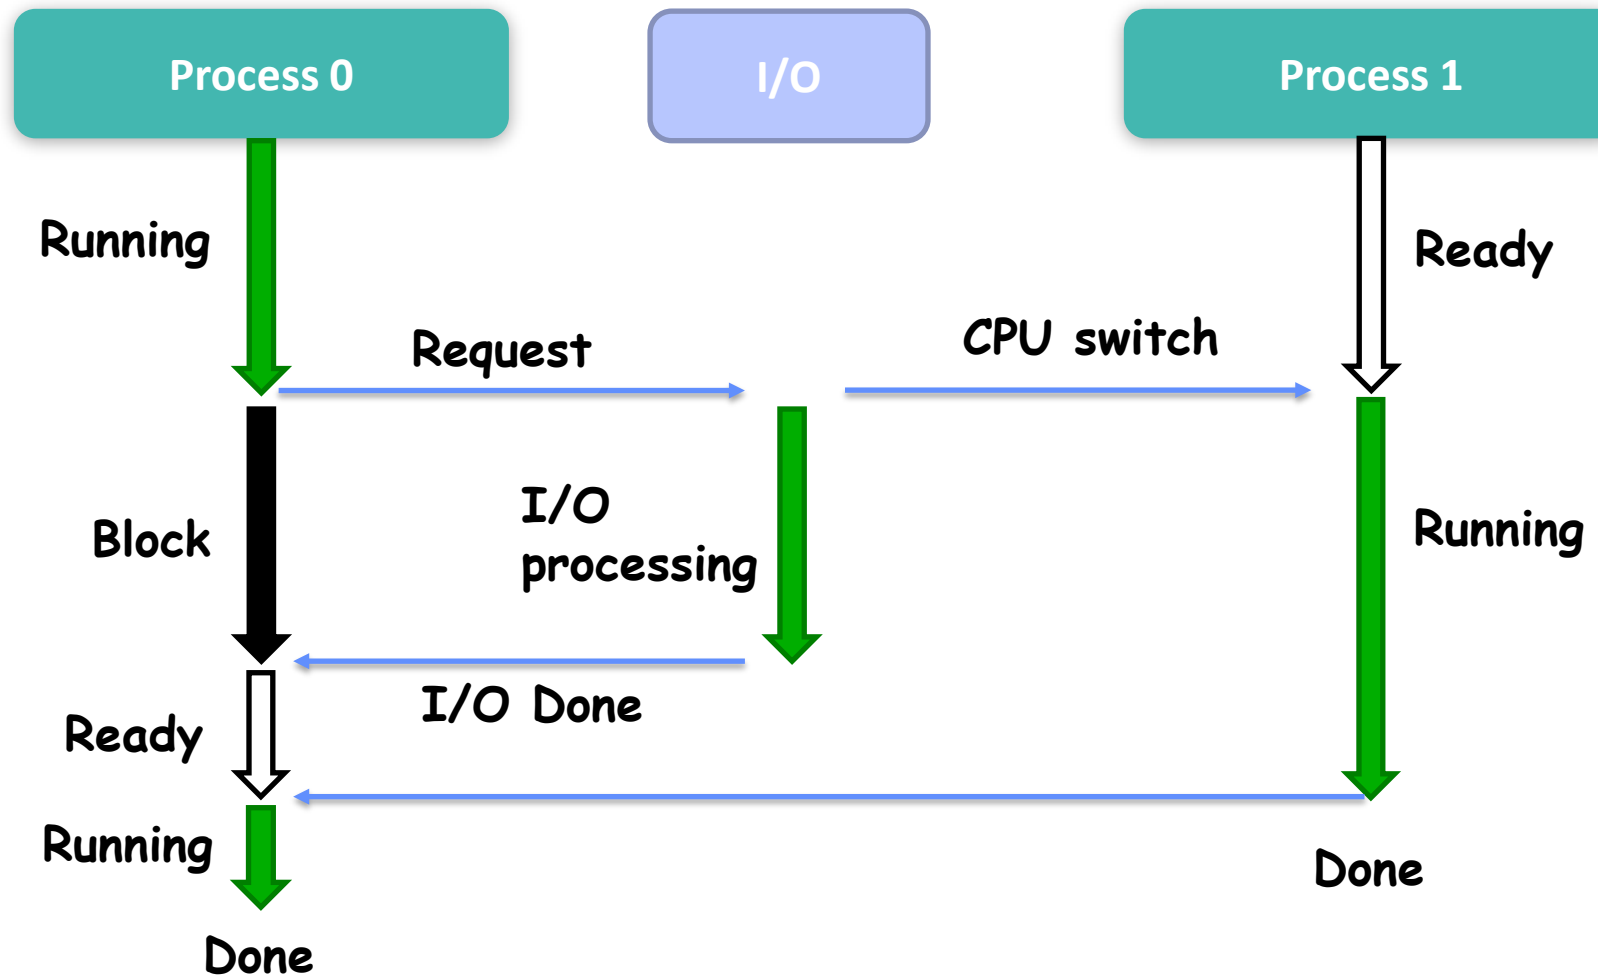
- » Being executed by OS

- **BLOCKED**

- » Suspended due to some other events, e.g., I/O requests



Process State



Process API

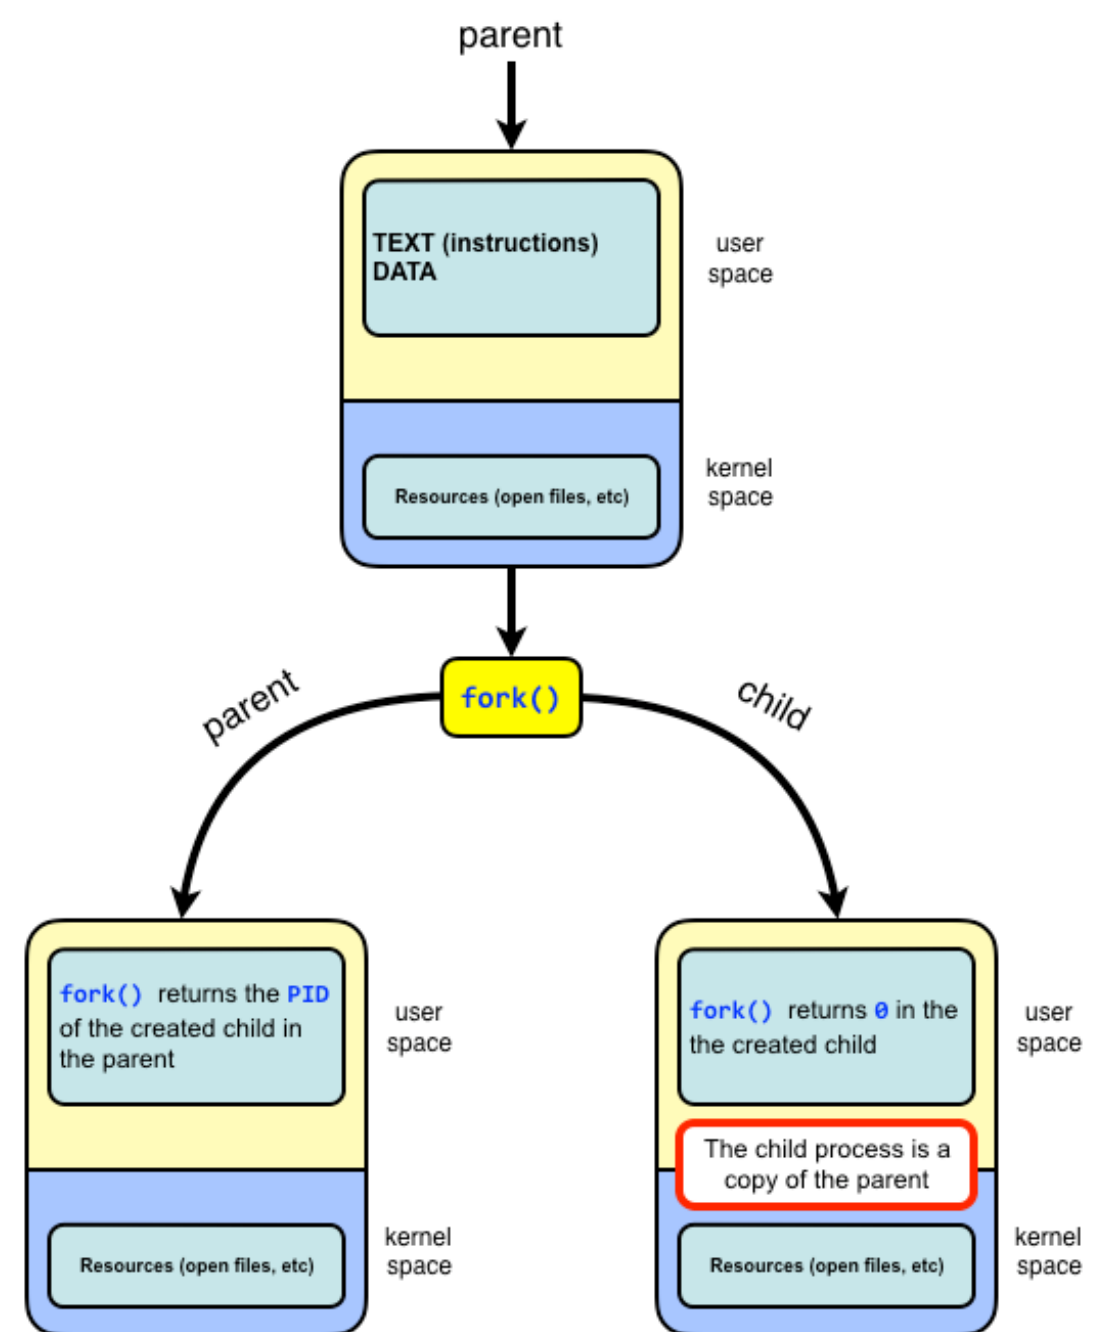
- Process API to manipulate processes
 - **CREATE**
 - » Create a new process, e.g., double click, a command in terminal
 - **WAIT**
 - » Wait for a process to stop
 - » Like I/O request
 - **DESTROY**
 - » Kill the processes
 - **STATUS**
 - » Obtain the information of a process
 - **OTHERS**
 - » Suspend or resume a process

Process Creation

- A process is created by another process, **parent process** or **calling process**
- Process creation relies on two system calls
 - **fork()**
 - » Create a new process and **clone** its parent process
 - **exec()**
 - » Overwrite the created process with a new program

fork()

- A function without any arguments
 - `pid = fork()`
- Both **parent process** and **child process** continue to execute **the instruction following the fork()**
- The return value indicates which process it is (**parent** or **child**)
 - **Non-0 pid** (pid of child process) : return value of the **parent** process,
 - **0** : return value of the new **child** process
 - **-1** : an error or failure occurs when creating new process
- Child process is a **duplicate** of its parent process and has same
 - **instructions, data, stack**
- Child and parents have **different**
 - **PIDs, memory spaces**



fork()

```
int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n"); exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {
        // parent goes down this path (original process)
        printf("hello, I am parent of %d (pid:%d)\n", rc, (int) getpid());
    }
    return 0;
}
```



Output


```
hello world (pid:96744)
hello, I am parent of 96745 (pid:96744)
hello, I am child (pid:96745)
```

Child Process

Parent Process

wait()

- Let the parent process wait for the completion of the child process
 - `pid = wait()`
- **wait()** suspends the execution of the calling process until one of its child processes terminates. It does not allow the parent to specify which child process to wait for. It will reap any terminated child arbitrarily.
- **waitpid(pid)** is an advanced version of wait. It allows the parent process to specify which child process (or group of processes) it wants to wait for.



wait()

```
int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        sleep(1);
    } else {
        // parent goes down this path (original process). wc stores pid of the child process that is waited for
        int wc = wait(NULL); //wc contains pid of the child process being waited for by parent process
        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n", rc, wc, (int) getpid());
    }
    return 0;
}
```

Child process sleeps for 1 second
Parent process waits for the child process to finish sleeping

Child Process

Parent Process

wait()

- **Without wait():** it is nondeterministic which process (parent or child) runs first

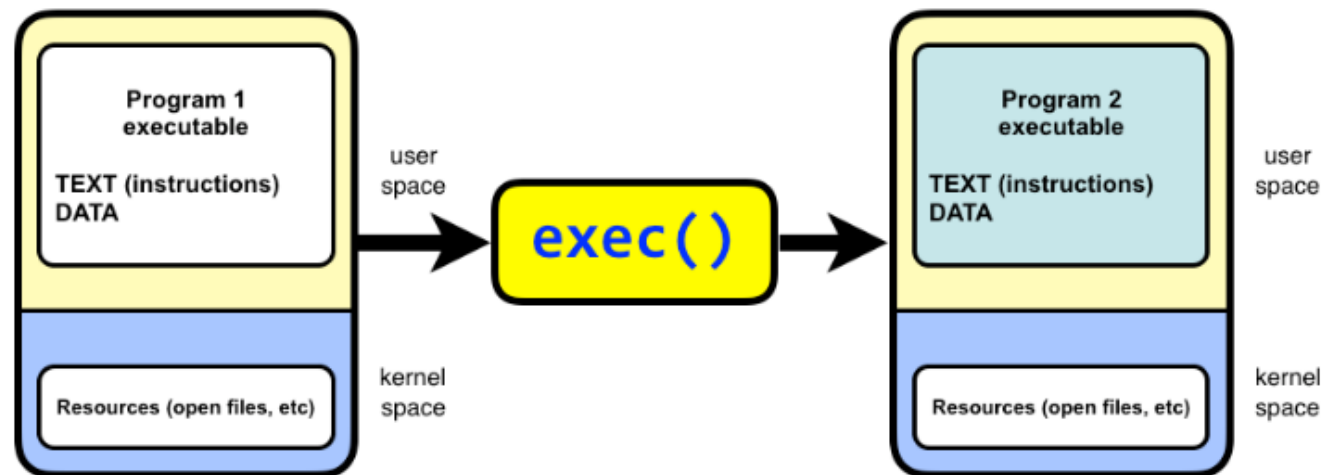
```
hello world (pid:96744)
hello, I am parent of 96745 (pid:96744)
hello, I am child (pid:96745)
```

- **With wait():** child runs first, and parents waits for child to finish

```
hello world (pid:96848)
hello, I am child (pid:96849)
hello, I am parent of 96849 (wc:96849) (pid:96848)
```

exec()

- **exec(cmd, argv)** replaces the current process image with a new process image specified by the path to an executable file.
 - It does not return. It starts to execute the new program.
- There is a family of **exec()**, e.g., **execl()**, **execvp()**
 - **execl()** takes a variable number of arguments that represent the program name and its arguments.
 - » `int execl(const char *path, const char *arg, ..., NULL);`
 - **execvp()** takes an array of arguments instead of a variable-length argument list
 - » `int execvp(const char *file, char *const argv[]);`




exec() Example

```
int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n"); exit(1);
    } else if (rc == 0) { // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = strdup("wc"); // program: "wc" (word count)
        myargs[1] = strdup("p3.c"); // argument: file to count
        myargs[2] = NULL; // marks end of array
        execvp(myargs[0], myargs); // run word count
        printf("this will be replaced, so not printed out");
    } else { // parent
        int rc_wait = wait(NULL);
        printf("hello, I am parent of %d (rc_wait:%d) (pid:%d)\n", rc, rc_wait, (int) getpid());
    }
    return 0;
}
```

wc: counts Lines, Words, and Bytes in a File:
Output format: [lines] [words] [bytes] [filename]

```
hello world (pid:97511)
hello, I am child (pid:97512)
      32      123      966 p3.c
hello, I am parent of 97512 (wc:97512) (pid:97511)
```



IO redirection and pipe

- By separating `fork()` and `exec()`, we can manipulate various settings just before executing a new program and **make the IO redirection and pipe possible**.

- IO redirection: output of the left command redirected to be written to the file on the right


```
% cat w3.c > newfile.txt
```

- Pipe: output of the left command passed as input to the right command

```
% echo hello world | wc
```

pipe

- A communication method between two processes



```
(base) dliu@dhcp-10-24-18-121 my_code % cat hello.c
```

```
#include <stdio.h>

int main()
{
    printf("Hello World!\n");
    return 0;
}
```




Command “cat” prints out content of hello.c file

Output of “cat” command passed through the pipe to command “grep” to search for any lines that contain “printf”



```
(base) dliu@dhcp-10-24-18-121 my_code % cat hello.c |grep printf
    printf("Hello World!\n");
(base) dliu@dhcp-10-24-18-121 my_code % █
```

Process Tree



Process Tree

- % pstree (to show the process tree in a hierarchy)

```
(base) dliu@dhcp-10-24-17-236 ~ % pstree
-+-= 00001 root /sbin/launchd
    |--= 00322 root /usr/libexec/logd
    |--= 00323 root /usr/libexec/smd
    |--= 00324 root /usr/libexec/UserEventAgent (System)
```

- % ps (to show all processes as a flat list)

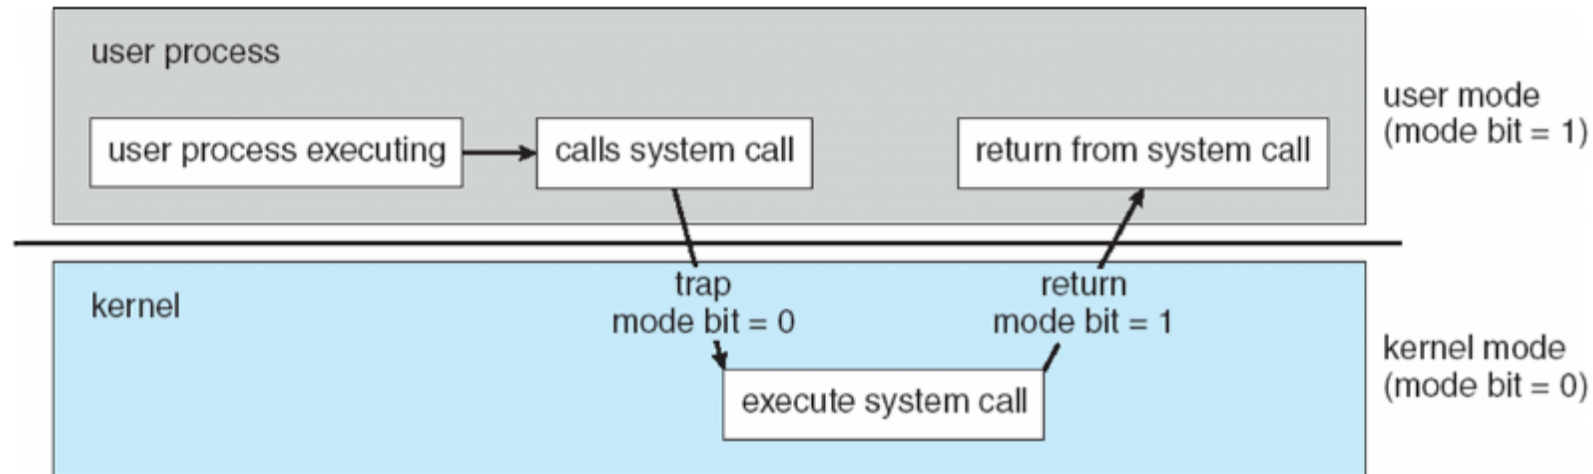
PID	TT	STAT	TIME	COMMAND
1	??	Ss	17:57.36	/sbin/launchd
322	??	Rs	6:29.86	/usr/libexec/logd
323	??	Ss	0:00.19	/usr/libexec/smd
324	??	Ss	0:19.58	/usr/libexec/UserEventAgent (System)

User/Kernel Mode Separation

- **User mode**: restricted, limited operations
 - Processes start in user mode
- **Kernel mode**: privileged, not restricted
 - OS starts in kernel mode
- What if a process wants to perform some restricted operations?
 - **System calls**: Allow the kernel services to provide some functionalities to user programs

User/Kernel Mode Separation

- A process starts in **user mode**
- If it needs to perform a restricted operation, it calls a system call by executing a **trap instruction**.
- The state and registers of the calling process are stored, the system enters **kernel mode**, OS completes the syscall work.
- **Return from syscall**, restore the states and registers of the process, and resume the execution of the process



Process Scheduling

- **Switching Between Processes**
 - Cooperative approach
 - Non-cooperative approach
- **Cooperative approach**
 - Trust process to relinquish CPU to OS through traps
 - » System calls
 - » Illegal operations, e.g., divided by zero
 - **Issue: if no system call**
- **Non-cooperative approach**
 - The OS takes control
 - OS obtains control periodically, e.g., timer interrupter

Summary

- In OS, process is a running program and has an address space
- We use process API to create and manage processes
- Fork() to duplicate a process, exec() to replace the command
- Process scheduling

What's in a process?

- A process consists of:
 - an address space
 - the code for the running program
 - the data for the running program
 - at least one thread
 - » Registers, IP
 - » Floating point state
 - » Stack and stack pointer
 - a set of OS resources
 - » open files, network connections, sound channels, ...
- Today: decompose process from threads of control

Concurrency

- Imagine a web server that handles multiple requests concurrently
 - While waiting for the credit card server to approve a purchase for one client, it could be retrieving the data requested by another client from disk, and assembling the response for a third client from cached information
- Imagine a web client (browser), which might like to initiate multiple requests concurrently
- Imagine a parallel program running on a multiprocessor, which might like to employ “physical concurrency”
 - For example, multiplying a large matrix – split the output matrix into k regions and compute the entries in each region concurrently using k processors


What's needed?

- In each of these examples of concurrency (web server, web client, parallel program):
 - Everybody wants to run the same code
 - Everybody wants to access the same data
 - Everybody has the same privileges
 - Everybody uses the same resources (open files, network connections, etc.)
- But you'd like to have multiple hardware execution states:
 - an execution stack and stack pointer (SP)
 - » traces state of procedure calls made
 - program counter (PC), indicating the next instruction
 - a set of general-purpose processor registers and their values
- Creating multiple processes is inefficient
- Key idea: separate the concept of a process (address space, etc.) from that of a minimal “thread of control” (execution state: PC, etc.)
- This execution state is usually called a **thread**, or sometimes, a **lightweight process**

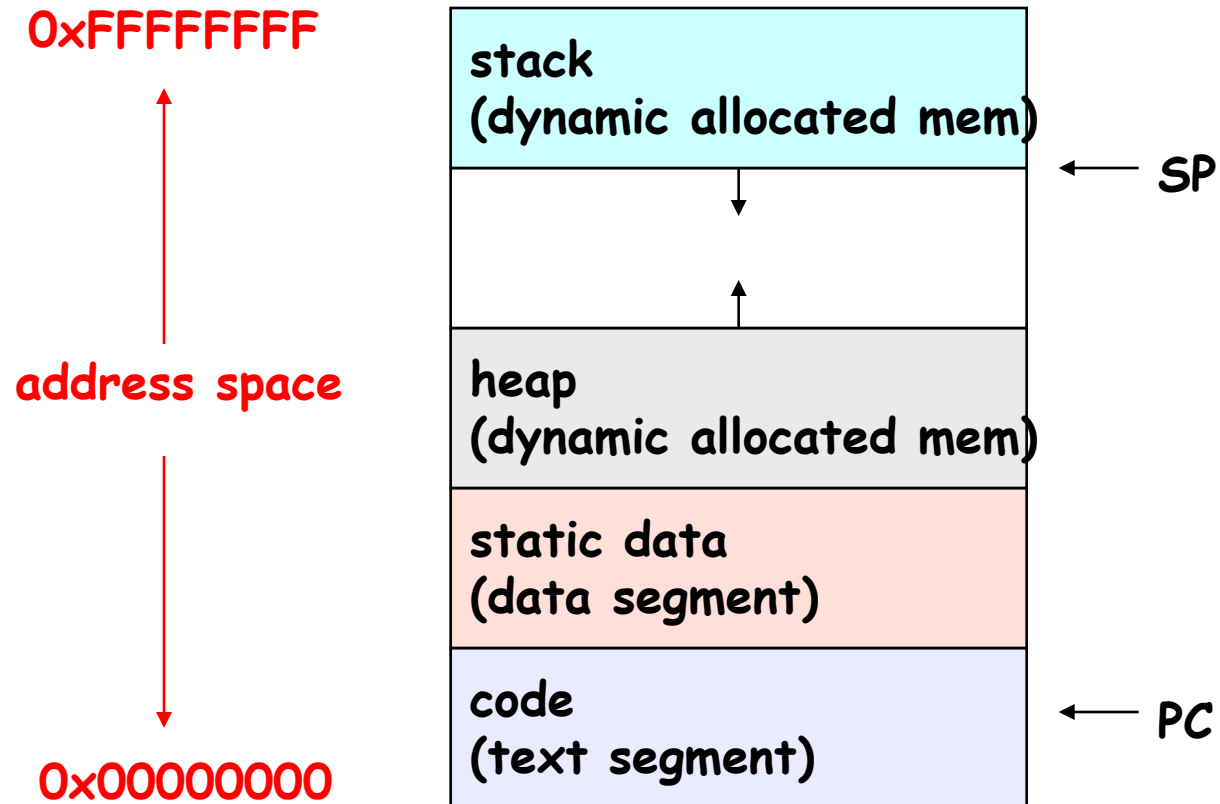
Processes and Threads

- Modern OSes support two entities:
 - the **process**, which defines the address space and general process attributes (such as open files, etc.)
 - the **thread**, which defines a sequential execution stream within a process
- A thread is bound to a single process / address space
 - address spaces, however, can have multiple threads executing within them
 - sharing data between threads is cheap: all see the same address space
 - creating threads is cheap too!
- Threads become the unit of scheduling
 - processes / address spaces are just **containers** in which threads execute

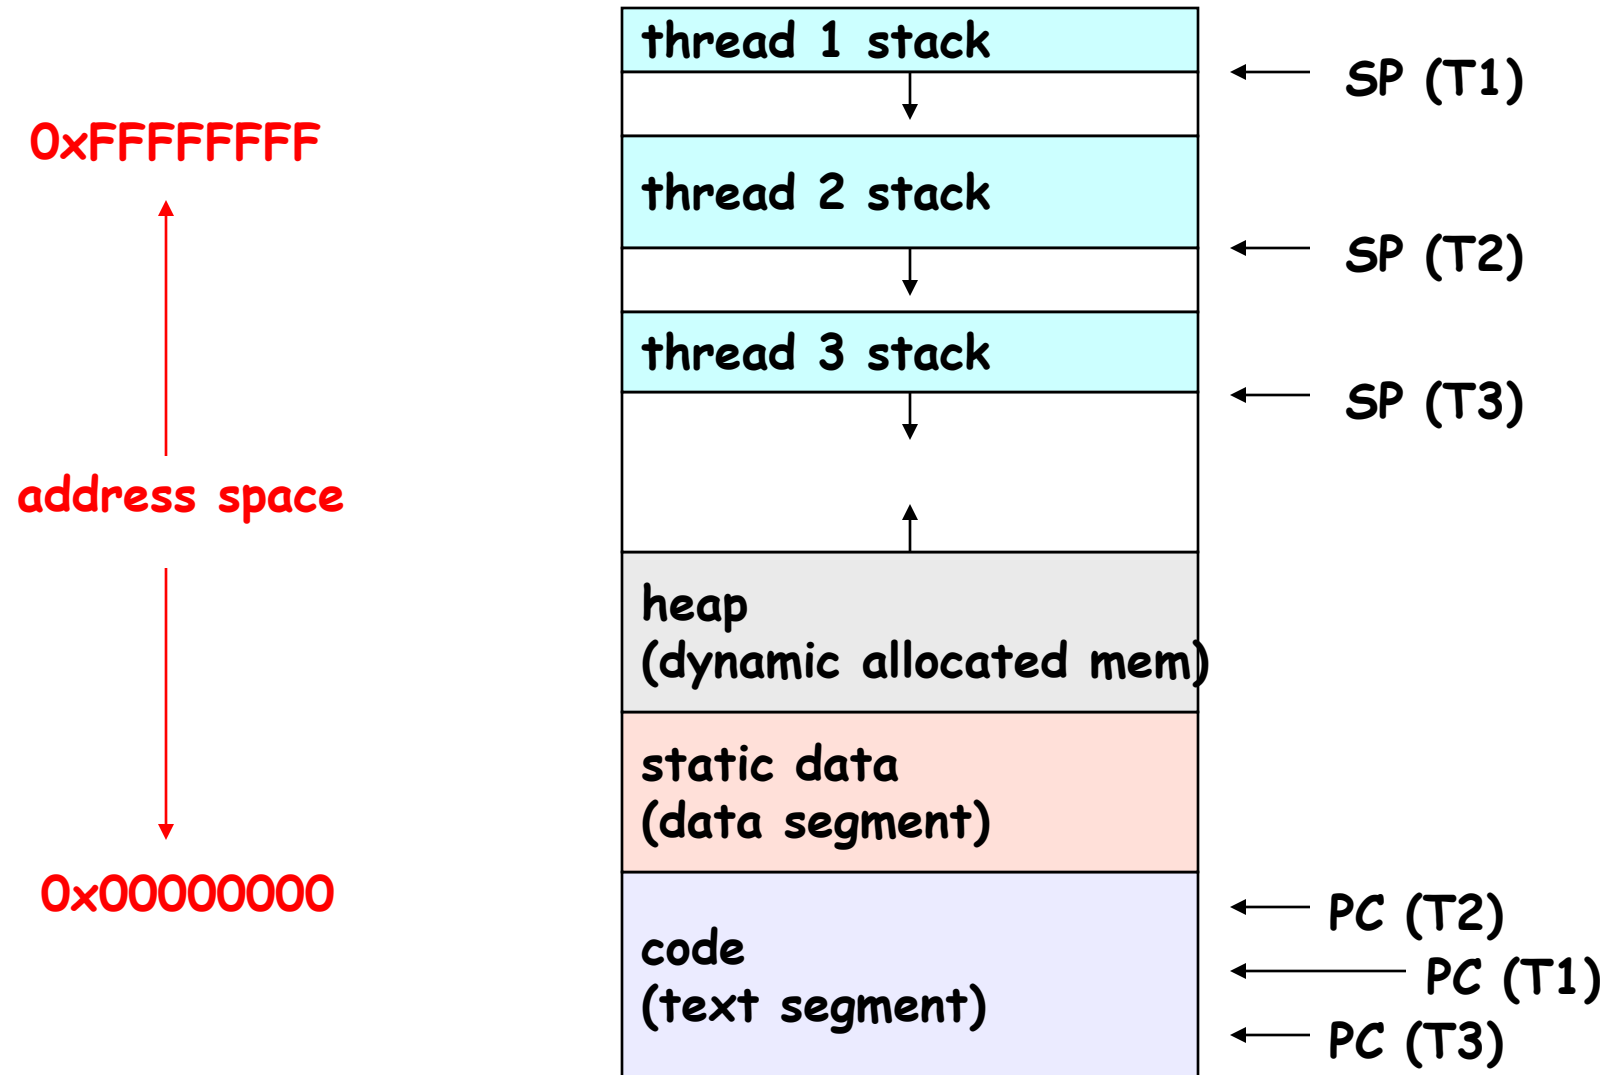
The design space



(old) Process address space



(new) Process address space with threads



Process/thread separation

- Concurrency (multithreading) is useful for:
 - handling concurrent events (e.g., web servers and clients)
 - building parallel programs (e.g., matrix multiply, ray tracing)
 - improving program structure (the Java argument)
- Multithreading is useful even on a uniprocessor
 - even though only one thread can run at a time
- Supporting multithreading – that is, separating the concept of a **process** (address space, files, etc.) from that of a minimal **thread of control** (execution state), is a big win
 - creating concurrency does not require creating new processes
 - “faster / better / cheaper”

“Where do threads come from?”

- The kernel is responsible for creating/managing threads
 - for example, the kernel call to create a new thread would
 - » allocate an execution stack within the process address space
 - » create and initialize a Thread Control Block
 - stack pointer, program counter, register values
 - » stick it on the ready queue
 - we call these **kernel threads**

“Where do threads come from?” (2)

- Threads can also be managed at the user level (that is, entirely from within the process)
 - a library linked into the program manages the threads
 - » because threads share the same address space, the thread manager doesn't need to manipulate address spaces (which only the kernel can do)
 - » threads differ (roughly) only in hardware contexts (PC, SP, registers), which can be manipulated by user-level code
 - » the Linux **thread package** multiplexes user-level threads on top of kernel thread(s), which it treats as “virtual processors”
 - we call these **user-level threads**

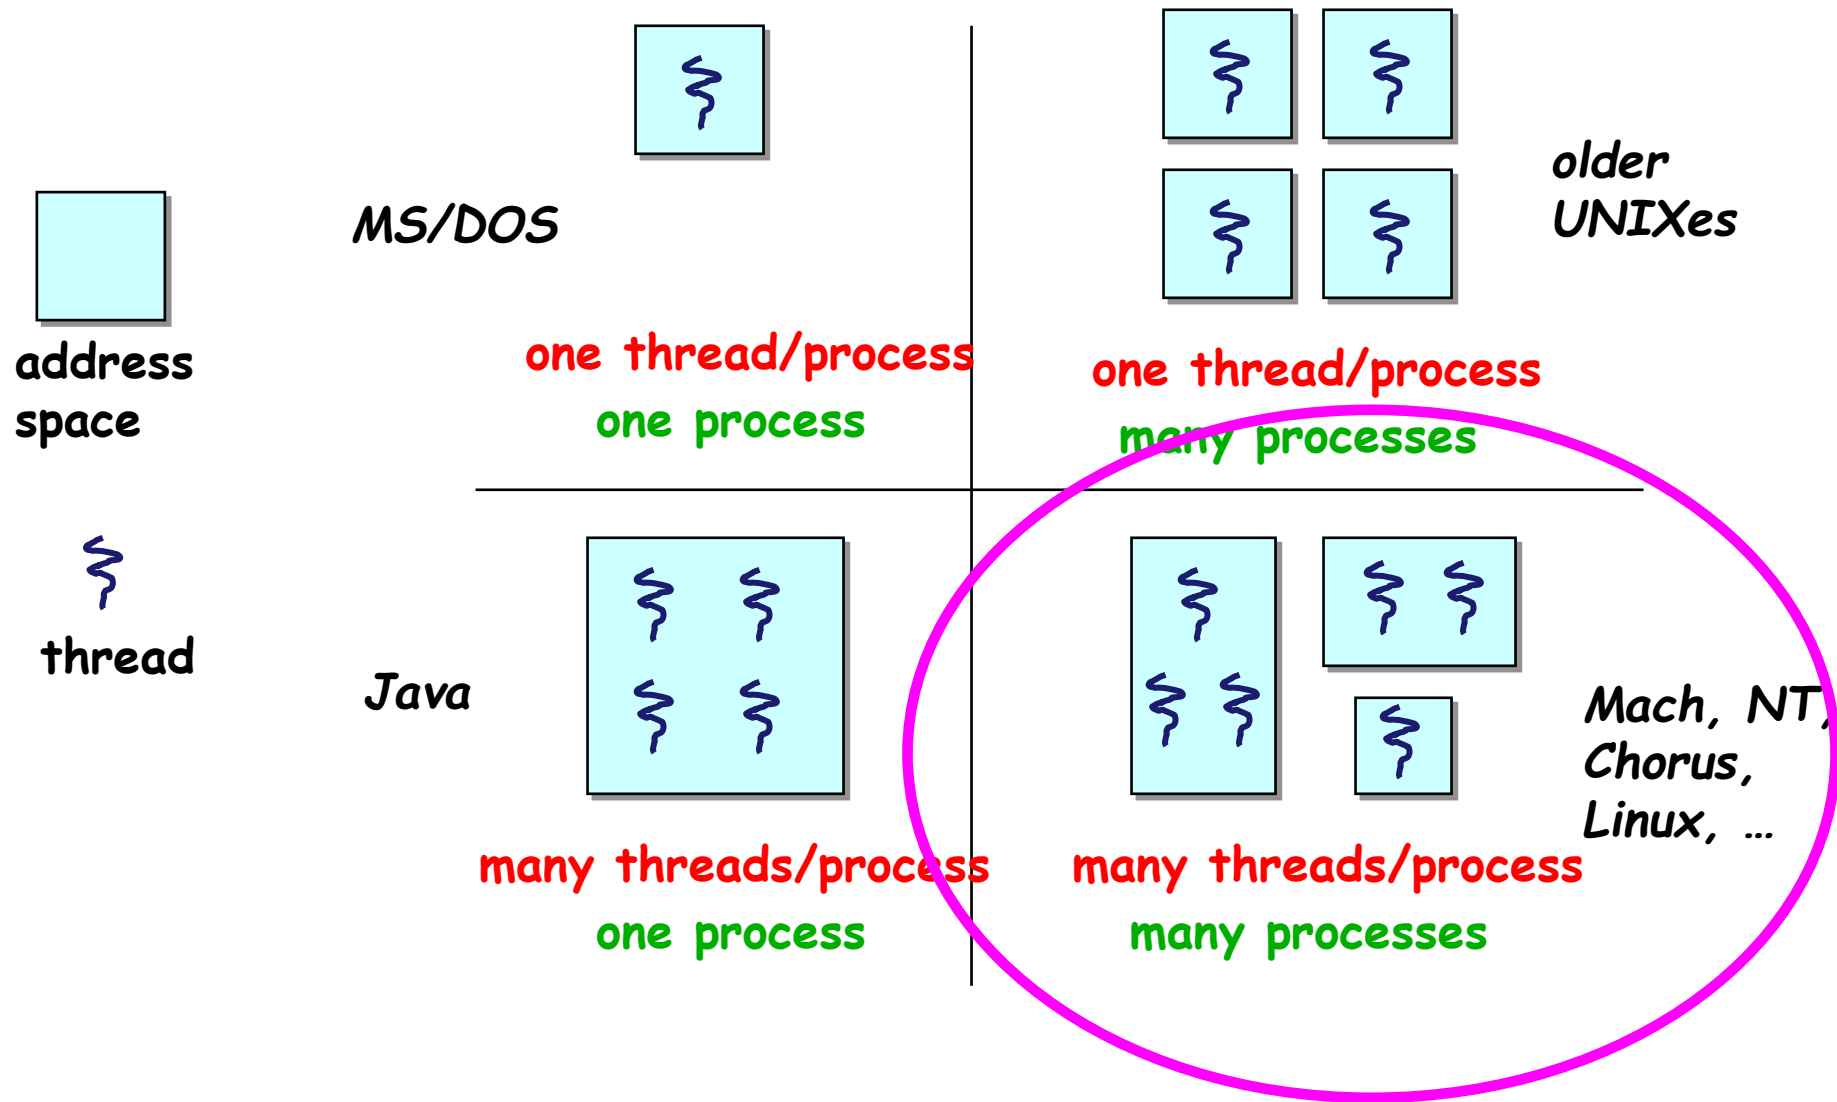
Kernel threads

- OS now manages threads *and* processes
 - all thread operations are implemented in the kernel
 - OS schedules all of the threads in a system
 - » if one thread in a process blocks (e.g., on I/O), the OS knows about it, and can run other threads from that process
 - » possible to overlap I/O and computation **inside** a process
- Kernel threads are cheaper than processes
 - less state to allocate and initialize
- But, they're still expensive for fine-grained use (e.g., orders of magnitude more expensive than a procedure call)
 - thread operations are all system calls
 - » context switch
 - » argument checks
 - must maintain kernel state for each thread


User-level threads

- To make threads cheap and fast, they may be implemented at the user level
 - managed entirely by user-level library, e.g., **libpthreads.a**
- User-level threads are small and fast
 - each thread is represented simply by a PC, registers, a stack, and a small **thread control block** (user-space TCB)
 - creating a thread, switching between threads, and synchronizing threads are done via procedure calls
 - » no kernel involvement is necessary!
 - user-level thread operations can be 10-100x faster than kernel threads as a result


The design space



Kernel threads



User-level threads



User-level thread implementation

- The kernel believes the user-level process is just a normal process running code
 - But, this code includes the thread support library and its associated thread scheduler
- The thread scheduler determines when a thread runs
 - it uses queues to keep track of what threads are doing: run, ready, wait
 - » just like the OS and processes
 - » but, implemented at user-level as a library
- Example implementations of user-level threads
 - Fibers, co-routines

Thread interface

- The POSIX pthreads API:
 - `t = pthread_create(attributes, start_procedure)`
 - » creates a new thread of control
 - » new thread begins executing at `start_procedure`
 - `pthread_cond_wait(condition_variable)`
 - » the calling thread blocks, sometimes called `thread_block()`
 - `pthread_signal(condition_variable)`
 - » starts the thread waiting on the condition variable
 - `pthread_exit()`
 - » terminates the calling thread
 - `pthread_wait(t)`
 - » waits for the named thread to terminate

How to prevent a user-level thread from hogging the CPU?

- Strategy 1: force everyone to cooperate
 - a thread willingly gives up the CPU by calling **yield()**
 - **yield()** calls into the scheduler, which context switches to another ready thread
 - what happens if a thread never calls **yield()**?
- Strategy 2: use preemption
 - scheduler requests that a timer interrupt be delivered by the OS periodically
 - » usually delivered as a UNIX signal (`man signal`)
 - » signals are just like software interrupts, but delivered to user-level by the OS instead of delivered to OS by hardware
 - at each timer interrupt, scheduler gains control and context switches as appropriate

Thread context switch

- Very simple for user-level threads:
 - save context of currently running thread
 - » push machine state onto thread stack
 - restore context of the next thread
 - » pop machine state from next thread's stack
 - return as the new thread
 - » execution resumes at PC of next thread
- This is all done by assembly language
 - it works at the level of the procedure calling convention
 - » thus, it cannot be implemented using procedure calls
 - » e.g., a thread might be preempted (and then resumed) in the middle of a procedure call

What if a thread tries to do I/O?

- The kernel thread is lost for the duration of the (synchronous) I/O operation!
- Could have one kernel thread for each user-level thread
 - no real difference from kernel threads – “common case” operations (e.g., synchronization) would be quick
- Could have a limited-size “pool” of kernel threads “powering” all the user-level threads in the address space
 - the kernel will be scheduling these threads, oblivious to what’s going on at user-level

Summary

- We want multiple threads per address space
- Kernel threads are much more efficient than processes, but they're still not cheap
 - all operations require a kernel call and parameter verification
- User-level threads are:
 - fast
 - great for common-case operations
 - » creation, synchronization, destruction
 - can suffer in uncommon cases due to kernel obliviousness
 - » I/O
 - » preemption of a lock-holder