

CSC 112: Computer Operating Systems

Lecture 6

Real-Time Scheduling II

Department of Computer Science,
Hofstra University

Outline

- Introduction to RTOS and Real-Time Scheduling
- Fixed-Priority Scheduling
- Earliest Deadline First Scheduling
- Least Laxity First (LLF) Scheduling
- Preemptive vs. Non-Preemptive Scheduling
- Multiprocessor Scheduling
- Resource Synchronization Protocols (for Fixed-Priority Scheduling)

Multiprocessor Scheduling

Multiprocessor models


- Identical multiprocessors:
 - each processor has the same computing capacity
- Uniform multiprocessors:
 - different processors have different computing capacities
- Heterogeneous multiprocessors:
 - each (task, processor) pair may have a different computing capacity
- MP scheduling
 - Many NP-hard problems, with few optimal results, mainly heuristic approaches
 - Only sufficient schedulability tests

Multiprocessor Models

Identical multiprocessors: each processor has the same speed

Task T1

Task T2






Multiprocessor Models

Uniform multiprocessors: different processors have different speeds


Task T1

Task T2



Multiprocessor Models


Heterogeneous multiprocessors: each (task, processor) pair may have a different relative speed, due to specialized processor architectures




Global vs partitioned scheduling

- Global scheduling
 - All ready jobs are kept in a common (global) queue; when selected for execution, a job can be dispatched to an arbitrary processor, even after being preempted
- Partitioned scheduling
 - Each task may only execute on a specific processor

**Global scheduling:
Single system-wide queue**



**Partitioned scheduling:
per-processor queues**



Global Scheduling vs. Partitioned Scheduling

- Global Scheduling

- Advantages:

- Runtime load-balancing across cores
 - » More effective utilization of processors and overload management
- Supported by most multiprocessor operating systems
 - » Windows, Linux, MacOS...

- Disadvantages:

- Low schedulable utilization
- Weak theoretical framework

- Partitioned Scheduling

- Advantages:

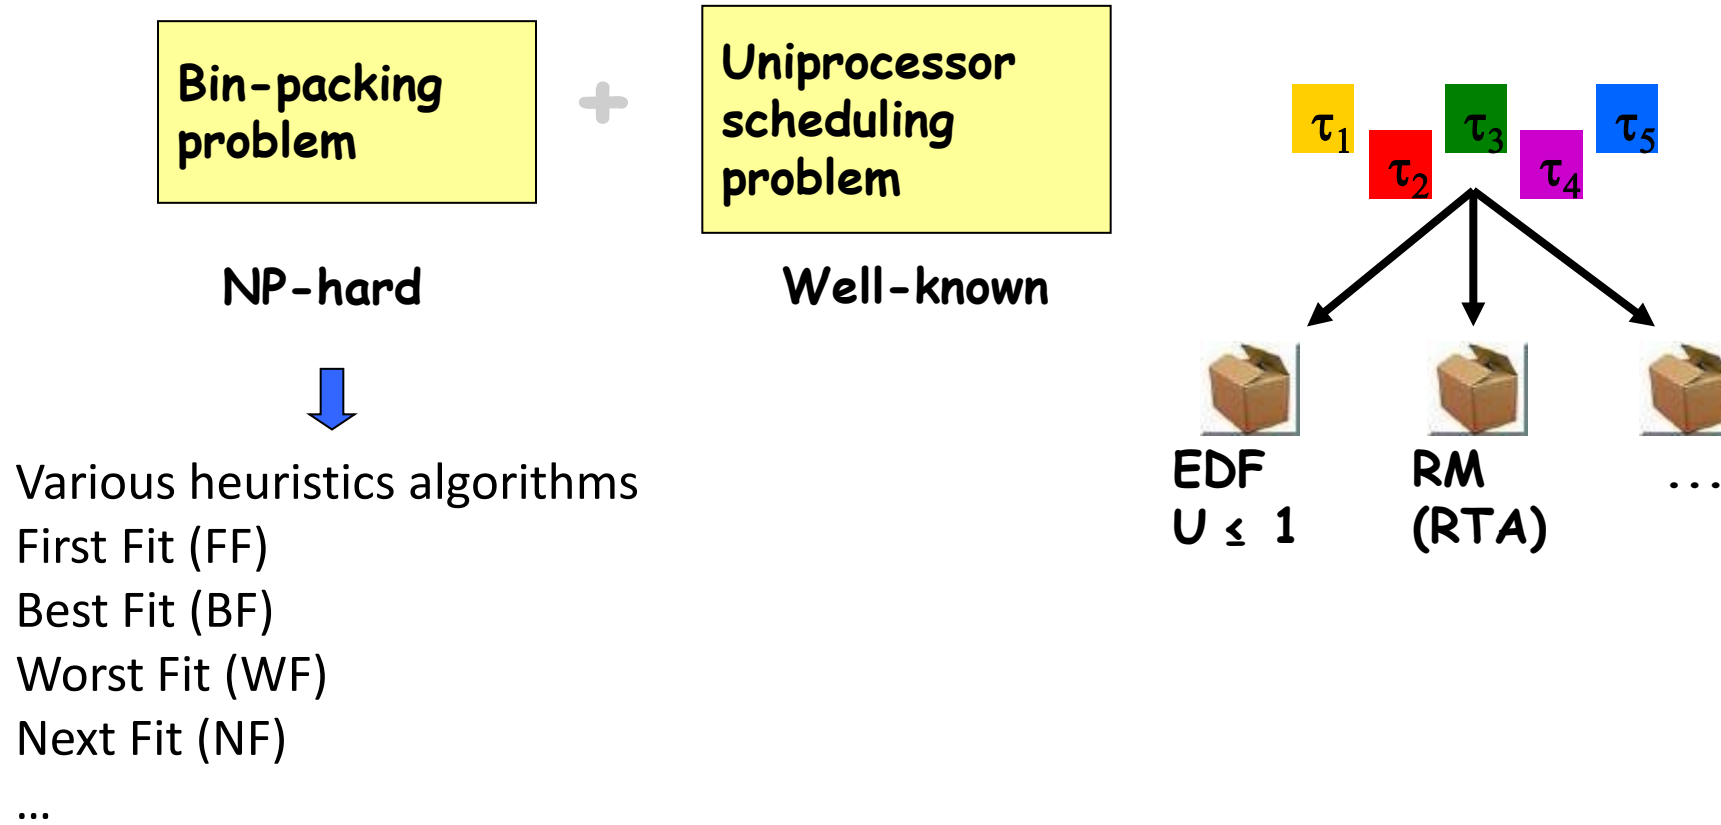
- Mature scheduling framework
- Uniprocessor scheduling theory scheduling are applicable on each core; uniprocessor resource access protocols (PIP, PCP...) can be used
- Partitioning of tasks can be done by efficient bin-packing algorithms

- Disadvantages:

- No runtime load-balancing; surplus CPU time cannot be shared among processors

Partitioned Scheduling

- Scheduling problem reduces to:



Partitioned Scheduling


- Bin-packing algorithms:
 - The problem concerns packing objects of varying sizes in boxes ("bins") with some optimization objective, e.g., minimizing number of used boxes (best-fit), or minimizing the maximum workload for each box (worst-fit)
- Application to multiprocessor scheduling:
 - Bins are represented by processors and objects by tasks
 - The decision whether a processor is "full" or not is derived from a utilization-based feasibility test.
- Since optimal bin-packing is a NP-complete problem, partitioned scheduling is also NP-complete
- Example: Rate-Monotonic-First-Fit (RMFF): (Dhall and Liu, 1978)
 - Let the processors be indexed as 1, 2, ...
 - Assign the tasks to processor in the order of increasing periods (that is, RM order)
 - For each task τ_i , choose the lowest previously-used processor j such that τ_i , together with all tasks that have already been assigned to processor j , can be feasibly scheduled according to the utilization-based schedulability test
 - Additional processors are added if needed

Assumptions for Global Scheduling

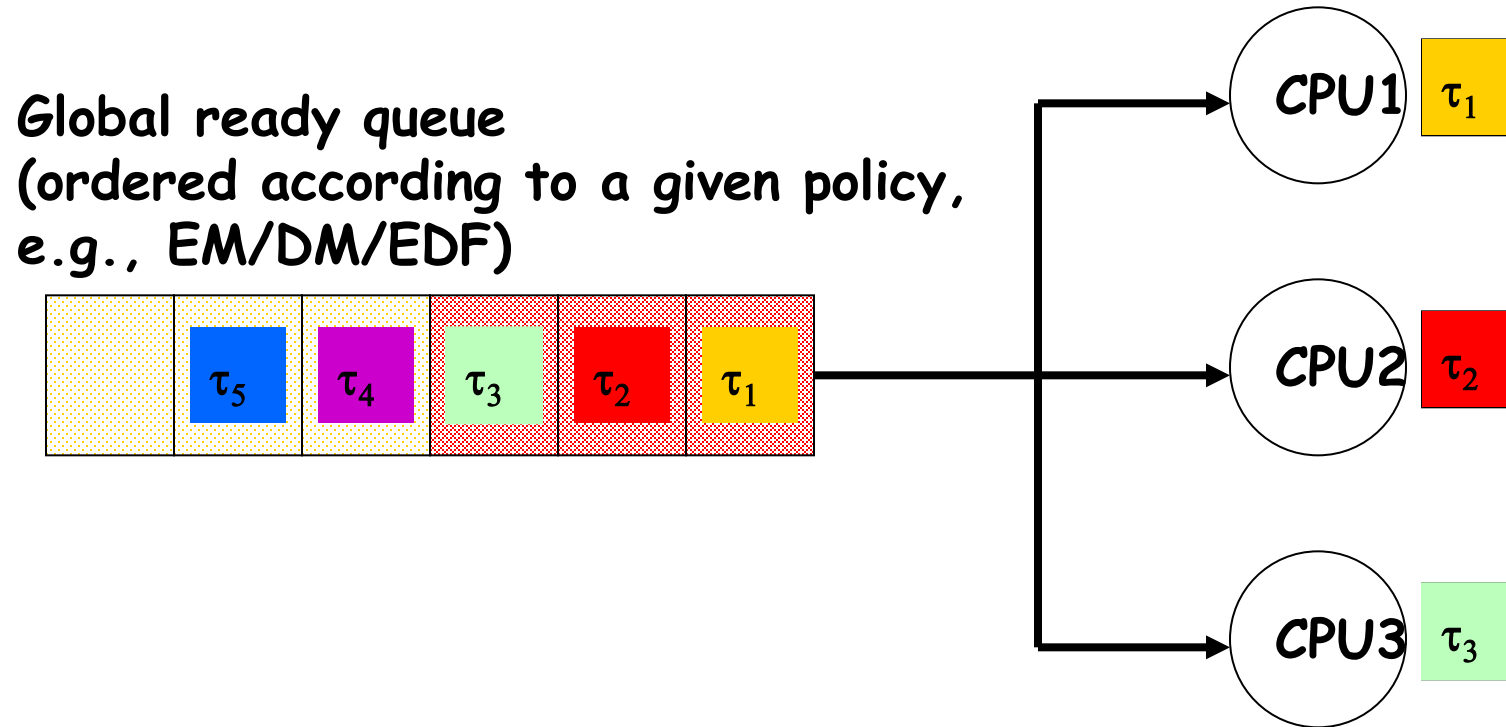
- Identical multiprocessors
- Work-conserving:
 - At each instant, the highest-priority jobs that are eligible to execute are selected for execution upon the available processors
 - No processor is ever idle when the ready queue is non-empty
- Preemption and Migration support
 - A preempted task can resume execution on a different processor with 0 overhead, as cost of preemption/migration is integrated into task WCET
- No job-level parallelism
 - the same job cannot be *simultaneously* executed on more than one processor, i.e., we do not consider parallel programs that can run on multiple processors in parallel

Source of Difficulty

- The “no job-level parallelism” assumption leads to difficult scheduling problems
- “The simple fact that a task can use only one processor even when several processors are free at the same time adds a surprising amount of difficulty to the scheduling of multiple processors” [Liu’69]

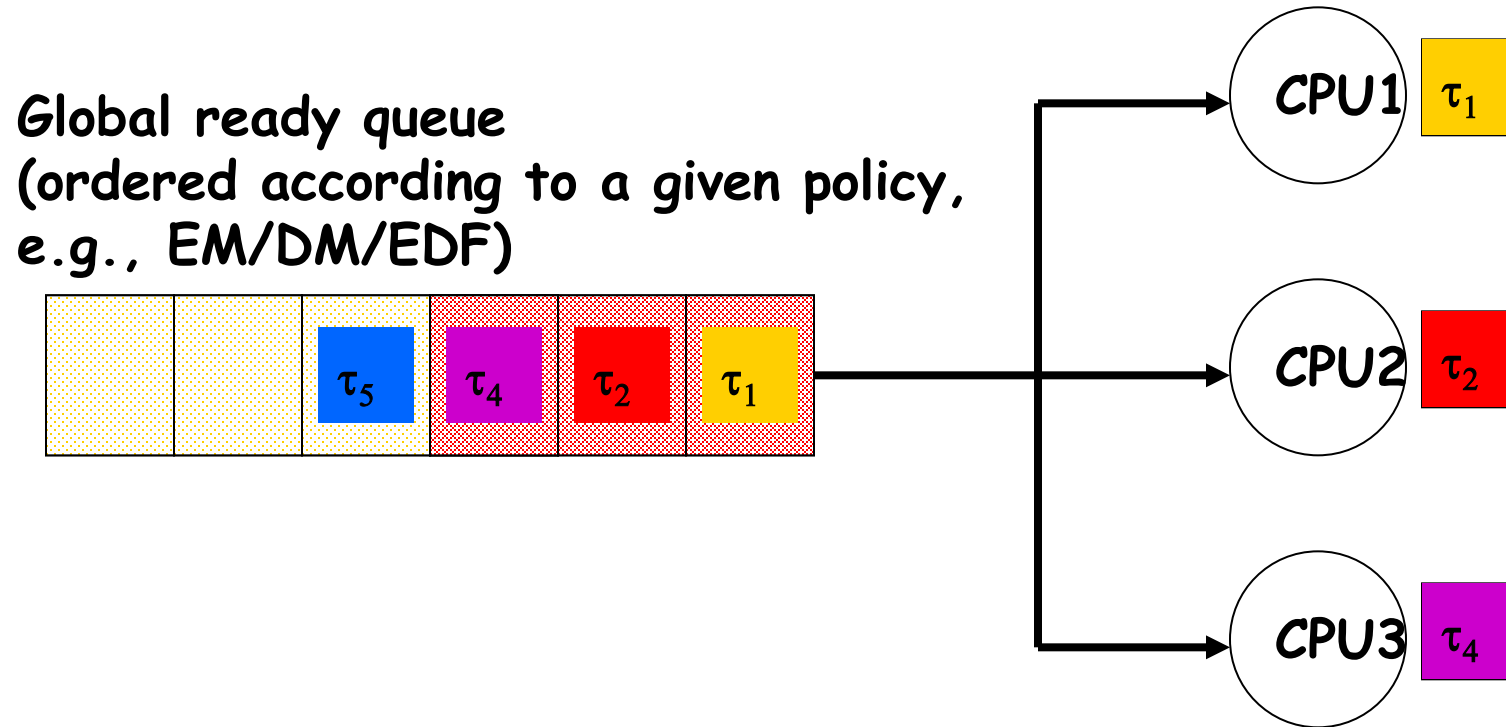


Global scheduling example



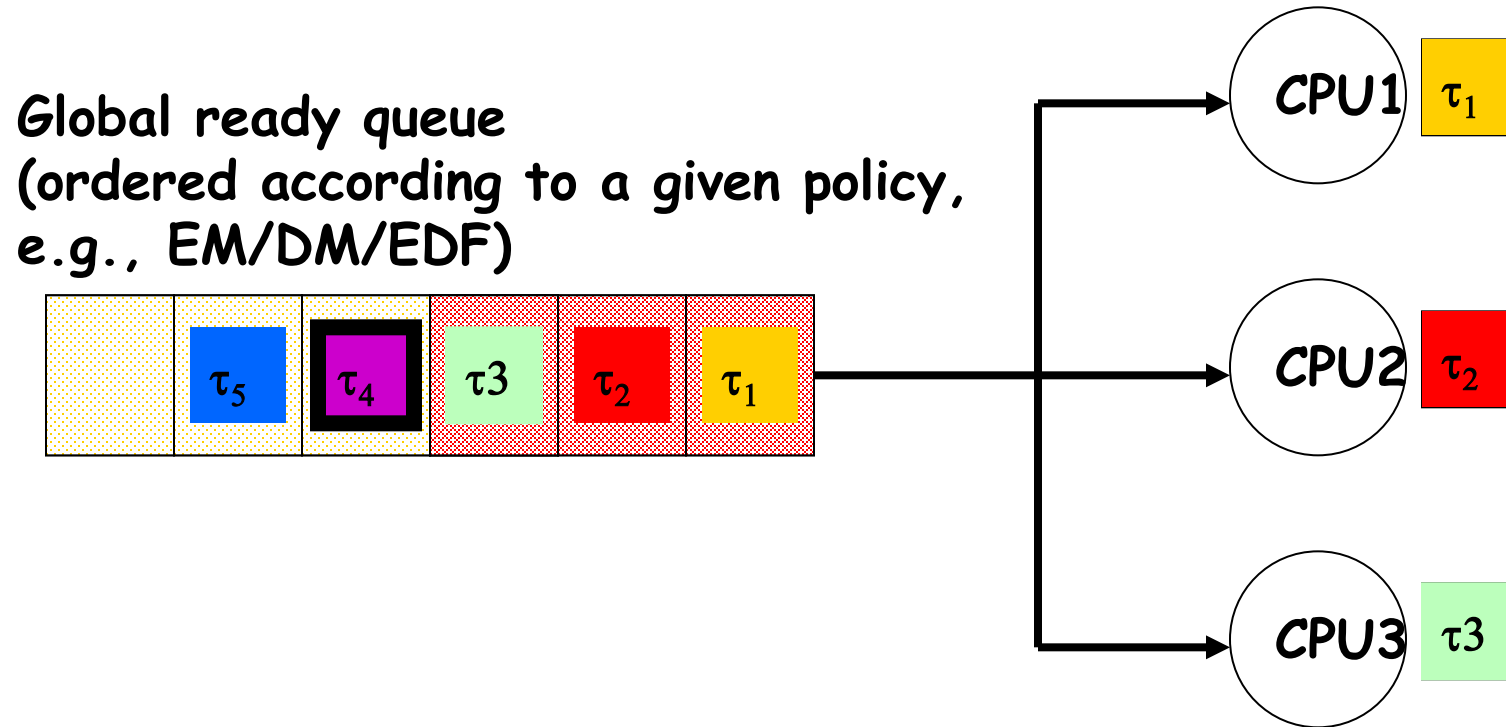
The first m jobs in the queue are scheduled upon the m CPUs

Global scheduling example



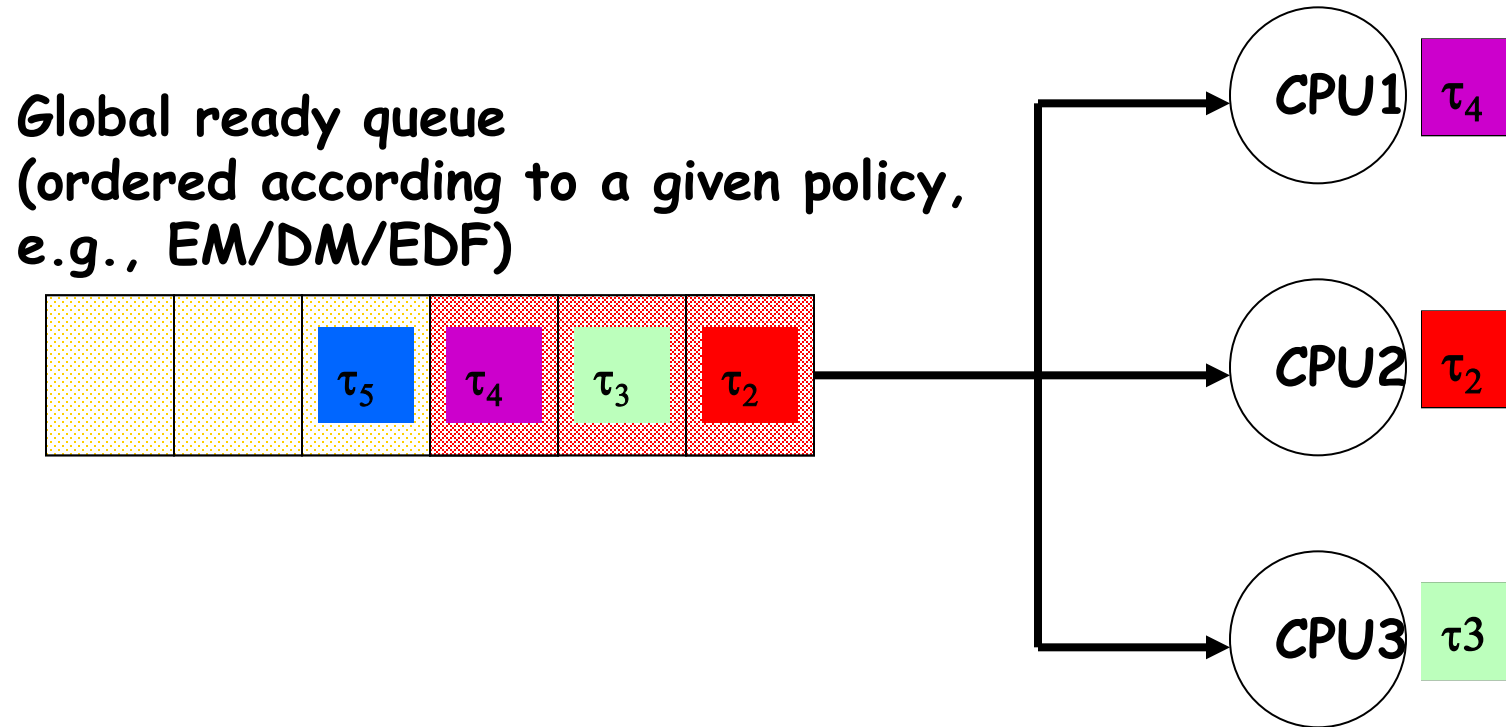
When a job τ_3 finishes its execution, the next job in the queue τ_4 is scheduled on the available CPU

Global scheduling example



When a new higher-priority job τ_3 arrives in its next period T_3 , it preempts the job with lowestpriority τ_4 among the executing ones

Global scheduling example



When another job τ_1 finishes its execution, the preempted job τ_4 can resume its execution. Net effect: τ_4 "migrated" from CPU3 to CPU1


Global vs. Partitioned

- Global (work-conserving) and partitioned scheduling algorithms are incomparable:
 - There are tasksets that are schedulable with a global scheduler, but not with a partitioned scheduler, and vice versa.

Global vs Partitioned (FP) Scheduling

Task	T=D	C	Prio
T1	2	1	H
T2	3	2	M
T3	3	2	L

- A taskset schedulable with global scheduling, but not partitioned scheduling. System utilization $U = \frac{1}{2} + \frac{2}{3} + \frac{2}{3} = 1.83$
- Global FP scheduling is schedulable with priority assignment $p_1 > p_2 > p_3$ (or $p_2 > p_1 > p_3$)
- Partitioned scheduling is unschedulable, since assigning any two tasks to the same processor will cause that processor's utilization to exceed 1, so the bin-packing problem has no feasible solution



A feasible execution trace under global scheduling

Global vs Partitioned (FP) Scheduling


Task	T=D	C	Prio
T1	6	4	4(H)
T2	12	7	3
T3	12	4	2
T4	24	10	1(L)

- A taskset schedulable with partitioned scheduling, but not global scheduling. System utilization $U = \frac{4}{6} + \frac{7}{12} + \frac{4}{12} + \frac{10}{24} = 2.0$, hence the two processors must be fully utilized with no possible idle intervals

- Partitioned FP scheduling with RM priority assignment ($p_1 > p_2 > p_3 > p_4$) is schedulable. T1, T3 assigned to Processor 1; T2, T4 assigned to Processor 2. Both processors have utilization 1.0, and harmonic task periods


- Global FP scheduling with RM priority assignment $p_1 > p_2 > p_3 > p_4$ is unschedulable. Compared to partitioned scheduling, the difference is at time 7, when T3 (with higher priority than T4) runs on Processor 2. This causes idle intervals on Processor 1 [10,12] and [22,24], since only one task T4 is ready during these time intervals. Since taskset $U = 2.0$ on 2 processors, any idle interval will cause the taskset to be unschedulable

At time 7, T4 runs on Processor 2



A feasible execution trace under partitioned scheduling

At time 7, T3 runs on Processor 2



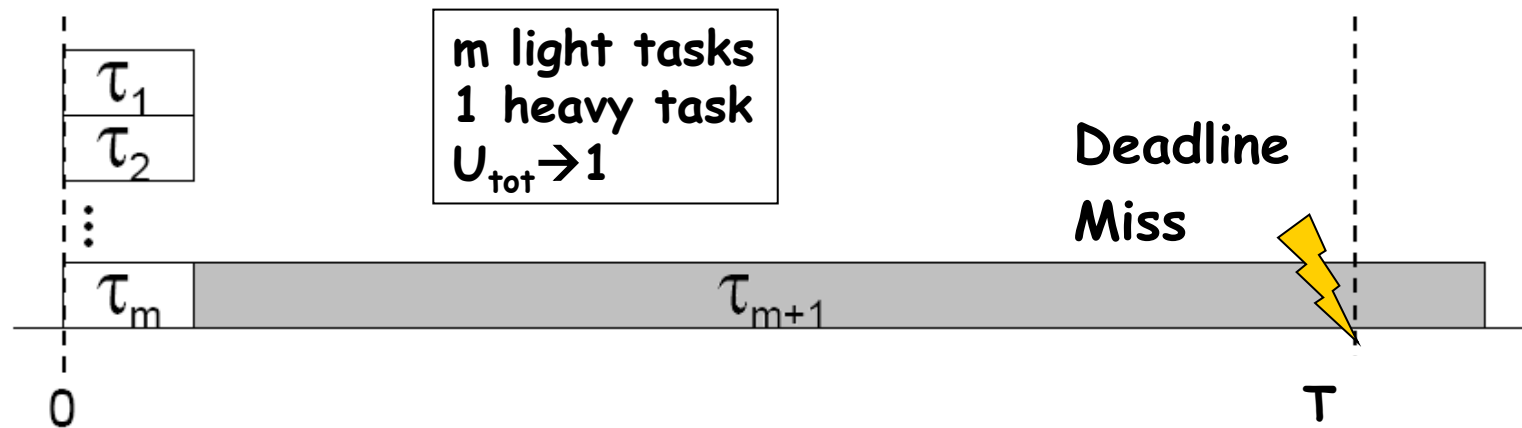
An infeasible execution trace under global scheduling

Difficulties of Global Scheduling

- Dhall's effect
 - With RM, DM and EDF, some low-utilization task sets can be unschedulable regardless of how many processors are used.
- Scheduling anomalies
 - Decreasing task execution time or increasing task period may cause deadline misses
- Hard-to-find worst-case
 - The worst-case does not always occur when a task arrives at the same time as all its higher-priority tasks
- Dependence on relative priority ordering (omitted)
 - Changing the relative priority ordering among higher-priority tasks may affect schedulability for a lower-priority task

Dhall's effect

- Global RM/DM/EDF can fail at very low utilization
- Example: m processors, $n=m+1$ tasks. Tasks τ_1, \dots, τ_m are light tasks, with small $C_i = 1$, $T_i = D_i = T - 1$; Task τ_{m+1} is a heavy task, with large $C_i = T$, $T_i = D_i = T$. $T > 1$ is some constant value
- For global RM/DM/EDF, Task τ_{m+1} has lowest priority, so τ_1, \dots, τ_m must run on m processors starting at time 0, causing τ_{m+1} to miss its deadline
- One solution: assign higher priority to heavy tasks
 - If heavy task τ_{m+1} is assigned the highest priority, then it runs from time 0 to T and meets its deadline; The light tasks can run on other processors and meet their deadlines as well



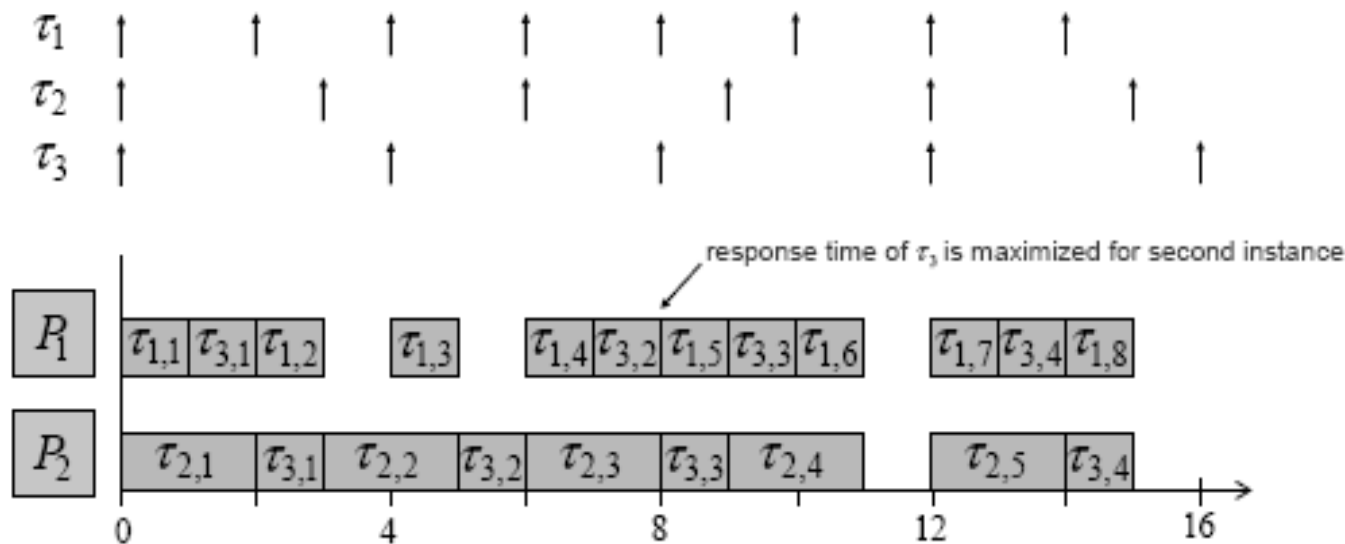
Hard-to-Find Worst-Case

- For uniprocessor scheduling, the worst case occurs when all tasks are initially released at time 0 simultaneously, called the critical instant (recall Slide [Response Time Analysis \(RTA\)](#))
- This is no longer true for multiprocessor scheduling, as the worst-case interference for a task does not always occur at time 0, when all tasks are initially released at time 0 simultaneously
 - Response time for task τ_3 is maximized for its 2nd job $\tau_{3,2}$ (8-4=4), which does not arrive at the same time as its higher priority tasks; not for its 1st job $\tau_{3,1}$ (3-0=3), which arrives at the same time as its higher priority tasks

Hard-to-find critical instant:

(RM scheduling)

$$\begin{aligned} \tau_1 &= \{C_1 = 1, T_1 = 2\} \\ \tau_2 &= \{C_2 = 2, T_2 = 3\} \\ \tau_3 &= \{C_3 = 2, T_3 = 4\} \end{aligned}$$




MP Scheduling Anomalies

- Decrease in processor demand (decreasing task execution time or increasing task period) may cause deadline misses!
- **Anomaly 1**
 - Decrease in processor demand from higher-priority tasks can *increase the interference on a lower-priority task because of change in the time when the tasks execute*
- **Anomaly 2**
 - Decrease in processor demand of a task *negatively affects the task itself* because change in the task arrival times cause it to suffer more interference


Scheduling Anomaly Example 1

- Three tasks on two processors under global scheduling
- With Task a's period $T_a = 3$, system utilization $\sum U_i = 1.83$. WCRT of task c is $R_c = 12 \leq D_c = 12$. $R_c = C_c + I_c = 8 + I_c$, where $I_c = 2 + 1 + 1 = 4$ is interference by higher priority tasks a and b. (Task c experiences inference when both processors are busy executing higher priority tasks a and b.) Task c is schedulable but saturated, as any increase in its WCET or interference would make it unschedulable.
- With Task a's period $T_a = 4$, system utilization $\sum U_i = 1.67$ is reduced. But WCRT of task c increases: $R_c = 14 > D_c = 12$. $R_c = 8 + I_c$ where $I_c = 2 + 2 + 2 = 6$, since execution segments of tasks a and b on two processors are aligned in time, thus causing more interference to task c

Task	T=D	C	Util	Prio
a	3	2	0.67	H
b	4	2	0.5	M
c	12	8	0.67	L



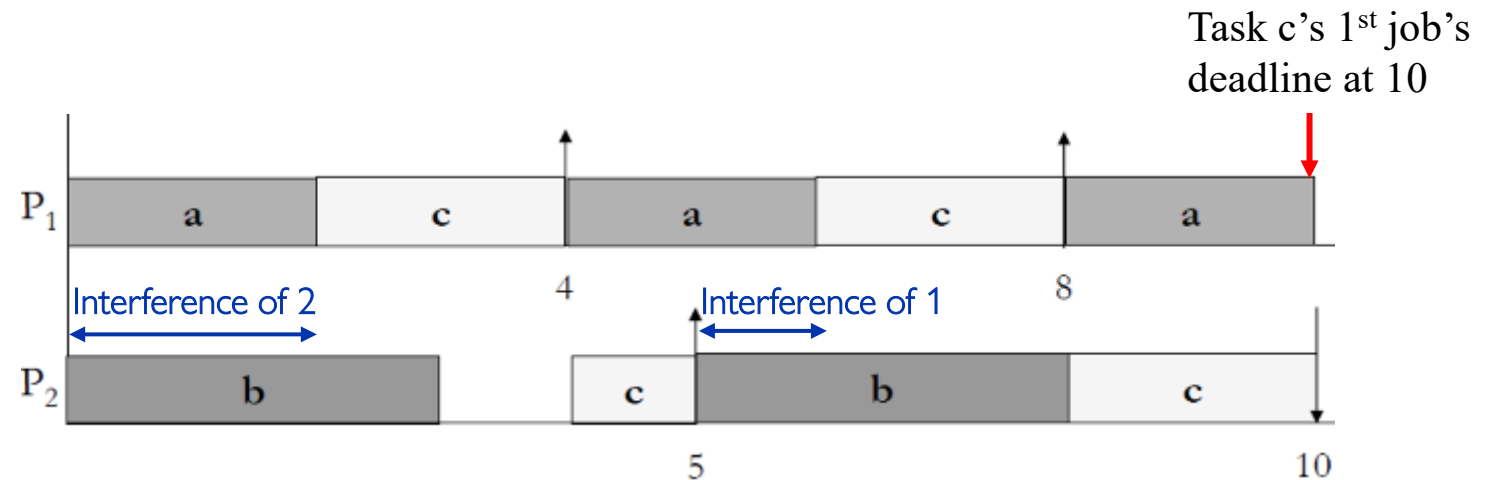
Task	T=D	C	Util	Prio
a	4	2	0.5	H
b	4	2	0.5	M
c	12	8	0.67	L



Scheduling Anomaly Example 2

- Three tasks on two processors under global scheduling
- With Task c's period $T_c = 10$, system utilization $\sum U_i = 1.8$. WCRT of task c is $R_c = 10 \leq D_c = 10$. $R_c = C_c + I_c = 7 + 3 = 10$, where $I_c = 2 + 1 = 3$ is interference by higher priority tasks a and b. Its 1st job meets its deadline at time 10. This schedule repeats in future periods, hence task c is schedulable but saturated, as any increase in its WCET or interference would make it unschedulable.


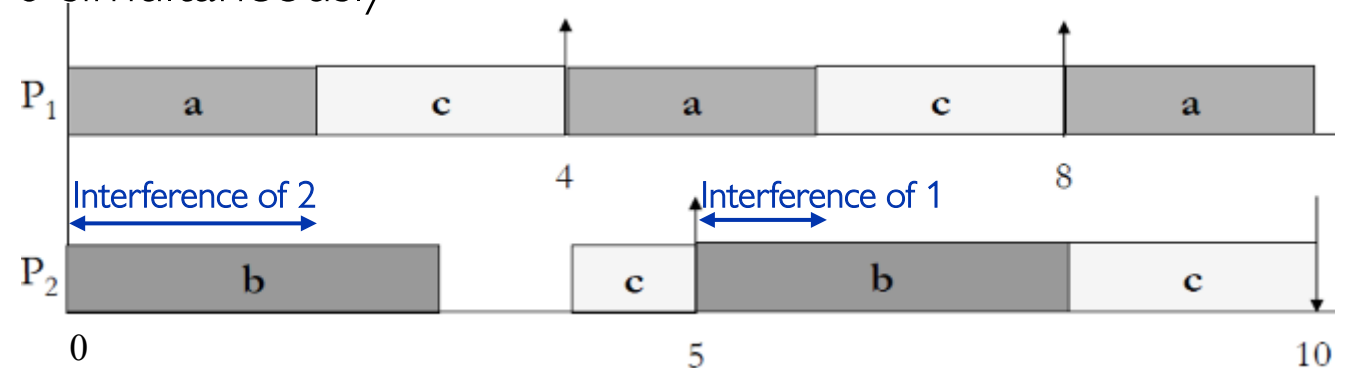
Task	T=D	C	Util	Prio
a	4	2	0.5	H
b	5	3	0.6	M
c	10	7	0.7	L



Scheduling Anomaly Example 2

- With Task c's period $T_c = 11$, system utilization $\sum U_i = 1.74$ is reduced. WCRT of task c is $R_c = 12 > D_c = 10$. Its 1st job has response time $C_c + I_c = 7 + 3 = 10 \leq D_c = 11$, where $I_c = 2 + 1 = 3$, but this is not task c's WCRT.
- Its 2nd job has response time $C_c + I_c = 7 + 5 = 12 > D_c = 11$, where $I_c = 1 + 2 + 2 = 5$. The 2nd job finishes at time $11+12=23$, and misses its deadline at time 22.
- Another example where the worst-case interference for task c does NOT occur at time 0, when all tasks are initially released at time 0 simultaneously


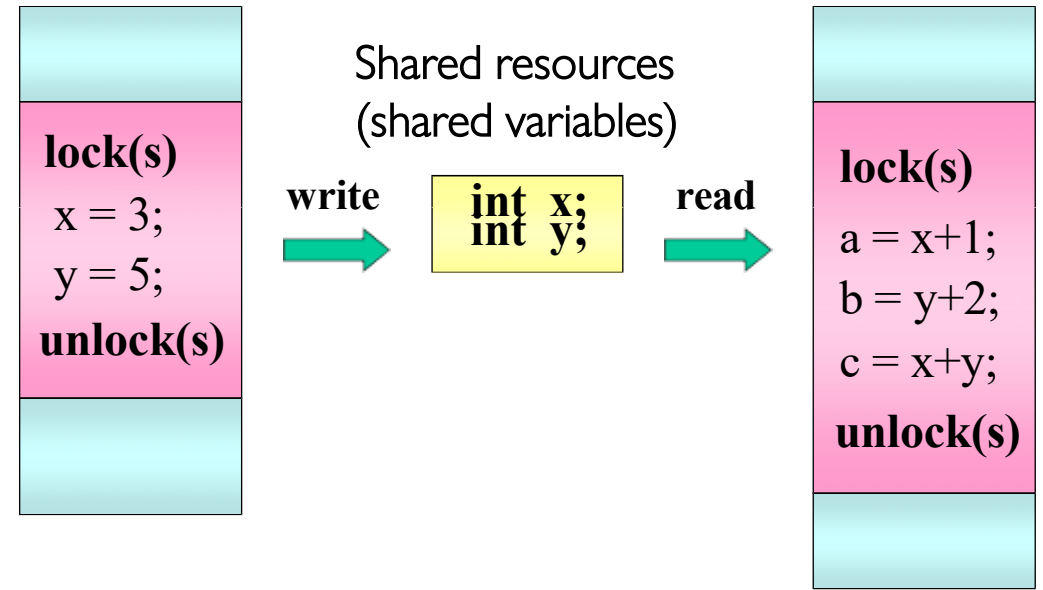
Task	T=D	C	Util	Prio
a	4	2	0.5	H
b	5	3	0.6	M
c	11	7	0.64	L



Resource Synchronization Protocols (for Fixed-Priority Scheduling)

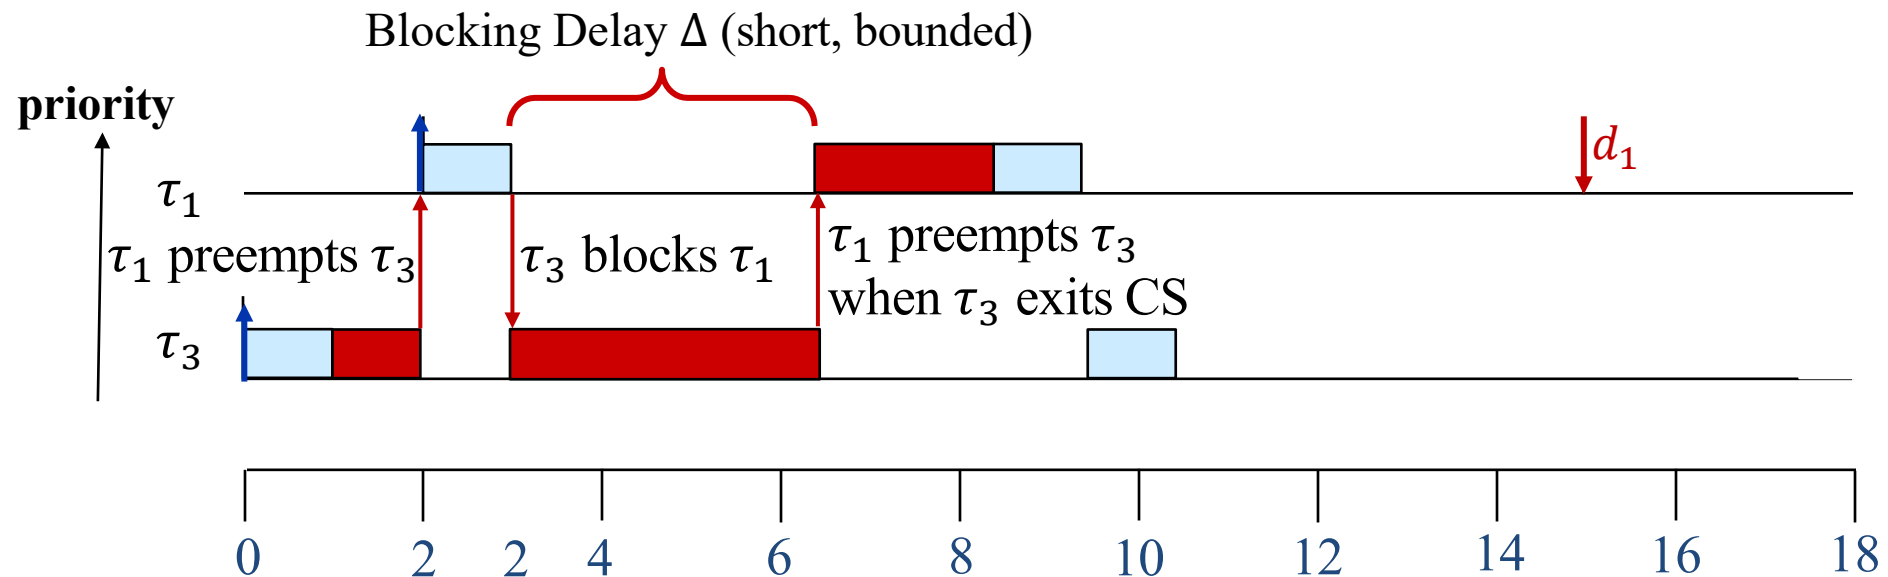
Resource Sharing

- When two tasks access shared resources (variables), mutexes (or binary semaphores) are used to protect critical sections. Each Critical Section (CS) must begin with `lock(s)` and end with `unlock(s)`
- A task waiting for a shared resource is blocked on that resource. Otherwise, it proceeds by entering the critical section and holds the resource
- Tasks blocked on the same resource are kept in a queue. When a running task invokes `lock(s)` when `s` is locked, it enters a waiting state, until another task unlocks `s`



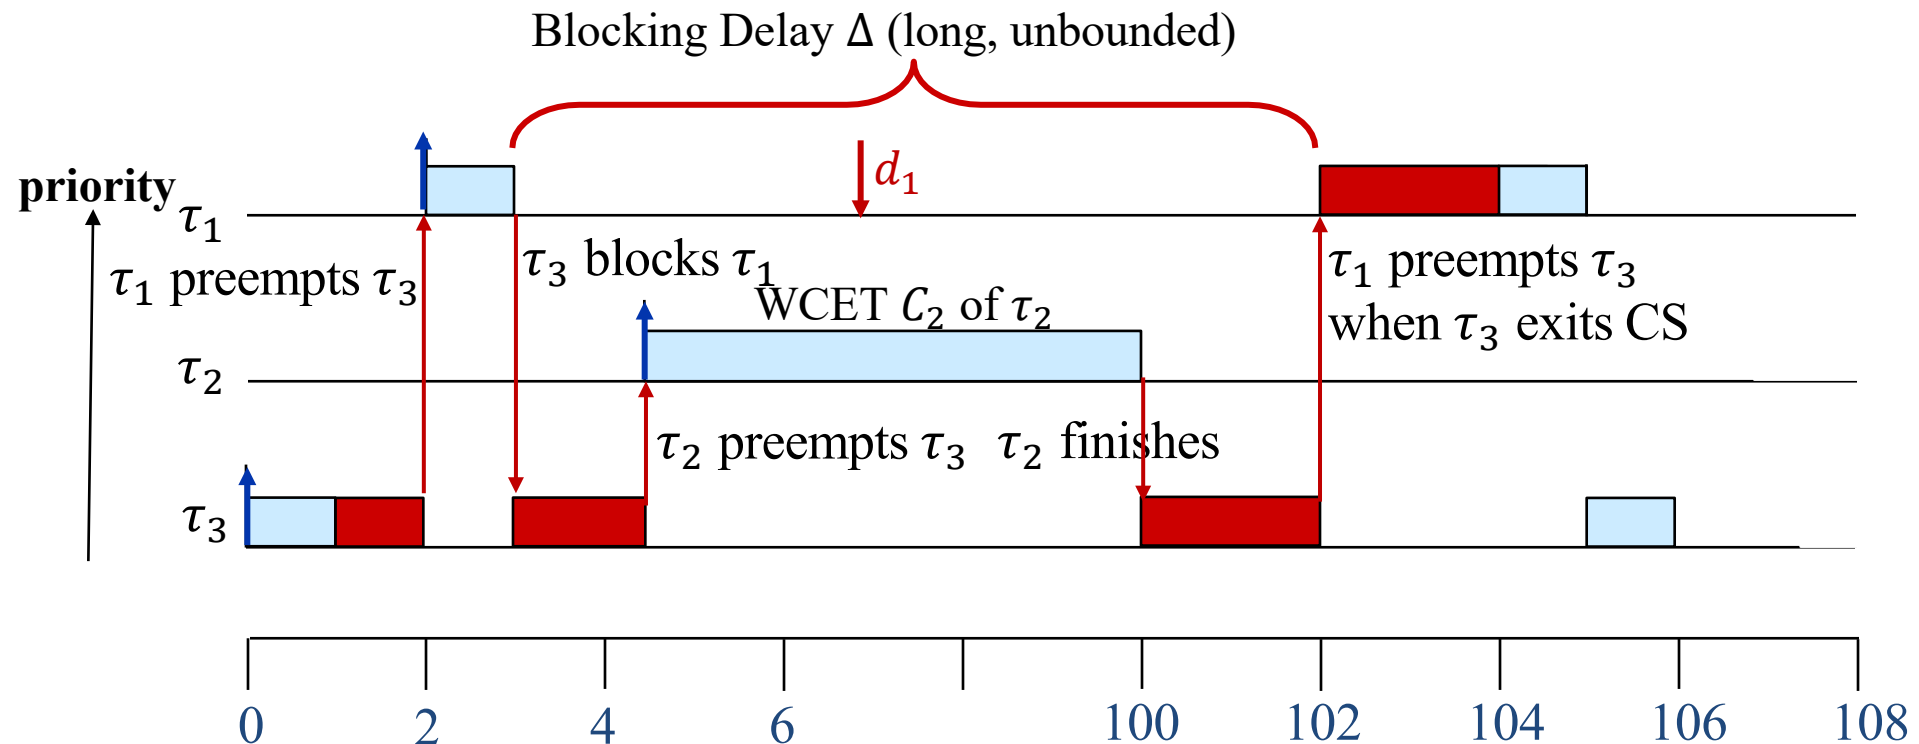
Blocking Delay

- Lower-priority (LP) tasks can cause **blocking delay** to higher-priority (HP) tasks due to resource sharing
 - HP tasks may cause **preemption delay** to LP tasks, but not blocking delay
- Example: Two tasks τ_1, τ_2 with priority ordering $P_1 > P_2$. They both require semaphore s
- If HP task τ_1 tries to lock s that is held by LP task τ_3 , τ_1 blocks until τ_3 unlocks s , so τ_1 experiences a blocking delay Δ .
 - Since CS is typically very short, it seems this blocking time delay Δ is bounded by the longest critical section in lower-priority tasks?
- No, blocking delay may be unbounded!



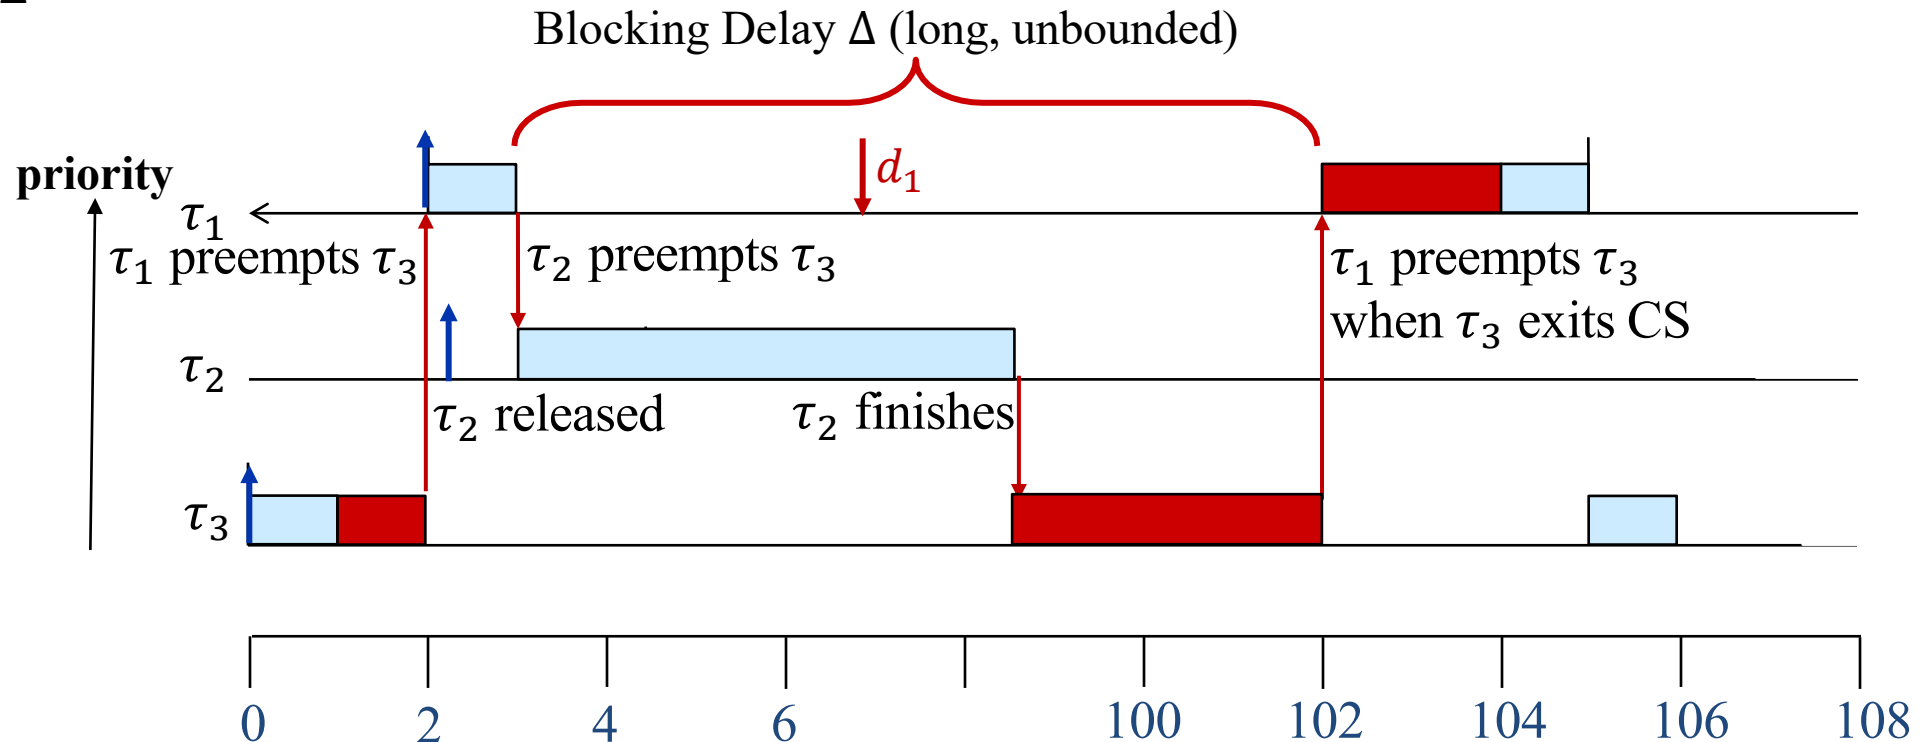
Priority Inversion I

- Three tasks τ_1, τ_2, τ_3 with priority ordering $P_1 > P_2 > P_3$. τ_1, τ_2 both require semaphore s , and τ_3 does not require any semaphore
- $t=1$: LP task τ_3 locks s and enters CS
- $t=2$: HP task τ_1 is released and preempts τ_3
- $t=3$: HP task τ_1 tries to lock s , but gets blocked by τ_3 holding s
- $t=4.2$: Medium Priority (MP) task τ_2 is released and preempts τ_3
- $t=100$: MP task τ_2 finishes execution after running for its WCET C_2 ; τ_3 resumes execution in CS
- $t=102$: LP task τ_3 unlocks s ; HP task τ_1 preempts τ_3 and finally locks s , after experiencing a long, unbounded blocking delay Δ , and misses its deadline d_1




Priority Inversion II

- (This scenario is more likely than previous one, as MP task τ_2 may be released anytime during τ_1 's first execution segment $t \in [2, 3]$, not during τ_3 's CS)
- $t=1$: LP task τ_3 locks s and enters CS
- $t=2$: HP task τ_1 is released and preempts τ_3
- $t \in [2, 3]$: MP task τ_2 is released, but cannot run since HP task τ_1 is running
- $t=3$: HP task τ_1 tries to lock s , but gets blocked by τ_3 holding s ; MP task τ_2 starts running
- $t=98.5$: MP task τ_2 finishes execution after running for its WCET C_2 ; τ_3 resumes execution in the CS
- $t=102$: LP task τ_3 unlocks s ; HP task τ_1 preempts τ_3 and finally locks s , after experiencing a long, unbounded blocking Delay Δ



Deadlocks

- Classic deadlock scenario: Two tasks τ_1 and τ_2 acquire two semaphores s_1 , s_2 in opposite order:
 - HP task τ_1 enters blue CS A before pink CS B: ...lock(s_1)...lock(s_2)... unlock(s_2)...unlock(s_1)...
 - LP task τ_2 enters pink CS B before blue CS A: ...lock(s_2)...lock(s_1)... unlock(s_1)...unlock(s_2)...
 - LP task τ_2 runs first and locks s_2
 - HP task τ_1 starts running and locks s_1 , then tries to lock s_2 , gets blocked by τ_2
 - τ_2 starts running and tries to lock s_1 , but τ_1 holds s_2 . Deadlocked!

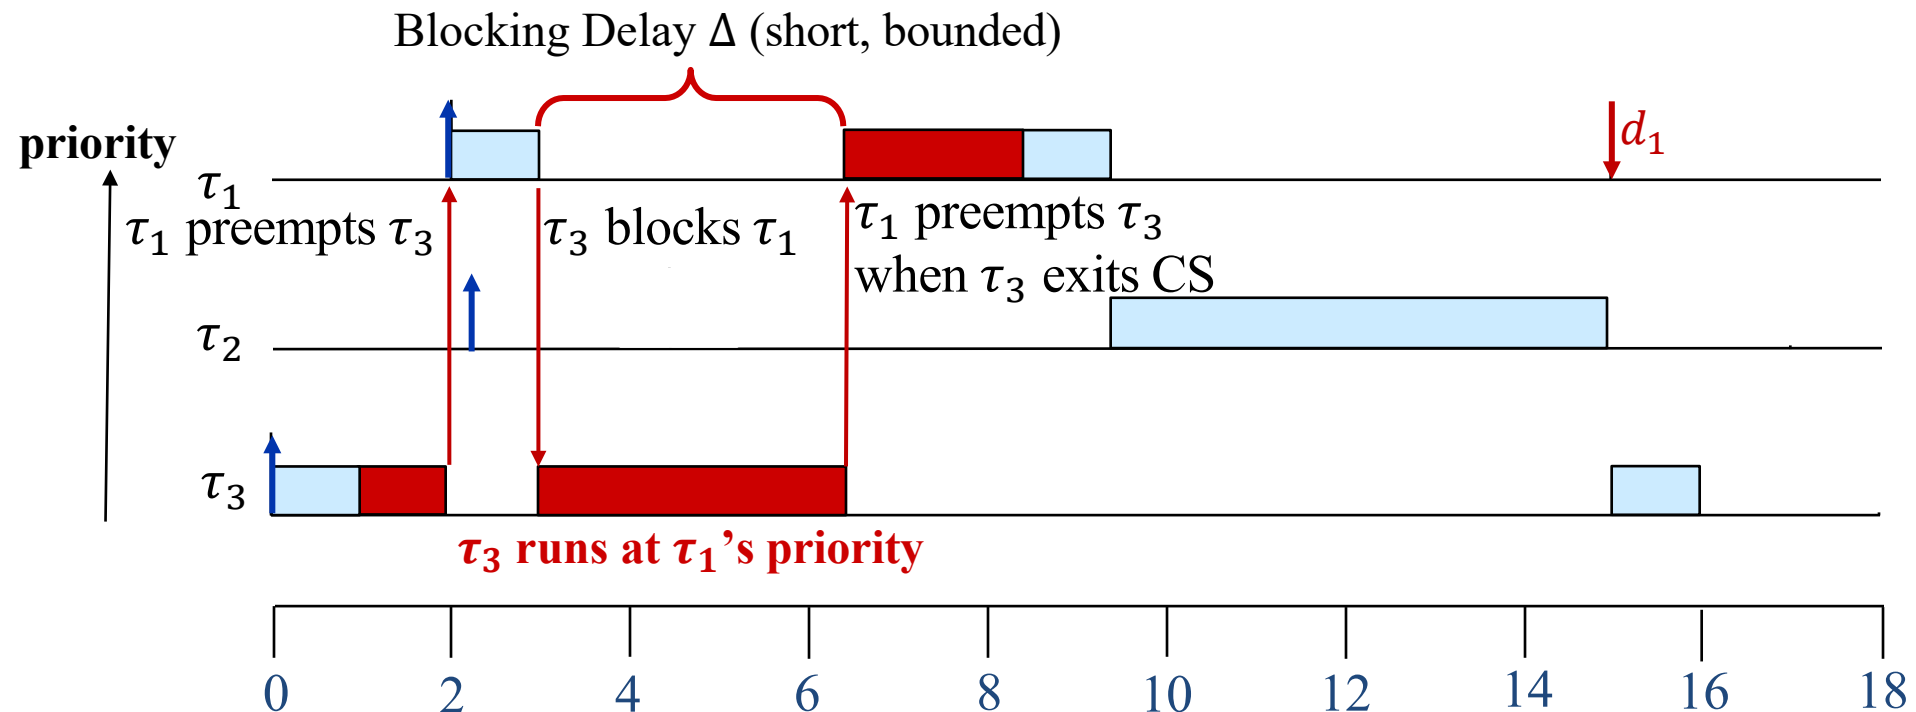


Priority Inversion and Priority Inheritance

- In 1997, this bug caused the Mars pathfinder to freeze up occasionally without explanation, and then starts working again
- Fixed by uploading a software patch enabling Priority-Inheritance Protocol (PIP)
- PIP: A task in a CS increases its priority if it is holding a lock s , and blocks other higher priority tasks, by inheriting the highest priority of all tasks blocked waiting for lock s
 - $P_{CS} = \max\{P_k | \tau_k \text{ blocked on CS}\}$


Priority Inheritance Protocol (PIP)

- In 1997, this bug caused the Mars pathfinder to freeze up occasionally and then starts working again. Fixed by uploading a software patch enabling Priority-Inheritance Protocol (PIP)
- A task τ_i in a CS increases its priority, if it is holding a lock s and blocks other higher priority tasks, by inheriting the highest priority of all tasks τ_k blocked waiting for lock s
 - $P_{\tau_i \text{ holding } s} = \max\{P_k | \tau_k \text{ blocked on } s\}$
- $t=3$: HP task τ_1 tries to enter CS, gets blocked since LP task τ_3 is in CS; τ_3 inherits τ_1 's high priority, and runs without pre-emption by MP task τ_2 (regardless of if τ_2 is released $t \in [2, 3]$ or $t > 3$)




Blocking Time under PIP

- Under PIP, task τ_i may experience two types of blocking delays:
 - **Direct blocking:** τ_i tries to lock s that is already locked
 - **Push-through blocking:** τ_i blocked by lower-priority task that has inherited a higher-priority (τ_i itself may not need any semaphores)
- Example:
 - HP task τ_1 experiences direct blocking by LP task τ_3 in time interval $[t_3, t_5]$
 - MP task τ_2 experiences push-through blocking by LP task τ_3 in time interval $[t_4, t_5]$



PIP Pros and Cons


- Pros:
 - It prevents priority inversion
 - It is transparent to the programmer
- Cons:
 - It does not prevent deadlocks and chained blocking



Deadlock still occurs under PIP

PCP Causes Chained Blocking

- **Chained blocking**: task τ_i can be blocked at most once by each lower priority task
- Theorem: Task τ_i can be blocked at most for the duration of $\min(n, m)$ critical sections
 - n is the number of tasks with priority lower than τ_i
 - m is the number of locks/semaphores on which τ_i can be blocked




Priority Ceiling Protocol (PCP)

- Assumptions: fixed-priority scheduling; resources required by all tasks are known a priori at design time (not required by PIP)
- PCP = PIP + ceiling blocking
- PIP still holds: When τ_i is blocked on s_k , the lower-priority task currently holding s_k inherits τ_i 's priority
- Each semaphore is assigned a ceiling $C(s_k) = \max\{P_j: \tau_j \text{ uses } s_k\}$
- Task τ_i can acquire s_k and enter CS only if $P_i > \max\{C(s_k): s_k \text{ locked by other tasks} \neq \tau_i\}$, that is,
 - Task τ_i 's priority P_i is higher than the ceilings of all semaphores currently held by other tasks; otherwise it is blocked on s_k
 - Corollary: If s_k is currently held by some task, then τ_i cannot lock s_k , since ceiling of s_k is at least the priority of τ_i , i.e., $P_i \leq C(s_k)$
- In addition to **direct blocking** and **push-through blocking** under PIP, a task may experience **ceiling blocking** under PCP: τ_i tried to lock s , but ceilings of currently locked semaphores are higher than P_i (s itself may be free)
 - Ceiling blocking helps to prevent deadlocks and chained blocking

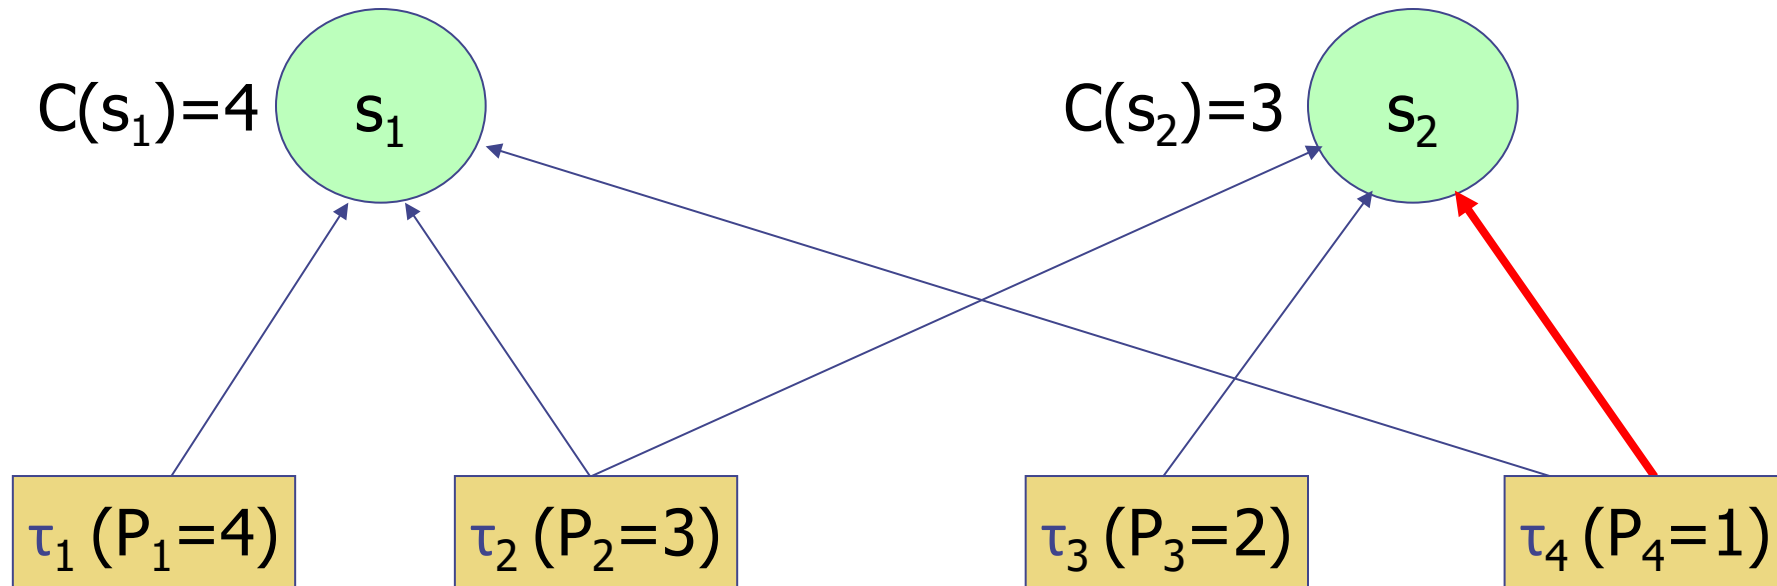
PCP Example I

- Three tasks τ_1, τ_2, τ_3 with priority ordering $P_1 > P_2 > P_3$. τ_1, τ_2 both require semaphore s , and τ_3 does not require any semaphore
 - $C(s) = \max\{P_j: \tau_j \text{ uses } s\} = \max\{P_1, P_2\} = P_1$
- The execution trace is the same as PIP:
 - When τ_3 is holding s , τ_1 cannot lock s , since $P_1 \leq C(s) = P_1$



PCP Example I


- Four tasks with priority ordering $P_1 = 4$ (highest) to $P_2 = 3$, $P_3 = 2$, $P_4 = 1$ (lowest) and two semaphores s_1, s_2
 - $C(s_1) = \max\{P_j: \tau_j \text{ uses } s_1\} = \max\{P_1, P_2, P_4\} = 4$
 - $C(s_2) = \max\{P_j: \tau_j \text{ uses } s_2\} = \max\{P_2, P_3, P_4\} = 3$
- When τ_4 is holding s_2 , τ_2 cannot lock s_1 , since $P_2 = 3 \leq C(s_2) = 3$ (ceiling blocking)
 - Ceiling blocking is “preventive blocking”, since τ_2 is blocked by τ_4 , even though τ_2 tries to lock s_1 which is free. This prevents any potential deadlocks in the future, when τ_2, τ_4 each holds one of s_1, s_2 and requests the other
- When τ_4 is holding s_2 , τ_1 can lock s_1 , since $P_1 = 4 > C(s_2) = 3$
 - Since τ_1 requires s_1 only, there is no cyclic dependency and no deadlock




PCP Prevents Deadlocks

- Classic deadlock scenario (with or without PIP): Two tasks τ_1 and τ_2 acquires two semaphores in opposite order:
 - LP task τ_2 runs first and locks s_2
 - HP task τ_1 starts running and locks s_1 , then tries to lock s_2 , gets blocked by τ_2
 - τ_2 starts running and tries to lock s_1 but τ_1 holds s_2 . Circular waiting \rightarrow deadlock
- Under PCP, $C(s_1) = C(s_2) = \max\{P_1, P_2\} = P_1$. Both semaphores s_1 and s_2 have ceiling equal to P_1 , since they are all required by the higher priority task τ_1 .
 - LP task τ_2 runs first and locks s_2
 - HP task τ_1 runs and preempts τ_2 . When τ_1 tries to lock s_1 , it is blocked since its priority does not exceed ceiling of s_2 , i.e., $P_1 \leq \text{ceil}(s_2) = P_1$
 - τ_2 will lock both s_2 and s_1 , and exit both CS before τ_1 can lock s_1 and s_2 . This prevents circular waiting and deadlock

Typical Deadlock




Deadlock avoidance with PCP



PCP Prevents Chained Blocking


- Example 1: $C(s_1) = \max\{P_1, P_3\} = P_1$, $C(s_2) = \max\{P_1, P_2\} = P_1$. Both semaphores s_1 and s_2 have ceiling equal to P_1 , since they are all required by the the highest priority task τ_1 .
- At time t_1 , LP task τ_3 is holding s_1 (in red CS). When MP task τ_2 tries to lock s_2 and enter yellow CS, it is blocked since its priority does not exceed ceiling of s_1 , $P_2 \leq C(s_1) = P_1$ (ceiling blocking)
- Hence τ_3 must unlock s_1 before τ_2 can lock s_2 . This prevents possible chained blocking



t_1 : τ_2 is blocked by the PCP, since $P_2 < C(s_1)$

PCP Prevents Chained Blocking

- Recall the example with chained blocking under PIP
- Under PCP: $C(s_1) = \max\{P_1, P_4\} = P_1$, $C(s_2) = \max\{P_1, P_3\} = P_1$, $C(s_3) = \max\{P_1, P_2\} = P_1$. All semaphores s_1, s_2, s_3 have ceiling equal to P_1 , since they are all required by the highest priority task τ_1 .
- While τ_4 is holding s_1 (in the red CS), τ_3 cannot lock s_2 , since $P_3 \leq C(s_1) = P_1$; and τ_2 cannot lock s_3 , since $P_2 \leq C(s_1) = P_1$ (ceiling blocking)
- Hence PCP prevents chained blocking, since task τ_1 is blocked at most once by a lower-priority task (either τ_4 , or τ_3 , or τ_2)



PCP Blocking Time

A given task i is blocked (or delayed) by at most one critical section of any lower priority task locking a semaphore with priority ceiling greater than or equal to the priority of task i . We can explain that mathematically using the notation:

$$B_i = \max_{\{k,s \mid k \in lp(i) \wedge s \in used_by(k) \wedge ceil(s) \geq pri(i)\}} CS_{k,s} \quad (4.5)$$

- Consider all lower-priority tasks ($k \in lp(i)$), and the semaphores they can lock (s)
- Select from those semaphores (s) with ceiling higher than or equal to $pri(i) = P_i$
- Take max length of all tasks (k)'s critical sections that lock semaphores (s)

PCP Pros and Cons

- Pros:
 - It prevents priority inversion, deadlocks, and chained blocking
 - Any given task is blocked at most once by a lower-priority task
- Cons:
 - It is not transparent to the programmer, as shared resources required by all tasks must be known a priori at design time, and designer needs to calculate priority ceilings of all semaphores and pass them to the OS (PIP does not need this step)

	Deadlock Prevention	Number of blockings	Programmer Transparency
PIP	No		Yes
PCP	Yes	1	No

Schedulability Analysis under PIP and PCP

- Let B_i denote the maximum blocking time experienced by task τ_i due to shared resources
- Schedulable utilization bound for RM scheduling with blocking time (sufficient condition):
 - A taskset is schedulable under RM scheduling with blocking time if
 - $\forall i$, priority level i utilization $U_i = \sum_{j \in hp(i)} \frac{C_j}{T_j} + \frac{C_i + B_i}{T_i} \leq i(2^{1/i} - 1)$
- Response Time Analysis (RTA) for RM scheduling with blocking time (necessary and sufficient condition):
 - WCRT R_i is computed by solving the following recursive equation:
 - $R_i = C_i + B_i + \sum_{j \in hp(i)} \left\lceil \frac{R_i}{T_j} \right\rceil C_j$

Example Taskset (without shared resources)

- System utilization $U = \frac{5}{50} + \frac{250}{500} + \frac{1000}{3000} = 0.933 > 0.780$
 - Since utilization exceeds the Utilization Bound of 0.780 of 3 tasks under RM scheduling, we cannot determine schedulability by the Utilization Bound test
- RTA shows that the taskset is schedulable by computing WCRT of each task:
 - $R_1 = C_1 + 0 = 5 + 0 = 5 \leq D_1 = 50$
 - $R_2 = C_2 + \left\lceil \frac{R_2}{T_1} \right\rceil \cdot C_1 = 250 + \left\lceil \frac{R_2}{50} \right\rceil \cdot 5 = 280 \leq D_2 = 500$
 - $R_3 = C_3 + \left\lceil \frac{R_3}{T_1} \right\rceil \cdot C_1 + \left\lceil \frac{R_3}{T_2} \right\rceil \cdot C_2 = 1000 + \left\lceil \frac{R_3}{50} \right\rceil \cdot 5 + \left\lceil \frac{R_3}{500} \right\rceil \cdot 250 = 2500 \leq D_3 = 3000$

Task	T	D	C	Prio
1	50	50	5	H
2	500	500	250	M
3	3000	3000	1000	L

Example Taskset (with shared resources under PCP) I

- 3 semaphores s_1, s_2, s_3
 - Task 1 requires semaphore s_1 , with critical section length 1
 - Task 2 requires semaphores s_2 and s_3 , with critical section lengths 2 and 5, respectively
 - Task 3 requires semaphores s_2 and s_3 , with critical section lengths 3 and 4, respectively
- Ceilings $C(s_1) = P_1 = H$; $C(s_2) = C(s_3) = \max(P_2, P_3) = M$
- Blocking times:
 - Task 1: $B_1 = 0$ (Task 1's priority is higher than ceilings of s_2 and s_3), so it remains schedulable
 - Task 2: $B_2 = \max(3, 4) = 4$ (maximum CS length of LP Task 3)
 - » Utilization $U_2 = \sum_{j \in hp(2)} \frac{C_j}{T_j} + \frac{C_2 + B_2}{T_2} = \frac{5}{50} + \frac{250+4}{500} = 0.608 \leq 0.828$ (utilization bound for 2 tasks under RM)
 - » Or WCRT: $R_2 = C_2 + B_2 + \left\lceil \frac{R_2}{T_1} \right\rceil \cdot C_1 = 250 + 4 + \left\lceil \frac{R_2}{50} \right\rceil \cdot 5 = 284 \leq D_2 = 500$
 - Task 3: $B_3 = 0$ (Task 3 is the lowest priority task, so it does not experience blocking), so it remains schedulable
- The taskset remains schedulable with shared resources under PCP

Task	T	D	C	Prio	sems	CS Len
1	50	50	5	H	s_1	1
2	500	500	250	M	s_2, s_3	2, 5
3	3000	3000	1000	L	s_2, s_3	3, 4



Example Taskset (with shared resources under PCP) II

- 3 semaphores s_1, s_2, s_3
 - Task 1 requires semaphores s_1, s_2 and s_3 with critical section lengths 1, 1, 1
 - Task 2 requires semaphores s_2 and s_3 , with critical section lengths 2 and 5, respectively
 - Task 3 requires semaphores s_2 and s_3 , with critical section lengths 3 and 4, respectively
- Ceilings $C(s_1) = C(s_2) = C(s_3) = \max(P_1, P_2, P_3) = H$
- Blocking times:
 - Task 1: $B_1 = \max(2, 5, 3, 4) = 5$ (maximum CS length of LP Tasks 2 and 3, associated with semaphores s_2 and s_3)
 - » Utilization $U_1 = \sum_{j \in hp(1)} \frac{C_j}{T_j} + \frac{C_1 + B_1}{T_1} = 0 + \frac{5+5}{50} = 0.2 \leq 1$ (Utilization bound for 1 task under RM), so Task 1 remains schedulable
 - » Or WCRT: $R_1 = C_1 + B_1 = 5 + 5 = 10 \leq D_1 = 50$
 - » (Task 1's CS lengths 1, 1, 1 do not matter since it is the highest priority task and does not block any other task)
 - Task 2: $B_2 = \max(3, 4) = 4$ (maximum CS length of LP Task 3)
 - » Same calculation of utilization and WCRT as before
 - Task 3: $B_3 = 0$ (Task 3 is the lowest priority task, so it does not experience blocking), so it remains schedulable
- The taskset remains schedulable with shared resources under PCP

Task	T	D	C	Prio	sems	CS Len
1	50	50	5	H	s_1, s_2, s_3	1, 1, 1
2	500	500	250	M	s_2, s_3	2, 5
3	3000	3000	1000	L	s_2, s_3	3, 4

Scheduling Anomaly w/ Resource Synchronization

- Doubling processor speed causes T1 to miss its deadline
 - (Yellow part denotes a critical section shared by T1 and T2)



Online Resources

- Priority-Driven Scheduling, Marilyn Wolf
 - https://www.youtube.com/watch?v=zSgr_oFmjql&list=PLzwefUCNStZsmz5fWVPVwVvTo1iPeGmG9M&index=4
- RMS and EDF, Marilyn Wolf
 - <https://www.youtube.com/watch?v=oHMC2aO8GII&list=PLzwefUCNStZsmz5fWVPVwVvTo1iPeGmG9M&index=5>
- Real-Time Scheduling Models, Marilyn Wolf (long)
 - <https://www.youtube.com/watch?v=WloSQ7ZEKXk>