

# CSC 112: Computer Operating Systems

## Lecture 6

### Real-Time Scheduling II

Department of Computer Science,  
Hofstra University

# Outline

---

- Part I
  - Introduction to RTOS and Real-Time Scheduling
  - Fixed-Priority Scheduling
  - Earliest Deadline First Scheduling
  - Least Laxity First (LLF) Scheduling
  - Preemptive vs. Non-Preemptive Scheduling
- Part II
  - Multiprocessor Scheduling
  - Resource Synchronization Protocols (for Fixed-Priority Scheduling)

---

# Multiprocessor Scheduling

# Multiprocessor models

---

- Identical multiprocessors:
  - each processor has the same computing capacity
- Uniform multiprocessors:
  - different processors have different computing capacities
- Heterogeneous multiprocessors:
  - each (task, processor) pair may have a different computing capacity
- MP scheduling
  - Many NP-hard problems, with few optimal results, mainly heuristic approaches
  - Only sufficient schedulability tests


# Multiprocessor Models

---

Identical multiprocessors: each processor has the same speed

Task T1

Task T2






# Multiprocessor Models

Uniform multiprocessors: different processors have different speeds


Task T1

Task T2



# Multiprocessor Models

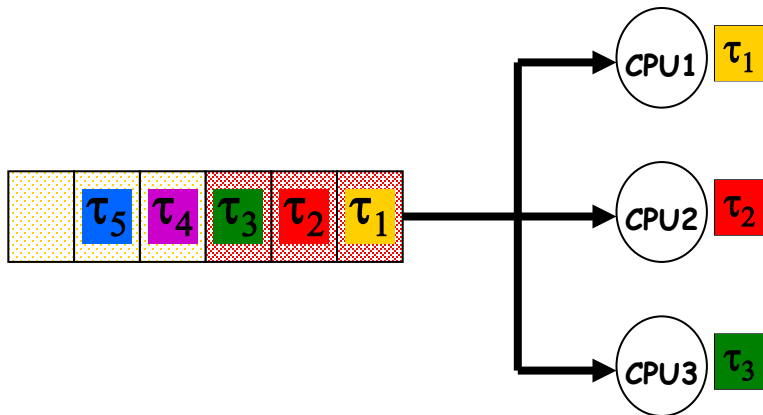
**Heterogeneous** multiprocessors: each (task, processor) pair may have a different relative speed, due to specialized processor architectures



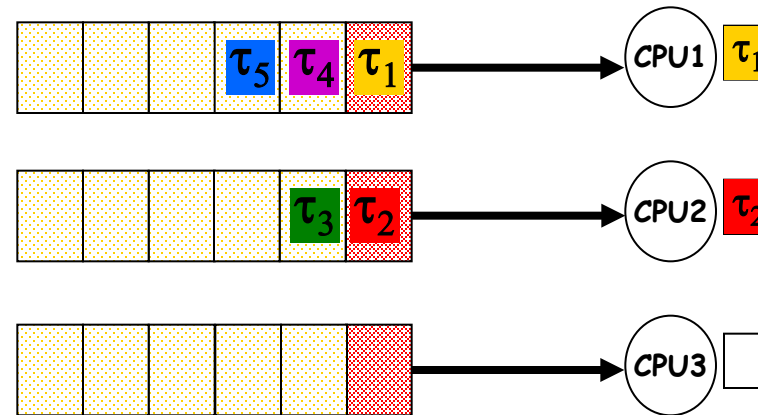
# Global vs partitioned scheduling

- Global scheduling
  - All ready jobs are kept in a common (global) queue; when selected for execution, a job can be dispatched to an arbitrary processor, even after being preempted
- Partitioned scheduling
  - Each task may only execute on a specific processor

**Global scheduling:  
Single system-wide queue**



**Partitioned scheduling:  
per-processor queues**





# Global Scheduling vs. Partitioned Scheduling

---

- Global Scheduling

- Pros:

- Runtime load-balancing across cores
  - » More effective utilization of processors and overload management
- Supported by most multiprocessor operating systems
  - » Windows, Linux, MacOS...

- Cons:

- Low schedulable utilization
- Weak theoretical framework

- Partitioned Scheduling

- Pros:

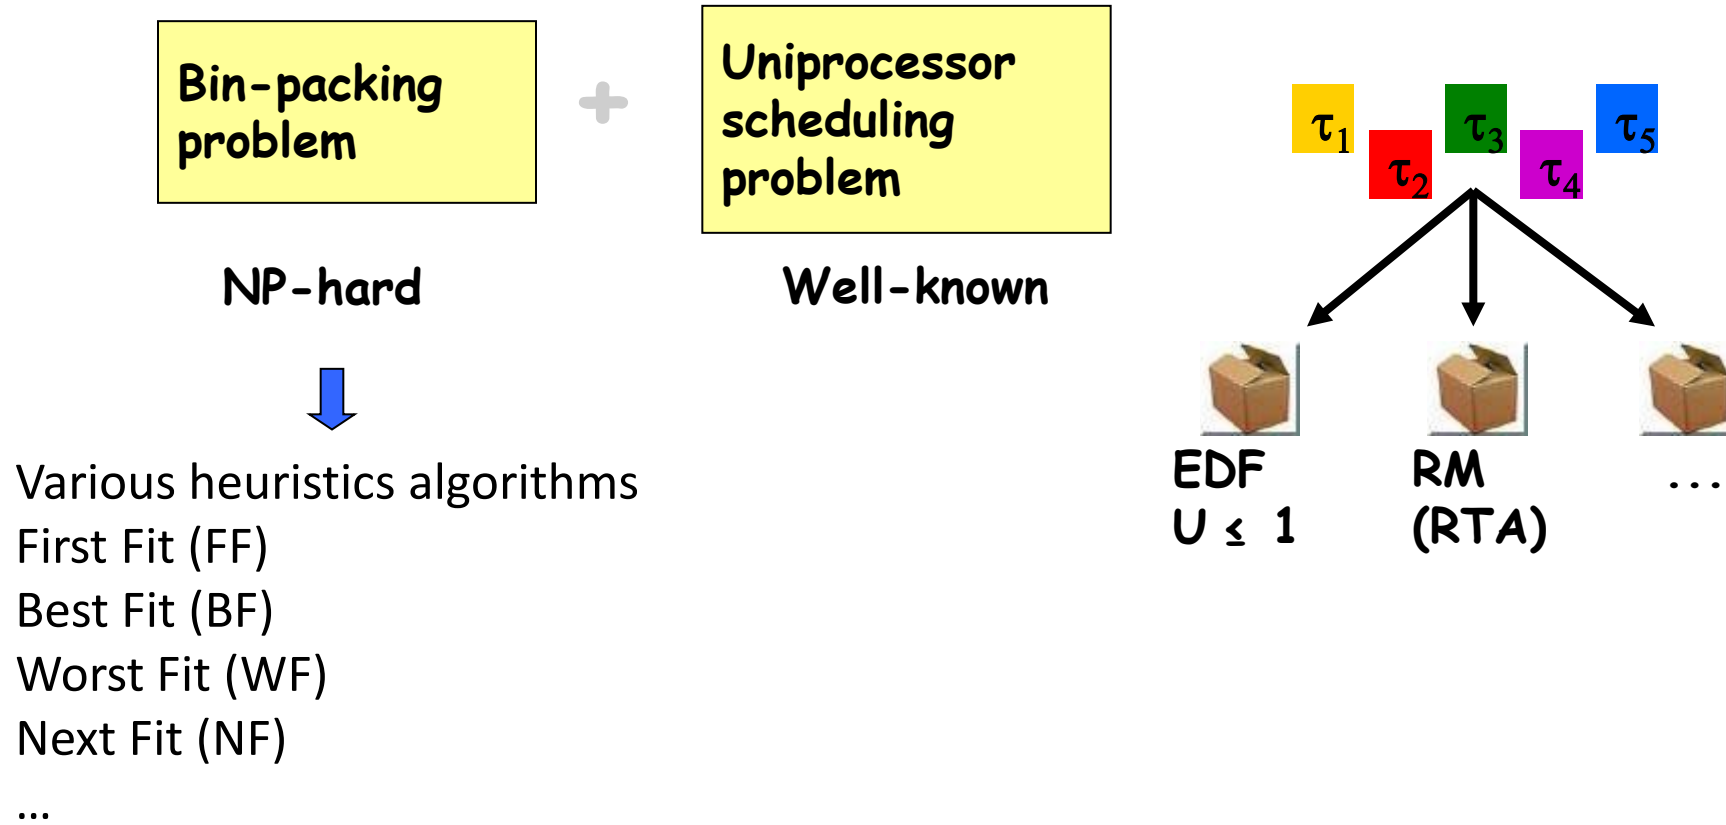
- Mature scheduling framework
- Uniprocessor scheduling theory scheduling are applicable on each core; uniprocessor resource access protocols (PIP, PCP...) can be used
- Partitioning of tasks can be done by efficient bin-packing algorithms

- Cons:

- No runtime load-balancing; surplus CPU time cannot be shared among processors

# Partitioned Scheduling

- Scheduling problem reduces to:



# Partitioned Scheduling

---

- Bin-packing algorithms:
  - The problem concerns packing objects of varying sizes in boxes ("bins") with some optimization objective, e.g., minimizing number of used boxes (best-fit), or minimizing the maximum workload for each box (worst-fit)
- Application to multiprocessor scheduling:
  - Bins are represented by processors and objects by tasks
  - The decision whether a processor is "full" or not is derived from a utilization-based feasibility test.
- Since optimal bin-packing is a NP-complete problem, partitioned scheduling is also NP-complete
- Example: Rate-Monotonic-First-Fit (RMFF): (Dhall and Liu, 1978)
  - Let the processors be indexed as 1, 2, ...
  - Assign the tasks to processor in the order of increasing periods (that is, RM order)
  - For each task  $\tau_i$ , choose the lowest previously-used processor  $j$  such that  $\tau_i$ , together with all tasks that have already been assigned to processor  $j$ , can be feasibly scheduled according to the utilization-based schedulability test
  - Additional processors are added if needed

# Assumptions for Global Scheduling

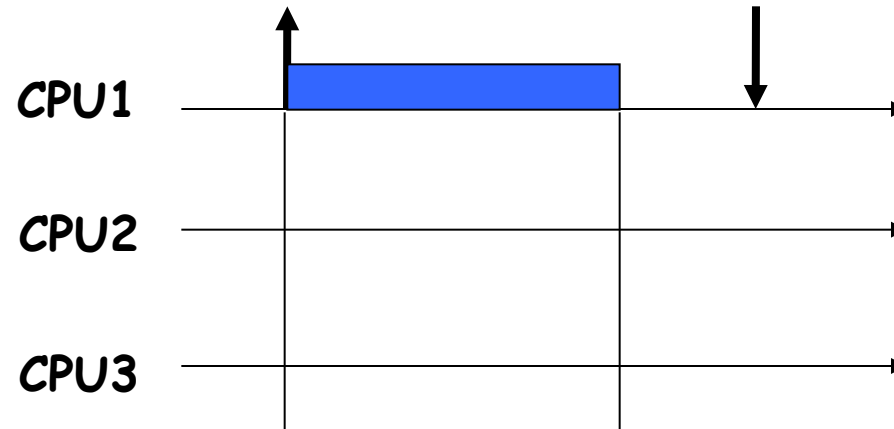
---

- Identical multiprocessors
- Work-conserving:
  - At each instant, the highest-priority jobs that are eligible to execute are selected for execution upon the available processors
  - No processor is ever idle when the ready queue is non-empty
- Preemption and Migration support
  - A preempted task can resume execution on a different processor with 0 overhead, as cost of preemption/migration is integrated into task WCET
- No job-level parallelism
  - the same job cannot be *simultaneously* executed on more than one processor, i.e., we do not consider parallel programs that can run on multiple processors in parallel

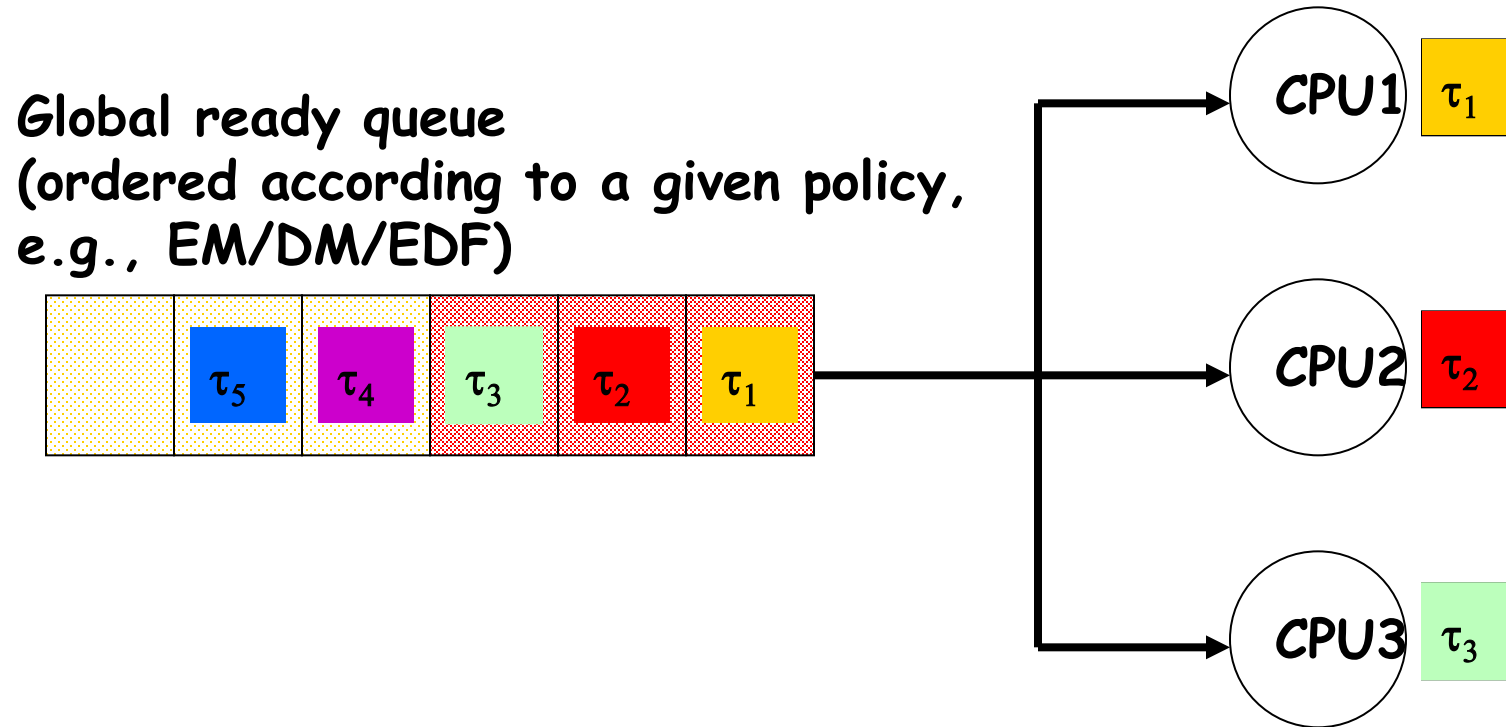
# Source of Difficulty

---

- The “no job-level parallelism” assumption leads to difficult scheduling problems
- “The simple fact that a task can use only one processor even when several processors are free at the same time adds a surprising amount of difficulty to the scheduling of multiple processors” [Liu’69]

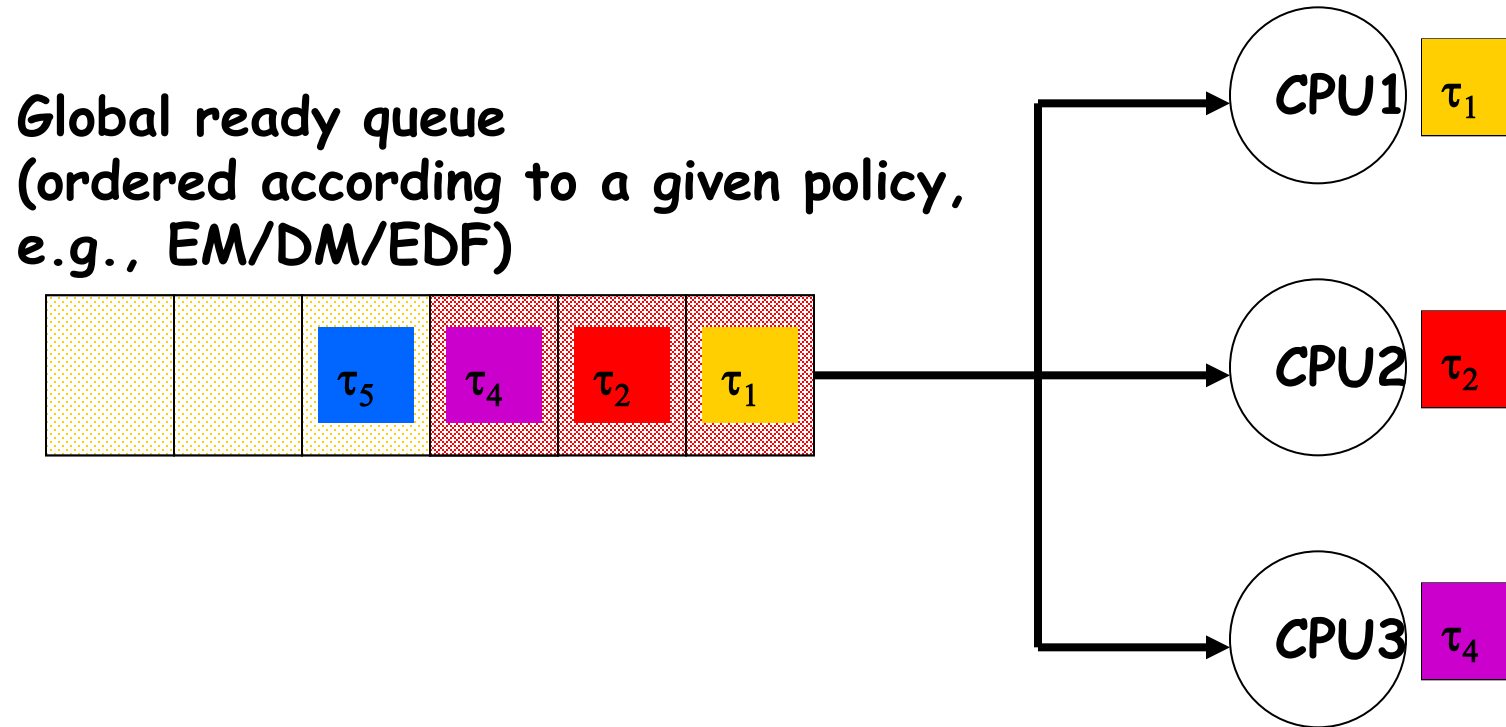


# Global scheduling example



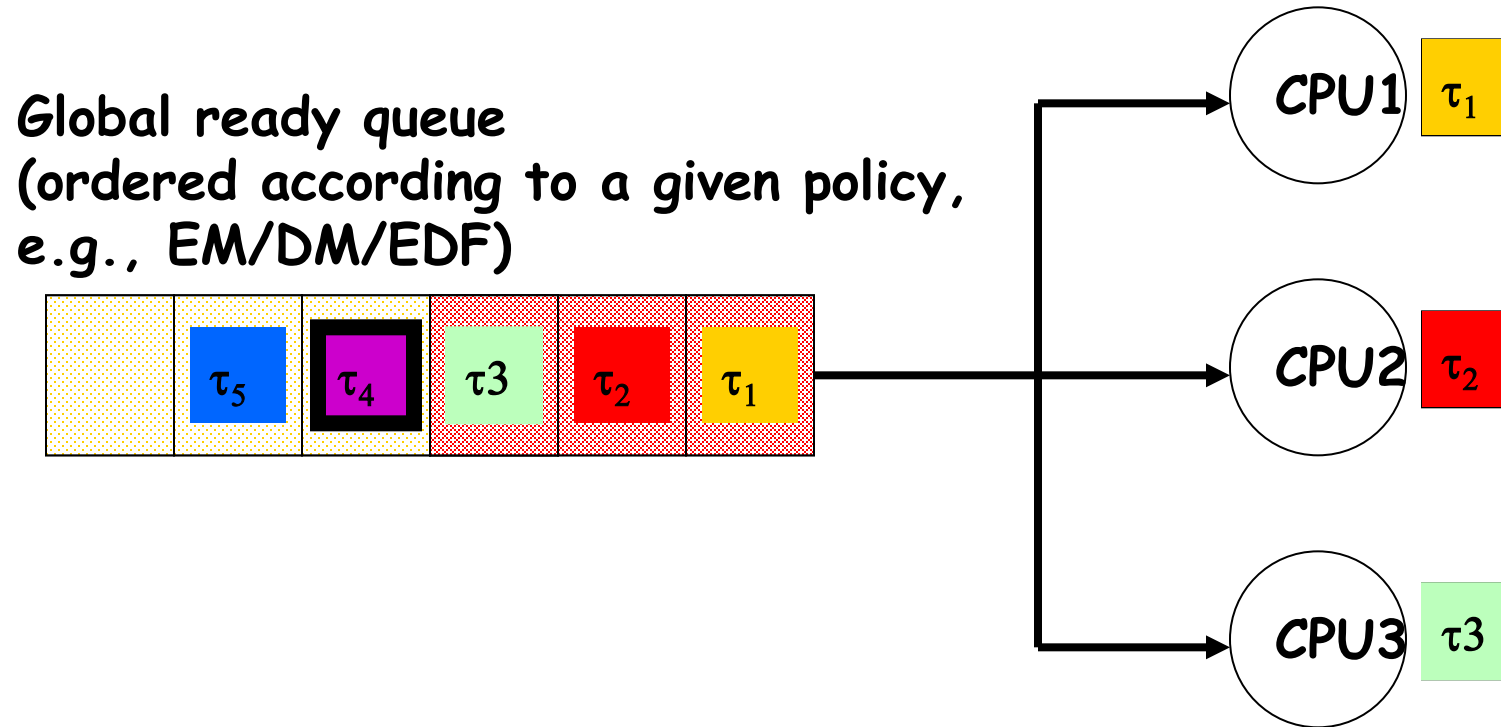
The first  $m$  jobs in the queue are scheduled upon the  $m$  CPUs

# Global scheduling example



When a job  $\tau_3$  finishes its execution, the next job in the queue  $\tau_4$  is scheduled on the available CPU

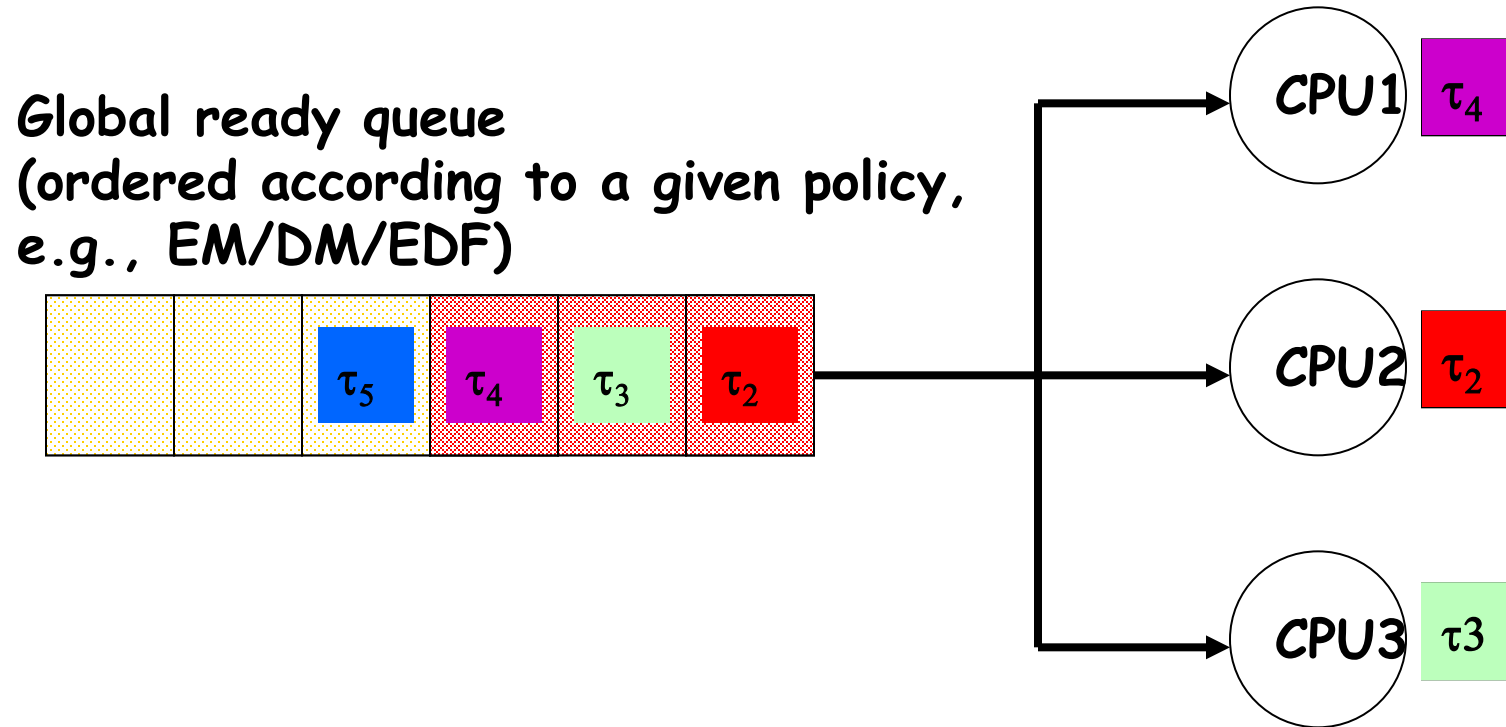
# Global scheduling example



When a new higher-priority job  $\tau_3$  arrives in its next period  $T_3$ , it preempts the job with lowestpriority  $\tau_4$  among the executing ones



# Global scheduling example



When another job  $\tau_1$  finishes its execution, the preempted job  $\tau_4$  can resume its execution. Net effect:  $\tau_4$  "migrated" from CPU3 to CPU1

## Global vs. Partitioned


---

- Global (work-conserving) and partitioned scheduling algorithms are incomparable:
  - There are tasksets that are schedulable with a global scheduler, but not with a partitioned scheduler, and vice versa.

# Global vs Partitioned (FP) Scheduling

Task	T=D	C	Prio
T1	2	1	H
T2	3	2	M
T3	3	2	L

- A taskset schedulable with global scheduling, but not partitioned scheduling. System utilization  $U = \frac{1}{2} + \frac{2}{3} + \frac{2}{3} = 1.83$
- Global FP scheduling is schedulable with priority assignment  $p_1 > p_2 > p_3$  (or  $p_2 > p_1 > p_3$ )
- Partitioned scheduling is unschedulable, since assigning any two tasks to the same processor will cause that processor's utilization to exceed 1, so the bin-packing problem has no feasible solution



A feasible execution trace under global scheduling

# Global vs Partitioned (FP) Scheduling

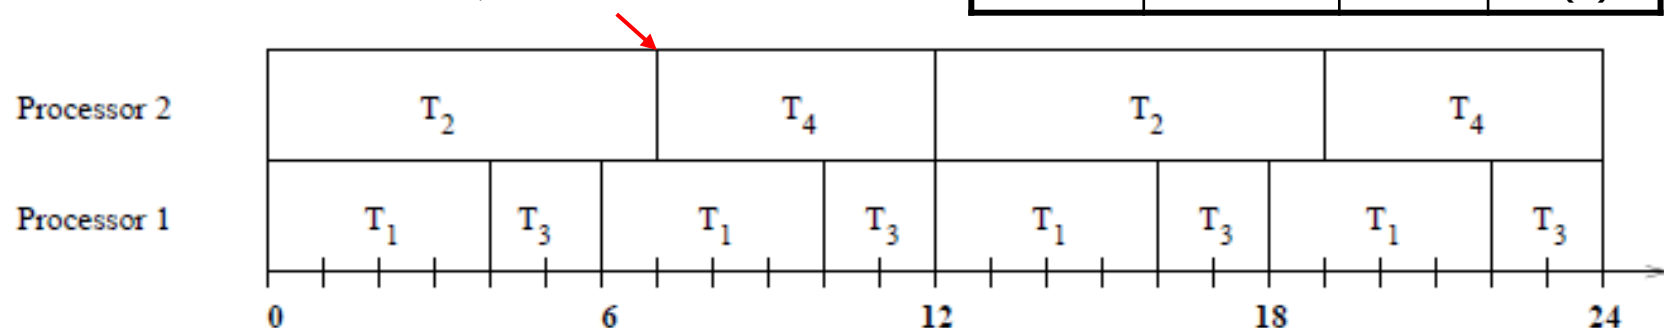
Task	T=D	C	Prio
T1	6	4	4(H)
T2	12	7	3
T3	12	4	2
T4	24	10	1(L)

- A taskset schedulable with partitioned scheduling, but not global scheduling. System utilization  $U = \frac{4}{6} + \frac{7}{12} + \frac{4}{12} + \frac{10}{24} = 2.0$ , hence the two processors must be fully utilized with no possible idle intervals

- Partitioned FP scheduling with RM priority assignment ( $p_1 > p_2 > p_3 > p_4$ ) is schedulable. T1, T3 assigned to Processor 1; T2, T4 assigned to Processor 2. Both processors have utilization 1.0, and harmonic task periods


- Global FP scheduling with RM priority assignment  $p_1 > p_2 > p_3 > p_4$  is unschedulable. Compared to partitioned scheduling, the difference is at time 7, when T3 (with higher priority than T4) runs on Processor 2. This causes idle intervals on Processor 1 [10,12] and [22,24], since only one task T4 is ready during these time intervals. Since taskset  $U = 2.0$  on 2 processors, any idle interval will cause the taskset to be unschedulable

At time 7, T4 runs on Processor 2



A feasible execution trace under partitioned scheduling

At time 7, T3 runs on Processor 2



An infeasible execution trace under global scheduling


# Difficulties of Global Scheduling

---

- Dhall's effect
  - With RM, DM and EDF, some low-utilization task sets can be unschedulable regardless of how many processors are used.
- Scheduling anomalies
  - Decreasing task execution time or increasing task period may cause deadline misses
- Hard-to-find worst-case
  - The worst-case does not always occur when a task arrives at the same time as all its higher-priority tasks
- Dependence on relative priority ordering (omitted)
  - Changing the relative priority ordering among higher-priority tasks may affect schedulability for a lower-priority task

# Dhall's effect

- Global RM/DM/EDF can fail at very low utilization
- Example:  $m$  processors,  $n=m+1$  tasks. Tasks  $\tau_1, \dots, \tau_m$  are light tasks, with small  $C_i = 1$ ,  $T_i = D_i = T - 1$ ; Task  $\tau_{m+1}$  is a heavy task, with large  $C_i = T$ ,  $T_i = D_i = T$ .  $T > 1$  is some constant value
- For global RM/DM/EDF, Task  $\tau_{m+1}$  has lowest priority, so  $\tau_1, \dots, \tau_m$  must run on  $m$  processors starting at time 0, causing  $\tau_{m+1}$  to miss its deadline
- One solution: assign higher priority to heavy tasks
  - If heavy task  $\tau_{m+1}$  is assigned the highest priority, then it runs from time 0 to  $T$  and meets its deadline; The light tasks can run on other processors and meet their deadlines as well



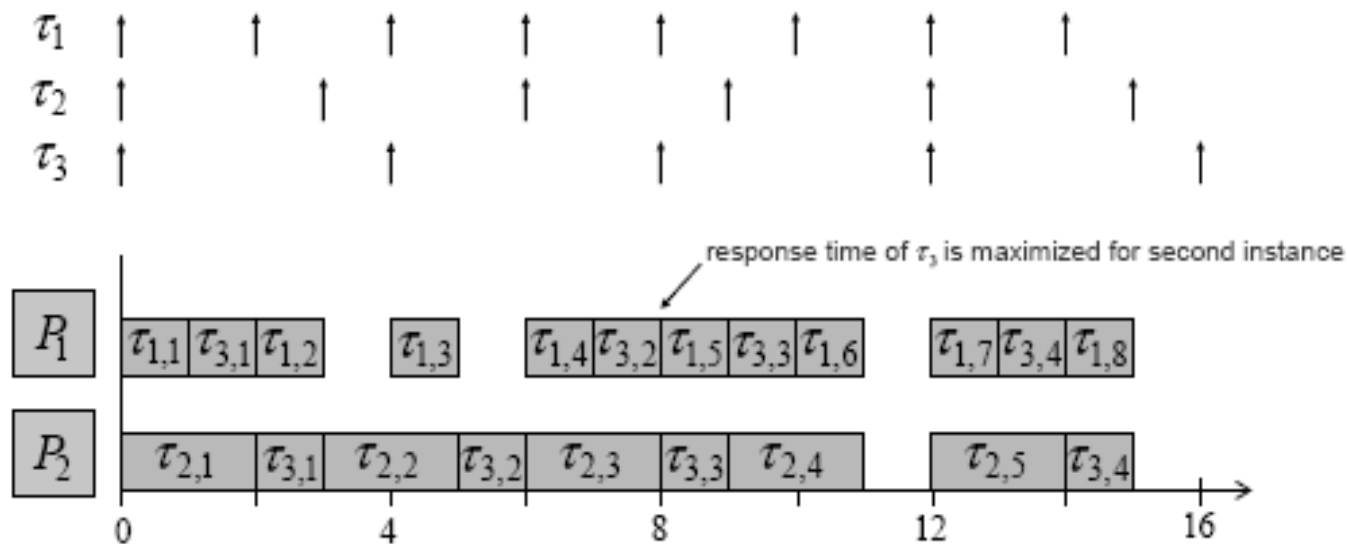
# Hard-to-Find Worst-Case

- For uniprocessor scheduling, the worst case occurs when all tasks are initially released at time 0 simultaneously, called the critical instant (recall Slide [Response Time Analysis \(RTA\)](#))
- This is no longer true for multiprocessor scheduling, as the worst-case interference for a task does not always occur at time 0, when all tasks are initially released at time 0 simultaneously
  - Response time for task  $\tau_3$  is maximized for its 2<sup>nd</sup> job  $\tau_{3,2}$  (8-4=4), which does not arrive at the same time as its higher priority tasks; not for its 1<sup>st</sup> job  $\tau_{3,1}$  (3-0=3), which arrives at the same time as its higher priority tasks

Hard-to-find critical instant:

(RM scheduling)

$$\begin{aligned}\tau_1 &= \{C_1 = 1, T_1 = 2\} \\ \tau_2 &= \{C_2 = 2, T_2 = 3\} \\ \tau_3 &= \{C_3 = 2, T_3 = 4\}\end{aligned}$$



# MP Scheduling Anomalies

---

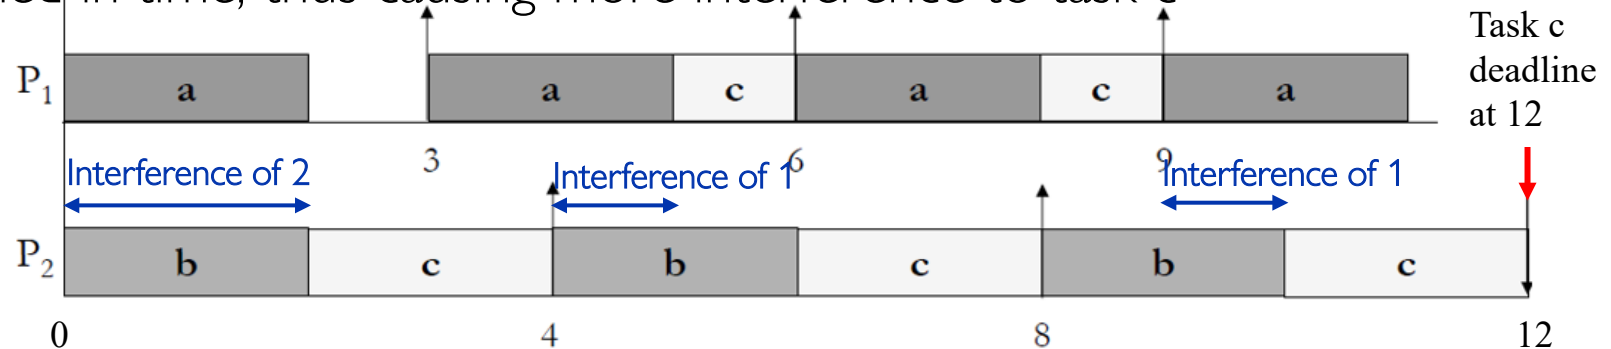
- Decrease in processor demand (decreasing task execution time or increasing task period) may cause deadline misses!
- **Anomaly 1**
  - Decrease in processor demand from higher-priority tasks can *increase the interference on a lower-priority task because of change in the time when the tasks execute*
- **Anomaly 2**
  - Decrease in processor demand of a task *negatively affects the task itself because change in the task arrival times cause it to suffer more interference*




# Scheduling Anomaly Example 1

- Three tasks on two processors under global scheduling
- With Task a's period  $T_a = 3$ , system utilization  $\sum U_i = 1.83$ . WCRT of task c is  $R_c = 12 \leq D_c = 12$ .  $R_c = C_c + I_c = 8 + I_c$ , where  $I_c = 2 + 1 + 1 = 4$  is interference by higher priority tasks a and b. (Task c experiences inference when both processors are busy executing higher priority tasks a and b.) Task c is schedulable but saturated, as any increase in its WCET or interference would make it unschedulable.
- With Task a's period  $T_a = 4$ , system utilization  $\sum U_i = 1.67$  is reduced. But WCRT of task c increases:  $R_c = 14 > D_c = 12$ .  $R_c = 8 + I_c$  where  $I_c = 2 + 2 + 2 = 6$ , since execution segments of tasks a and b on two processors are aligned in time, thus causing more interference to task c

Task	T=D	C	Util	Prio
a	3	2	0.67	H
b	4	2	0.5	M
c	12	8	0.67	L



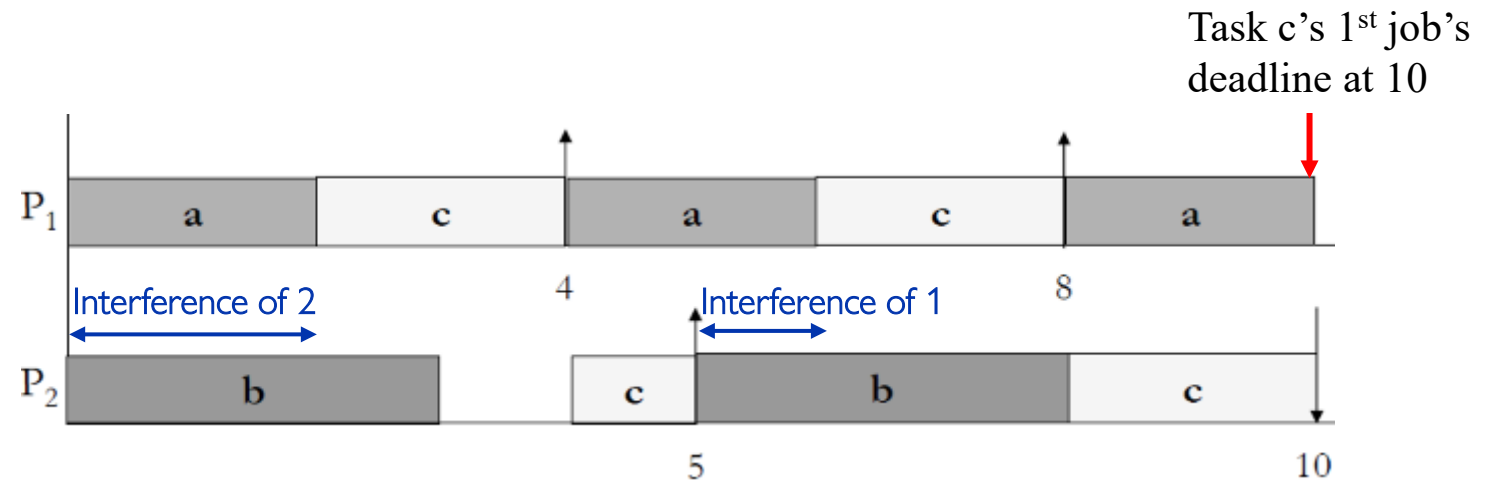
Task	T=D	C	Util	Prio
a	4	2	0.5	H
b	4	2	0.5	M
c	12	8	0.67	L



## Scheduling Anomaly Example 2

- Three tasks on two processors under global scheduling
- With Task c's period  $T_c = 10$ , system utilization  $\sum U_i = 1.8$ . WCRT of task c is  $R_c = 10 \leq D_c = 10$ .  $R_c = C_c + I_c = 7 + 3 = 10$ , where  $I_c = 2 + 1 = 3$  is interference by higher priority tasks a and b. Its 1st job meets its deadline at time 10. This schedule repeats in future periods, hence task c is schedulable but saturated, as any increase in its WCET or interference would make it unschedulable.


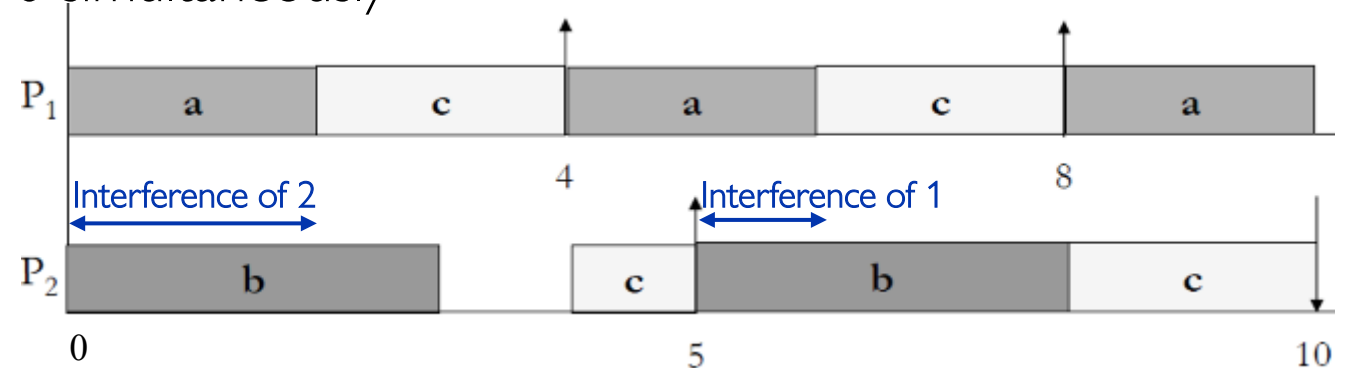
Task	T=D	C	Util	Prio
a	4	2	0.5	H
b	5	3	0.6	M
c	10	7	0.7	L



## Scheduling Anomaly Example 2

- With Task c's period  $T_c = 11$ , system utilization  $\sum U_i = 1.74$  is reduced. WCRT of task c is  $R_c = 12 > D_c = 10$ . Its 1<sup>st</sup> job has response time  $C_c + I_c = 7 + 3 = 10 \leq D_c = 11$ , where  $I_c = 2 + 1 = 3$ , but this is not task c's WCRT.
- Its 2<sup>nd</sup> job has response time  $C_c + I_c = 7 + 5 = 12 > D_c = 11$ , where  $I_c = 1 + 2 + 2 = 5$ . The 2<sup>nd</sup> job finishes at time  $11+12=23$ , and misses its deadline at time 22.
- Another example where the worst-case interference for task c does NOT occur at time 0, when all tasks are initially released at time 0 simultaneously

Task	T=D	C	Util	Prio
a	4	2	0.5	H
b	5	3	0.6	M
c	11	7	0.64	L

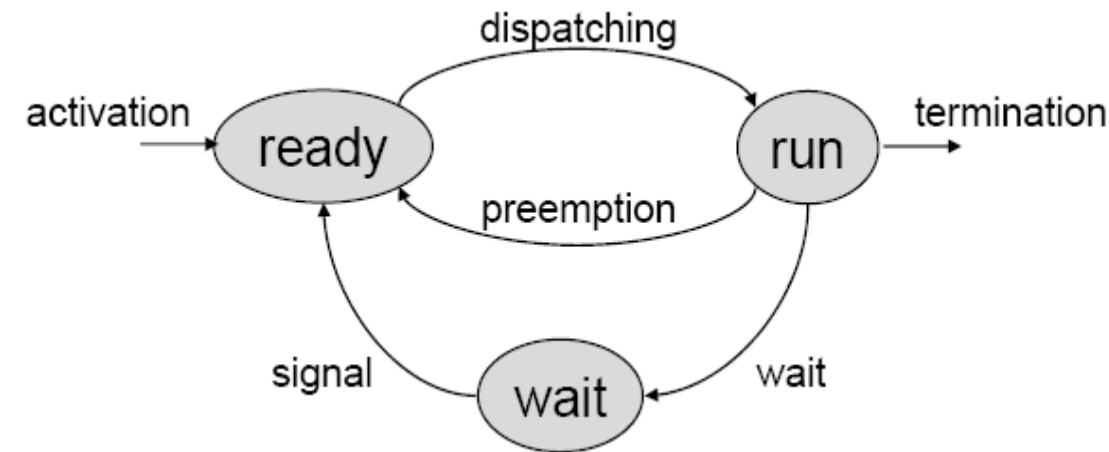
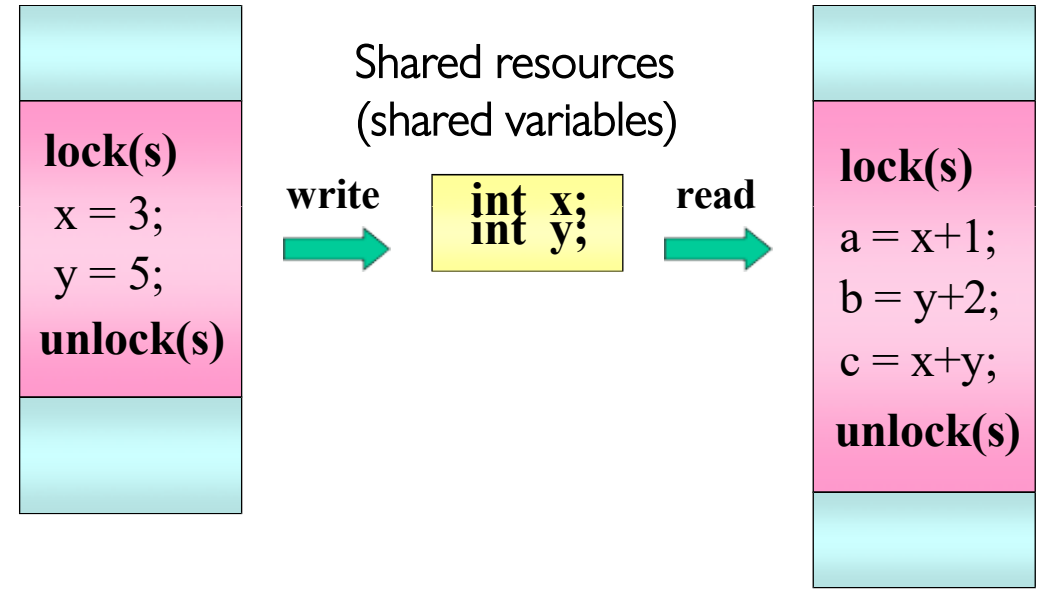


---

# Resource Synchronization Protocols (for Fixed-Priority Scheduling)

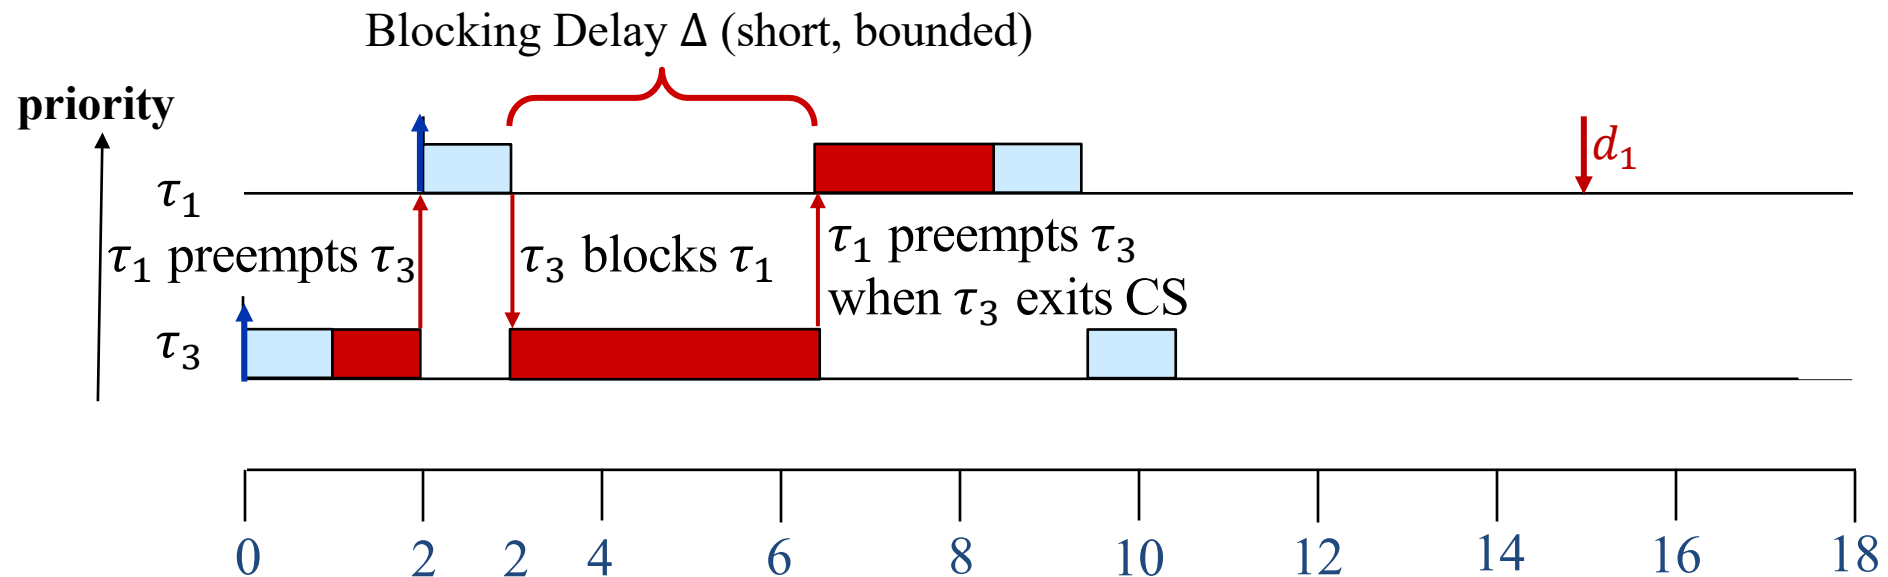
# Resource Sharing

- When two tasks access shared resources (variables), mutexes (or binary semaphores) are used to protect critical sections. Each Critical Section (CS) must begin with `lock(s)` and end with `unlock(s)`
- A task waiting for a shared resource is blocked on that resource. Otherwise, it proceeds by entering the critical section and holds the resource
- Tasks blocked on the same resource are kept in a queue. When a running task invokes `lock(s)` when `s` is already locked, it enters the waiting state, until another task unlocks `s`



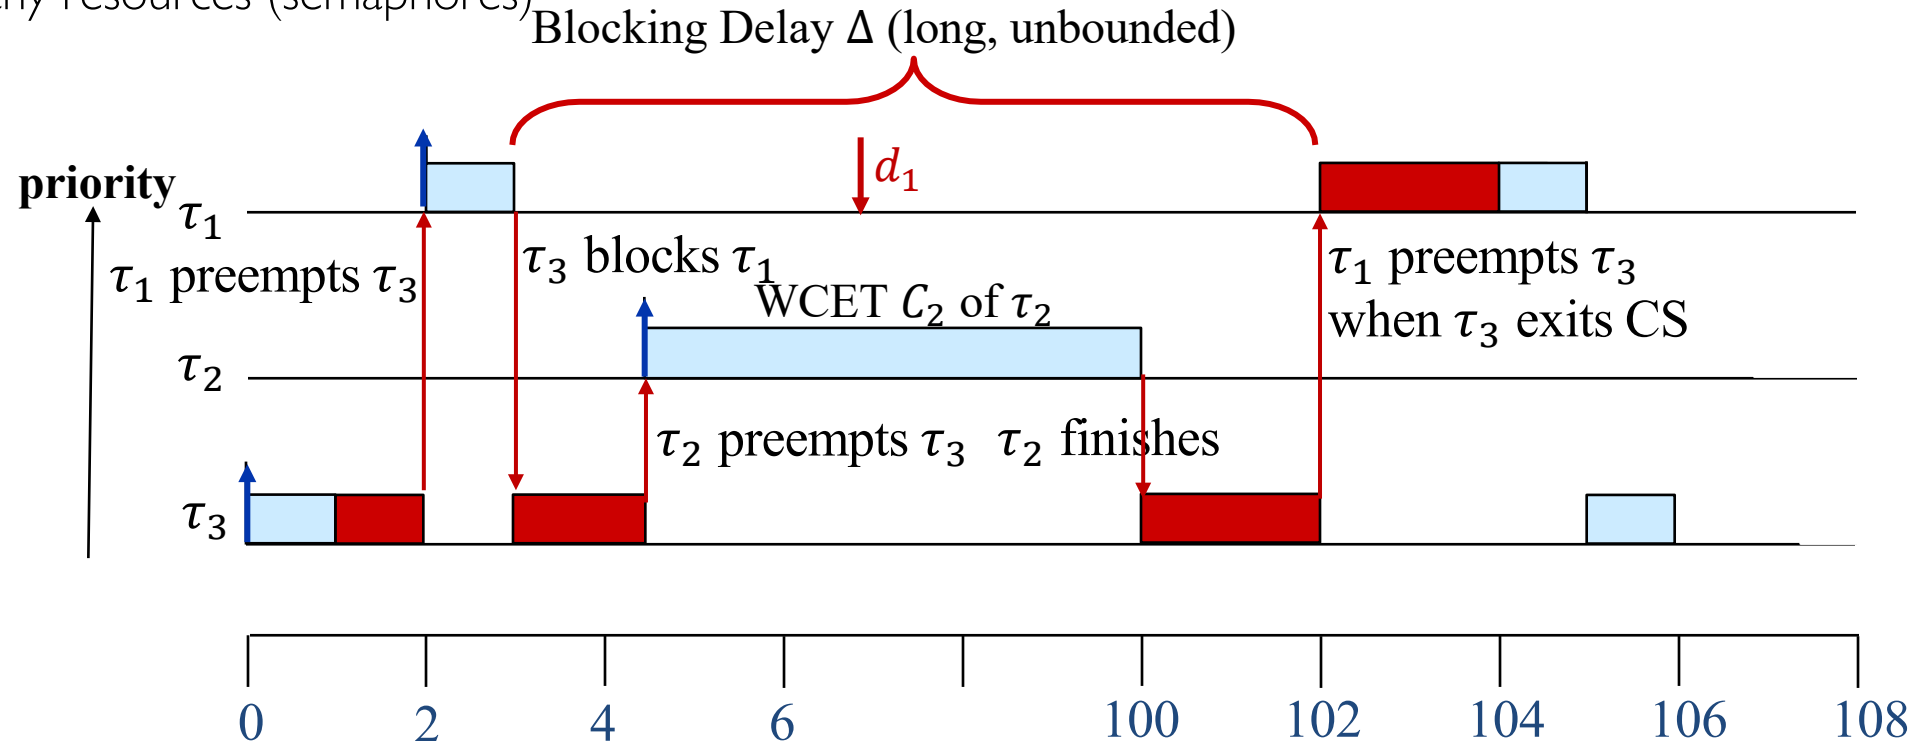
# Blocking Delay

- Lower Priority (LP) tasks can cause **blocking delay** to Higher Priority (HP) tasks due to resource sharing
  - HP tasks may cause **preemption delay** to LP tasks, but not blocking delay
- Example: Two tasks  $\tau_1, \tau_3$  with priority ordering  $P_1 > P_3$ . They both require semaphore  $s$  (which protects the red CS)
- If HP task  $\tau_1$  tries to lock  $s$  that is held by LP task  $\tau_3$ ,  $\tau_1$  is blocked until  $\tau_3$  unlocks  $s$ , so  $\tau_1$  experiences a blocking delay  $\Delta$ .
  - Since CS is typically very short, it seems this blocking time delay  $\Delta$  is bounded by the longest critical section in lower-priority tasks?
- No, blocking delay may be unbounded!



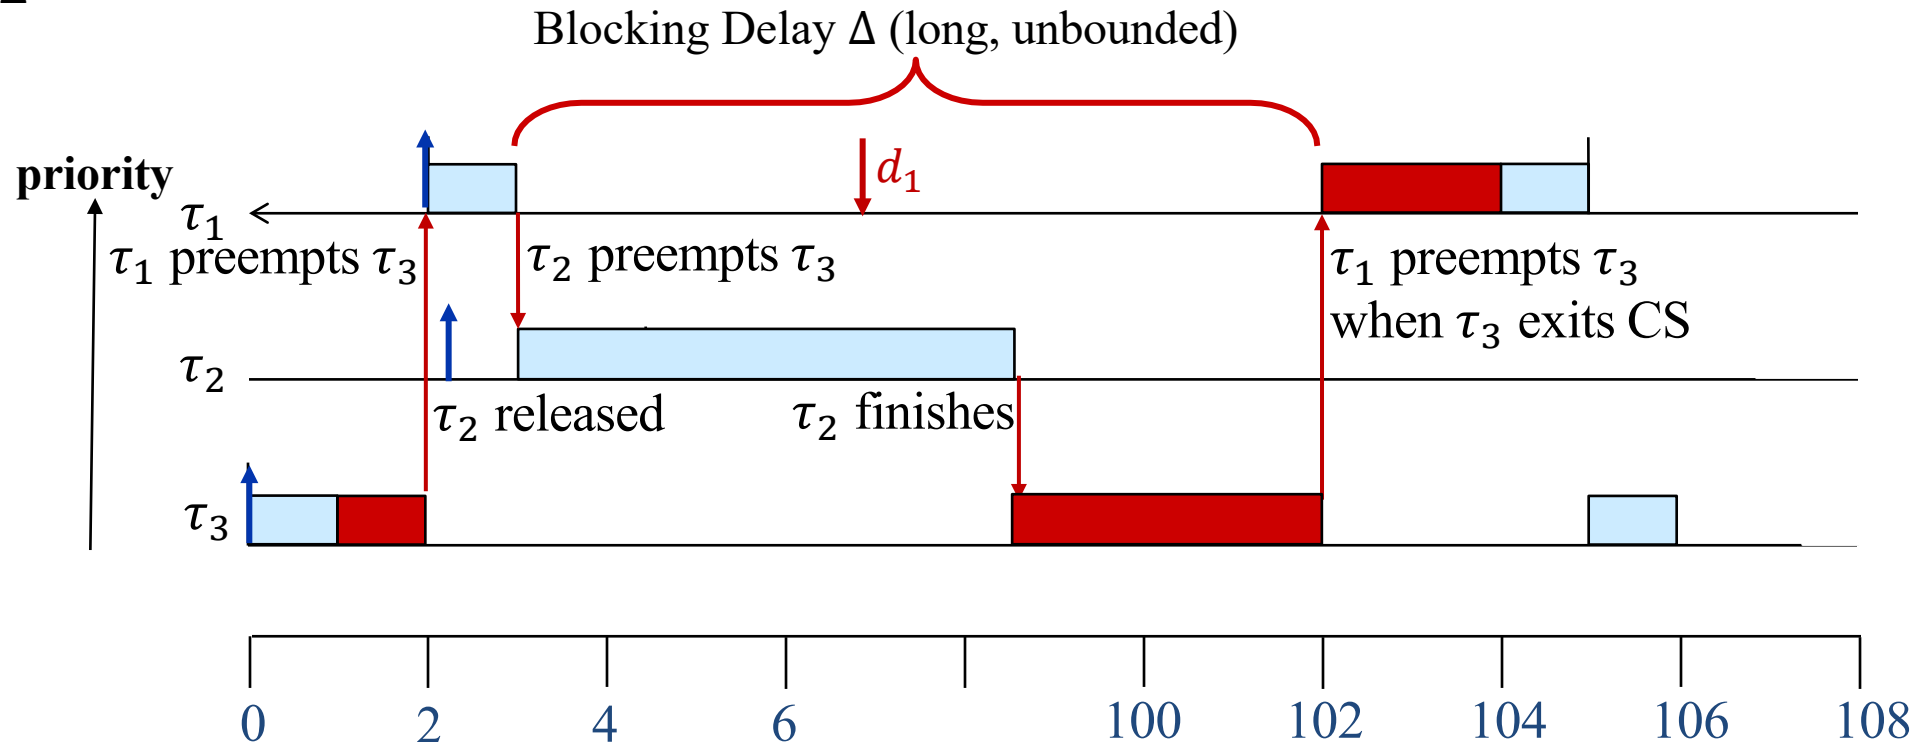
# Priority Inversion I

- Three tasks  $\tau_1, \tau_2, \tau_3$  with priority ordering  $P_1 > P_2 > P_3$ .  $\tau_1, \tau_3$  both require semaphore  $s$ , and  $\tau_2$  does not require any semaphore
- $t=1$ : LP task  $\tau_3$  locks  $s$  and enters CS
- $t=2$ : HP task  $\tau_1$  is released and preempts  $\tau_3$
- $t=3$ : HP task  $\tau_1$  tries to lock  $s$ , but gets blocked by  $\tau_3$  holding  $s$
- $t=4.2$ : Medium Priority (MP) task  $\tau_2$  is released and preempts  $\tau_3$
- $t=100$ : MP task  $\tau_2$  finishes execution after running for its WCET  $C_2$ ;  $\tau_3$  resumes execution in CS
- $t=102$ : LP task  $\tau_3$  unlocks  $s$ ; HP task  $\tau_1$  preempts  $\tau_3$  and finally locks  $s$ , after experiencing a long, unbounded blocking delay  $\Delta$ , and misses its deadline  $d_1$
- This is priority inversion, since MP task  $\tau_2$  causes a long blocking delay to HP task  $\tau_1$ , even though they do not share any resources (semaphores)



# Priority Inversion II

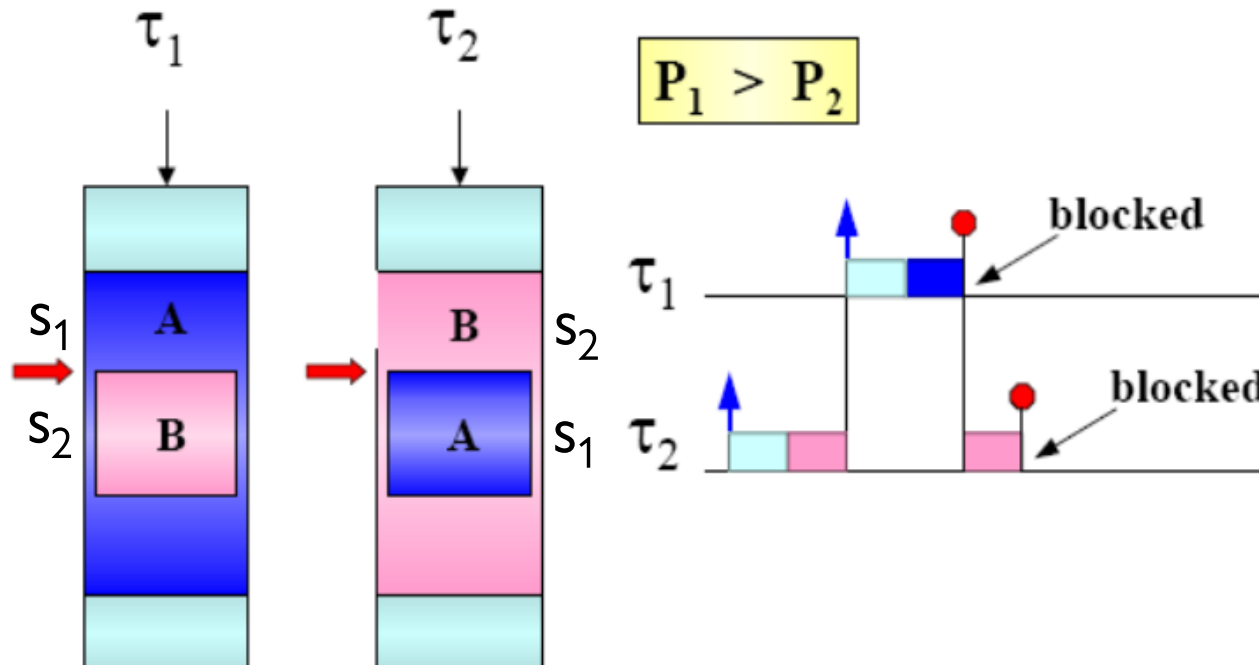
- (This scenario is more realistic and likely than previous one, as MP task  $\tau_2$  may be released anytime during  $\tau_1$ 's execution after it preempts  $\tau_3$ )
- $t=1$ : LP task  $\tau_3$  locks  $s$  and enters CS
- $t=2$ : HP task  $\tau_1$  is released and preempts  $\tau_3$
- $t \in [2, 3]$ : MP task  $\tau_2$  is released, but cannot run since HP task  $\tau_1$  is running
- $t=3$ : HP task  $\tau_1$  tries to lock  $s$ , but gets blocked by  $\tau_3$  holding  $s$ ; MP task  $\tau_2$  starts running
- $t=98.5$ : MP task  $\tau_2$  finishes execution after running for its WCET  $C_2$ ;  $\tau_3$  resumes execution in the CS
- $t=102$ : LP task  $\tau_3$  unlocks  $s$ ; HP task  $\tau_1$  preempts  $\tau_3$  and finally locks  $s$ , after experiencing a long, unbounded blocking Delay  $\Delta$






# Deadlocks

- Classic deadlock scenario: Two tasks  $\tau_1$  and  $\tau_2$  lock two semaphores  $s_1$ ,  $s_2$  in opposite order. ( $s_1$  protects blue CS A and  $s_2$  protects pink CS B)
  - HP task  $\tau_1$  enters blue CS A before pink CS B: ...lock( $s_1$ )...lock( $s_2$ )... unlock( $s_2$ )...unlock( $s_1$ )...
  - LP task  $\tau_2$  enters pink CS B before blue CS A: ...lock( $s_2$ )...lock( $s_1$ )... unlock( $s_1$ )...unlock( $s_2$ )...
  - LP task  $\tau_2$  runs first and locks  $s_2$
  - HP task  $\tau_1$  starts running and locks  $s_1$ , then tries to lock  $s_2$ , gets blocked by  $\tau_2$
  - $\tau_2$  starts running and tries to lock  $s_1$ , but  $\tau_1$  holds  $s_1$ . Circular waiting  $\rightarrow$  deadlock




# Priority Inheritance Protocol (PIP)

- In 1997, this bug caused the Mars pathfinder to freeze up occasionally and then starts working again. Fixed by uploading a software patch enabling Priority-Inheritance Protocol (PIP)
- A task  $\tau_i$  in a CS increases its priority, if it is holding a lock  $s$  and blocks other higher priority tasks, by inheriting the highest priority of all higher-priority tasks  $\tau_k$  blocked waiting for lock  $s$ 
  - $P_{\tau_i \text{ holding } s} = \max\{P_k | \tau_k \text{ blocked on } s\}$
- $t=3$ : HP task  $\tau_1$  tries to enter CS, gets blocked since LP task  $\tau_3$  is in CS;  $\tau_3$  inherits  $\tau_1$ 's high priority, and runs without preemption by MP task  $\tau_2$  (regardless of if  $\tau_2$  is released at  $t \in [2, 3]$  or  $t > 3$ )



# Blocking Time under PIP

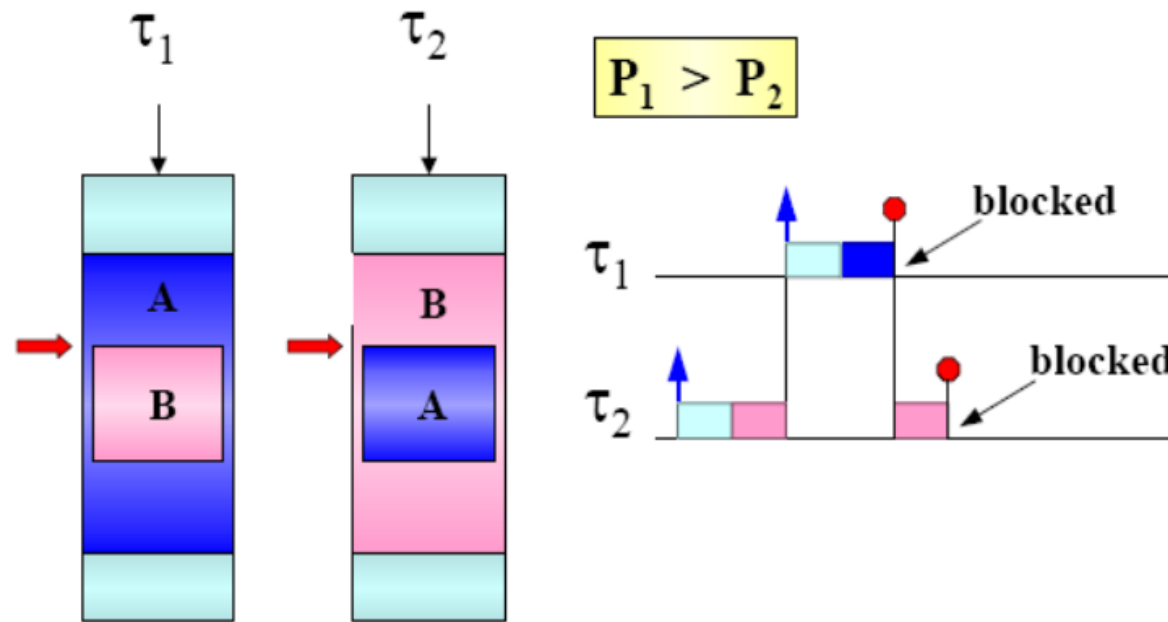
- Under PIP, task  $\tau_i$  may experience two types of blocking delays:
  - **Direct blocking:**  $\tau_i$  tries to lock semaphore  $s$  that is already locked
  - **Push-through blocking:**  $\tau_i$  blocked by lower-priority task that has inherited a higher-priority ( $\tau_i$  itself may not need any semaphores)
- Example:
  - HP task  $\tau_1$  experiences direct blocking by LP task  $\tau_3$  in time interval  $[t_3, t_5]$
  - MP task  $\tau_2$  experiences push-through blocking by LP task  $\tau_3$  in time interval  $[t_4, t_5]$
- PIP analogy: suppose you have checked out a book from library and planned to read it in your spare time. But you got a message from the library that some VIP, say the university president, just got in the waiting queue for the book. You should then hurry up, give the book-reading task a high priority so it is not preempted by other daily chores, finish reading it, and return it to the library quickly, so the VIP is not delayed for a long time.



LP task  $\tau_3$ 's priority is increased at time  $t_3$  when HP task  $\tau_1$  tries to lock semaphore  $s$  but is blocked by  $\tau_3$ ; NOT when  $\tau_3$  lock  $s$  at time  $t_1$

# PIP Pros and Cons

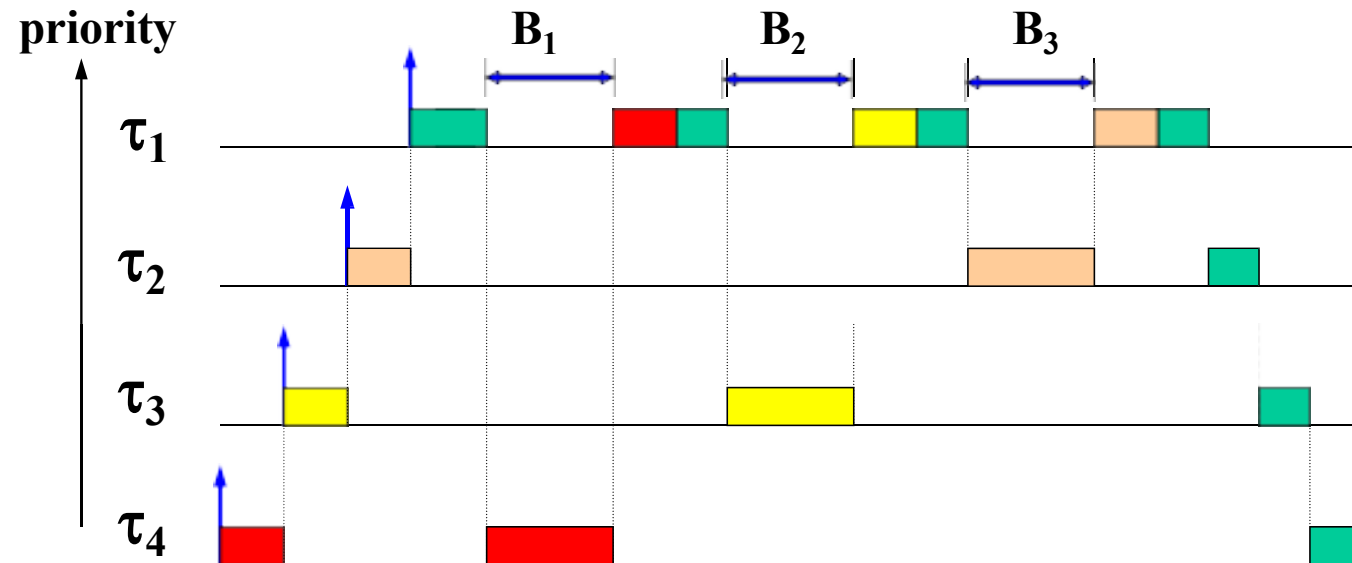
- Pros:
  - It prevents priority inversion
  - It is transparent to the programmer
- Cons:
  - It does not prevent deadlocks and chained blocking



Deadlock still occurs under PIP

# PIP Causes Chained Blocking

- **Chained blocking**: task  $\tau_i$  can be blocked at most once by each lower priority task
- Theorem: Task  $\tau_i$  can be blocked at most for the duration of  $\min(n, m)$  critical sections
  - $n$  is the number of tasks with priority lower than  $\tau_i$
  - $m$  is the number of locks/semaphores on which  $\tau_i$  can be blocked
- In this example, Four tasks and three semaphores ( $s_1$  protects red CS,  $s_2$  protects yellow CS,  $s_3$  protects beige CS). Task  $\tau_i$  is blocked for the duration of  $\min(3, 3) = 3$  critical sections



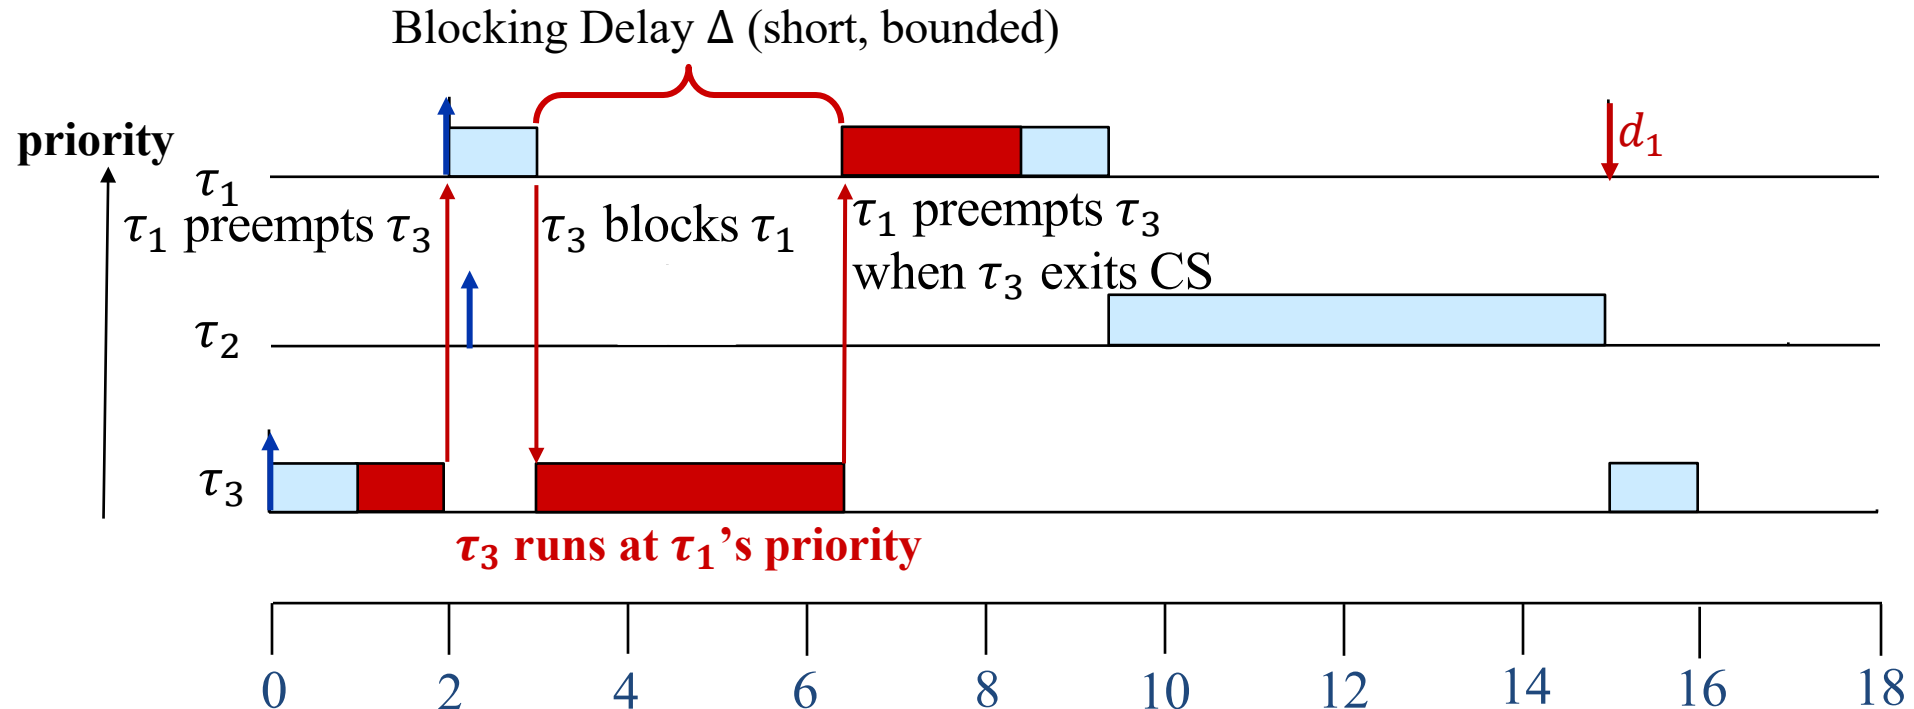
# Priority Ceiling Protocol (PCP)

---

- Assumptions: fixed-priority scheduling; resources required by all tasks are known a priori at design time (not required by PIP)
- Priority Ceiling Protocol PCP = PIP + ceiling blocking
- PIP still holds: When  $\tau_i$  is blocked on  $s_k$ , the lower-priority task currently holding  $s_k$  inherits  $\tau_i$ 's priority
- Each semaphore is assigned a ceiling, equal to maximum priority of all tasks that require it:  $C(s_k) = \max\{P_j: \tau_j \text{ uses } s_k\}$
- Task  $\tau_i$  can acquire  $s_k$  and enter CS only if
  - $P_i > \max\{C(s_k): s_k \text{ locked by other tasks} \neq \tau_i\}$ , that is, task  $\tau_i$ 's priority  $P_i$  is higher than the ceilings of all semaphores currently held by other tasks; otherwise it is blocked on  $s_k$
  - Corollary: If  $s_k$  is currently held by some task, then  $\tau_i$  cannot lock  $s_k$ , since ceiling of  $s_k$  is at least the priority of  $\tau_i$ , i.e.,  $P_i \leq C(s_k)$
- In addition to direct blocking and push-through blocking under PIP, a task may experience ceiling blocking under PCP:  $\tau_i$  tried to lock  $s$ , but ceilings of currently locked semaphores are higher than  $P_i$  ( $s$  itself may be free)
  - Ceiling blocking helps to prevent deadlocks and chained blocking

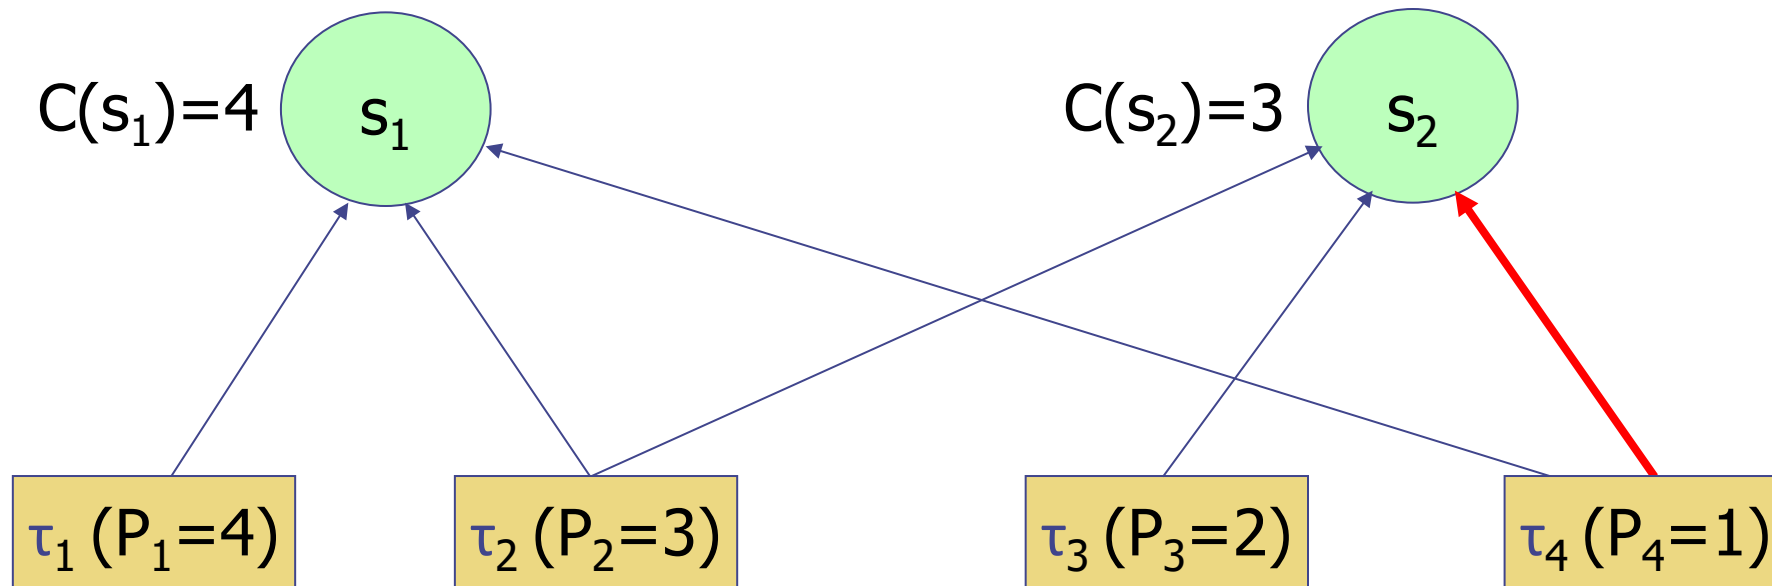
# PCP Example I

- Three tasks  $\tau_1, \tau_2, \tau_3$  with priority ordering  $P_1 > P_2 > P_3$ .  $\tau_1, \tau_2$  both require semaphore  $s$ , and  $\tau_3$  does not require any semaphore
  - $C(s) = \max\{P_j: \tau_j \text{ uses } s\} = \max\{P_1, P_2\} = P_1$
- The execution trace is the same as PIP, since PCP includes PIP as part of the protocol



## PCP Example II

- Four tasks with priority ordering  $P_1 = 4$  (highest) to  $P_2 = 3$ ,  $P_3 = 2$ ,  $P_4 = 1$  (lowest) and two semaphores  $s_1, s_2$ 
  - $C(s_1) = \max\{P_j: \tau_j \text{ uses } s_1\} = \max\{P_1, P_2, P_4\} = 4$
  - $C(s_2) = \max\{P_j: \tau_j \text{ uses } s_2\} = \max\{P_2, P_3, P_4\} = 3$
- When  $\tau_4$  is holding  $s_2$ ,  $\tau_2$  cannot lock  $s_1$ , since  $P_2 = 3 \leq C(s_2) = 3$  (ceiling blocking)
  - Ceiling blocking is “preventive blocking”, since  $\tau_2$  is blocked by  $\tau_4$ , even though  $\tau_2$  tries to lock  $s_1$  which is free. This prevents any potential deadlocks in the future, when  $\tau_2, \tau_4$  each holds one of  $s_1, s_2$  and tries to lock the other
- When  $\tau_4$  is holding  $s_2$ ,  $\tau_1$  can lock  $s_1$ , since  $P_1 = 4 > C(s_2) = 3$ 
  - Since  $\tau_1$  requires  $s_1$  only, there is no cyclic dependency and no deadlock







# PCP Prevents Deadlocks

- Classic deadlock scenario (with or without PIP): Two tasks  $\tau_1$  and  $\tau_2$  lock two semaphores in opposite order:
  - LP task  $\tau_2$  runs first and locks  $s_2$
  - HP task  $\tau_1$  starts running and locks  $s_1$ , then tries to lock  $s_2$ , gets blocked by  $\tau_2$
  - $\tau_2$  starts running and tries to lock  $s_1$  but  $\tau_1$  holds  $s_1$ . Circular waiting  $\rightarrow$  deadlock
- Under PCP,  $C(s_1) = C(s_2) = \max\{P_1, P_2\} = P_1$ . Both semaphores  $s_1$  and  $s_2$  have ceiling equal to  $P_1$ , since they are all required by the higher priority task  $\tau_1$ .
  - LP task  $\tau_2$  runs first and locks  $s_2$
  - HP task  $\tau_1$  runs and preempts  $\tau_2$ . When  $\tau_1$  tries to lock  $s_1$ , it is blocked since its priority does not exceed ceiling of  $s_2$ , i.e.,  $P_1 \leq \text{ceil}(s_2) = P_1$
  - $\tau_2$  will lock both  $s_2$  and  $s_1$ , and exit both CSes before  $\tau_1$  can lock  $s_1$  and  $s_2$ . This prevents circular waiting and deadlock

## Typical Deadlock




## Deadlock avoidance with PCP



# PCP Prevents Chained Blocking


- Three tasks and two semaphores ( $s_1$  protects red CS and  $s_2$  protects yellow CS)
- $C(s_1) = \max\{P_1, P_3\} = P_1$ ,  $C(s_2) = \max\{P_1, P_2\} = P_1$ . Both semaphores  $s_1$  and  $s_2$  have ceiling equal to  $P_1$ , since they are all required by the highest priority task  $\tau_1$ . At time  $t_1$ , LP task  $\tau_3$  is holding  $s_1$  (in red CS). When MP task  $\tau_2$  tries to lock  $s_2$  and enter yellow CS, it is blocked since its priority does not exceed ceiling of  $s_1$ ,  $P_2 \leq C(s_1) = P_1$  (ceiling blocking)
- Hence  $\tau_3$  must unlock  $s_1$  before  $\tau_2$  can lock  $s_2$ . This prevents possible chained blocking



$t_1$ :  $\tau_2$  is blocked by the PCP, since  $P_2 < C(s_1)$

# PCP Prevents Chained Blocking

- Recall the example with chained blocking under PIP
- Four tasks and three semaphores ( $s_1$  protects red CS,  $s_2$  protects yellow CS,  $s_3$  protects beige CS)
- Under PCP:  $C(s_1) = \max\{P_1, P_4\} = P_1$ ,  $C(s_2) = \max\{P_1, P_3\} = P_1$ ,  $C(s_3) = \max\{P_1, P_2\} = P_1$ . All semaphores  $s_1, s_2, s_3$  have ceiling equal to  $P_1$ , since they are all required by the highest priority task  $\tau_1$ .
- While  $\tau_4$  is holding  $s_1$  (in the red CS),  $\tau_3$  cannot lock  $s_2$ , since  $P_3 \leq C(s_1) = P_1$ ; and  $\tau_2$  cannot lock  $s_3$ , since  $P_2 \leq C(s_1) = P_1$  (ceiling blocking)
- Hence PCP prevents chained blocking, since task  $\tau_1$  is blocked at most once by a lower-priority task (either  $\tau_4$ , or  $\tau_3$ , or  $\tau_2$ )



## PCP Blocking Time

---

A given task  $i$  is blocked (or delayed) by at most one critical section of any lower priority task locking a semaphore with priority ceiling greater than or equal to the priority of task  $i$ . We can explain that mathematically using the notation:

$$B_i = \max_{\{k,s \mid k \in lp(i) \wedge s \in used\_by(k) \wedge ceil(s) \geq pri(i)\}} CS_{k,s} \quad (4.5)$$

- Consider all lower-priority tasks ( $k \in lp(i)$ ), and the semaphores they can lock ( $s$ )
- Select from those semaphores ( $s$ ) with ceiling higher than or equal to  $pri(i) = P_i$
- Take max length of all tasks ( $k$ )'s critical sections that lock semaphores ( $s$ )
- (The blocking time is valid even for a task that does not require any semaphores/critical sections, as it may experience push-through blocking.)

# PCP Pros and Cons

---

- Pros:
  - It prevents priority inversion, deadlocks, and chained blocking
  - Any given task is blocked at most once by a lower-priority task
- Cons:
  - It is not transparent to the programmer, as shared resources required by all tasks must be known a priori at design time, and programmer needs to calculate priority ceilings of all semaphores and pass them to the OS (PIP does not need this step)

	Deadlock Prevention	Number of blockings	Programmer Transparency
PIP	No	$\min(n, m)$	Yes
PCP	Yes	1	No

# blockings under PIP:  $n$  is the number of tasks with priority lower than  $\tau_i$ ;  $m$  is the number of locks/semaphores on which  $\tau_i$  can be blocked

# Schedulability Analysis under PIP and PCP

---

- Let  $B_i$  denote the maximum blocking time experienced by task  $\tau_i$  due to shared resources
- Schedulable utilization bound for RM scheduling with blocking time (sufficient condition):
  - A taskset is schedulable under RM scheduling with blocking time if
  - $\forall i$ , priority level  $i$  utilization  $U_i = \sum_{j \in hp(i)} \frac{C_j}{T_j} + \frac{C_i + B_i}{T_i} \leq i(2^{1/i} - 1)$
- Response Time Analysis (RTA) for RM scheduling with blocking time (necessary and sufficient condition):
  - Task  $\tau_i$ 's WCRT  $R_i$  is computed by solving the following recursive equation to find the minimum fixed-point solution:
    - $R_i = C_i + B_i + \sum_{j \in hp(i)} \left\lceil \frac{R_i}{T_j} \right\rceil C_j$
    - $\tau_i$  is schedulable iff  $R_i \leq D_i$

## Example Taskset (without shared resources)

---

- System utilization  $U = \frac{5}{50} + \frac{250}{500} + \frac{1000}{3000} = 0.933 > 0.780$ 
  - Since utilization exceeds the Utilization Bound of 0.780 of 3 tasks under RM scheduling, we cannot determine schedulability by the Utilization Bound test
- RTA shows that the taskset is schedulable by computing WCRT of each task:
  - $R_1 = C_1 + 0 = 5 + 0 = 5 \leq D_1 = 50$
  - $R_2 = C_2 + \left\lceil \frac{R_2}{T_1} \right\rceil \cdot C_1 = 250 + \left\lceil \frac{R_2}{50} \right\rceil \cdot 5 = 280 \leq D_2 = 500$
  - $R_3 = C_3 + \left\lceil \frac{R_3}{T_1} \right\rceil \cdot C_1 + \left\lceil \frac{R_3}{T_2} \right\rceil \cdot C_2 = 1000 + \left\lceil \frac{R_3}{50} \right\rceil \cdot 5 + \left\lceil \frac{R_3}{500} \right\rceil \cdot 250 = 2500 \leq D_3 = 3000$

Task	T	D	C	Prio
1	50	50	5	H
2	500	500	250	M
3	3000	3000	1000	L

# Example Taskset (with shared resources under PCP) I

- 3 semaphores  $s_1, s_2, s_3$ 
  - Task 1 requires semaphore  $s_1$ , with critical section length 1
  - Task 2 requires semaphores  $s_2$  and  $s_3$ , with critical section lengths 2 and 5, respectively
  - Task 3 requires semaphores  $s_2$  and  $s_3$ , with critical section lengths 3 and 4, respectively
- Ceilings  $C(s_1) = P_1 = H$ ;  $C(s_2) = C(s_3) = \max(P_2, P_3) = M$
- Blocking times:
  - Task 1:  $B_1 = 0$  (Task 1's priority is higher than ceilings of  $s_2$  and  $s_3$ ), so it remains schedulable
  - Task 2:  $B_2 = \max(3, 4) = 4$  (maximum CS length of LP Task 3)
    - » Utilization  $U_2 = \sum_{j \in hp(2)} \frac{C_j}{T_j} + \frac{C_2 + B_2}{T_2} = \frac{5}{50} + \frac{250+4}{500} = 0.608 \leq 0.828$  (utilization bound for 2 tasks under RM)
    - » Or WCRT:  $R_2 = C_2 + B_2 + \left\lceil \frac{R_2}{T_1} \right\rceil \cdot C_1 = 250 + 4 + \left\lceil \frac{R_2}{50} \right\rceil \cdot 5 = 284 \leq D_2 = 500$
  - Task 3:  $B_3 = 0$  (Task 3 is the lowest priority task, so it does not experience any blocking), so it remains schedulable
- The taskset remains schedulable with shared resources under PCP

Task	T	D	C	Prio	sems	CS Len
1	50	50	5	H	$s_1$	1
2	500	500	250	M	$s_2, s_3$	2, 5
3	3000	3000	1000	L	$s_2, s_3$	3, 4



# Example Taskset (with shared resources under PCP) II

- 3 semaphores  $s_1, s_2, s_3$ 
  - Task 1 requires semaphores  $s_1, s_2$  and  $s_3$  with CS lengths 1, 1, 1
  - Task 2 requires semaphores  $s_2$  and  $s_3$ , with CS lengths 2 and 5, respectively
  - Task 3 requires semaphores  $s_2$  and  $s_3$ , with CS lengths 3 and 4, respectively
- Ceilings  $C(s_1) = C(s_2) = C(s_3) = \max(P_1, P_2, P_3) = H$
- Blocking times:
  - Task 1:  $B_1 = \max(2, 5, 3, 4) = 5$  (maximum CS length of LP Tasks 2 and 3, associated with semaphores  $s_2$  and  $s_3$ )
    - » Utilization  $U_1 = \sum_{j \in hp(1)} \frac{C_j}{T_j} + \frac{C_1 + B_1}{T_1} = 0 + \frac{5+5}{50} = 0.2 \leq 1$  (Utilization bound for 1 task under RM), so Task 1 remains schedulable
    - » Or WCRT:  $R_1 = C_1 + B_1 = 5 + 5 = 10 \leq D_1 = 50$
    - » (Task 1's CS lengths (1, 1, 1) do not matter since it is the highest priority task and does not block any other task)
  - Task 2:  $B_2 = \max(3, 4) = 4$  (maximum CS length of LP Task 3)
    - » Same calculation of utilization and WCRT as before
  - Task 3:  $B_3 = 0$  (Task 3 is the lowest priority task, so it does not experience blocking), so it remains schedulable
- The taskset remains schedulable with shared resources under PCP

Task	T	D	C	Prio	sems	CS Len
1	50	50	5	H	$s_1, s_2, s_3$	1, 1, 1
2	500	500	250	M	$s_2, s_3$	2, 5
3	3000	3000	1000	L	$s_2, s_3$	3, 4

# Scheduling Anomaly w/ Resource Synchronization

- Doubling processor speed causes T1 to miss its deadline
  - (Yellow part denotes a critical section shared by T1 and T2)

