

# CSC 112: Computer Operating Systems

## Lecture 6

### Real-Time Scheduling

Department of Computer Science,  
Hofstra University

# Outline

---


- Introduction to RTOS and Real-Time Scheduling
- Fixed-Priority Scheduling
- Earliest Deadline First Scheduling
- Least Laxity First (LLF) Scheduling
- Resource Synchronization Protocols (for Fixed-Priority Scheduling)
- Preemptive vs. Non-Preemptive Scheduling
- Multiprocessor Scheduling

---

# Introduction to RTOS and Real-Time Scheduling

# Embedded Control Systems

- An embedded control system consists of:
  - The system-under-control (SUT)
    - » may include sensors and actuators
  - The controller/computer
    - » sends signals to the system according to a predetermined control objective
- In the old days, each control task runs on a dedicated CPU
  - No RTOS, bare metal
  - No need for scheduling
  - Just make sure that task execution time < deadline
- Now, multiple control tasks share one CPU
  - Multitasking RTOS
  - Need scheduling to make sure all tasks meet deadlines



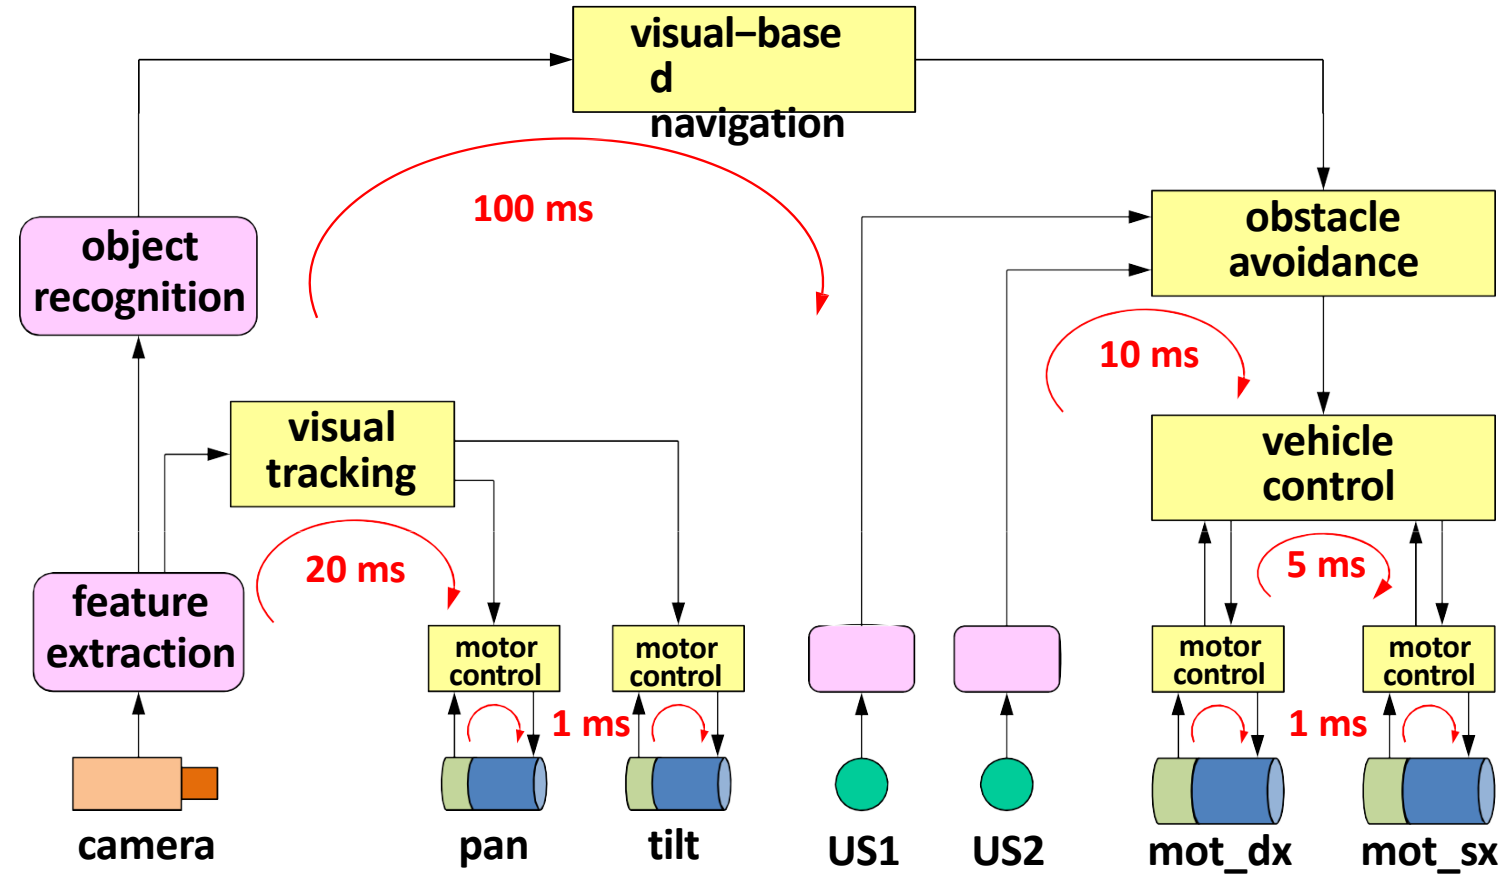

# Requirements

---

- The tight interaction with the environment requires the system to react to events within precise timing constraints
- Timing constraints are imposed by the dynamics of the environment
- The real-time operating system (RTOS) must be able to execute tasks within timing constraints

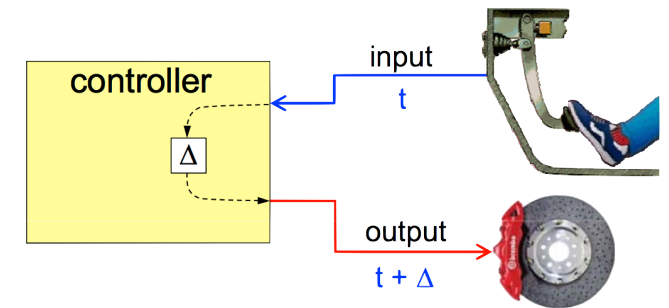

# A Robot Control Example

- Consider a robot equipped with:
  - two actuated wheels
  - two proximity (US) sensors
  - a mobile (pan/tilt) camera
  - a wireless transceiver
- Goal:
  - follow a path based on visual feedback
  - avoid obstacles



# Real-Time Systems

- A computer system that is able to respond to events within precise timing constraints
- A system where the correctness depends not only on the output values, but also on the time at which results are produced
- A real-time system is not necessarily a real fast system
  - Speed is always relative to a specific environment
  - Running faster is good, but does not guarantee hard real-time constraints
- The objective of a real-time system is to guarantee the worst-case timing behaviour of each individual task
- The objective of a fast system is to optimize the average-case performance
  - A system with fast average-case performance may not meet worst-case requirements, e.g.,
  - There was a person who drowned in a river with average depth of 6 inches



# RTOS Requirements

---

- Timeliness: results must be correct not only in their value but also in the time domain
  - provide kernel mechanism for time management and for handling tasks with explicit timing constraints and different criticality
- Predictability: system must be analyzable to predict the consequences of any scheduling decision
  - if some task cannot be guaranteed within time constraints, system must notify this in advance, to handle the exception (plan alternative actions)
- Efficiency: operating system should optimize the use of available resources (computation time, memory, energy)
- Robustness: must be resilient to peak-load conditions
- Fault tolerance: single software/hardware failures should not cause the system to crash
- Maintainability: modular architecture to ensure that modifications are easy to perform





# Sources of Nondeterminism

---

- Architecture
  - cache, pipelining, interrupts, DMA
- Operating System (our focus in this lecture)
  - scheduling, synchronization, communication
- Language
  - lack of explicit support for time
- Design Methodologies
  - lack of analysis and verification techniques


# Task

- The concept of concurrent tasks reflects the intuition about the functionality of embedded systems.
  - Task here can refer to either process or thread, depending on the underlying RTOS support
- Tasks help us manage timing complexity:
  - multiple execution rates
    - » multimedia
    - » automotive
  - asynchronous input
    - » user interfaces
    - » communication systems



# Schedule


- A specific assignment of tasks to the processor that determines the task execution sequence. Formally:
- Given a task set  $\Gamma = \{\tau_1, \dots, \tau_n\}$ , a schedule is a function  $\sigma: R^+ \rightarrow N$  that associates an integer  $k$  to each time slice  $[t_i, t_{i+1})$  with the meaning:
  - $k = 0$ : in  $[t_i, t_{i+1})$  the processor is idle
  - $k > 0$ : in  $[t_i, t_{i+1})$  the processor executes  $\tau_k$



At times  $t_1, t_2, \dots$ : context switch to a different task

# Preemptive vs. Nonpreemptive Scheduling

- A scheduling algorithm is:
  - preemptive: if the active job can be temporarily suspended to execute a more important job
  - non-preemptive: if the active job cannot be suspended, i.e., always runs to completion



Preemptive scheduling example

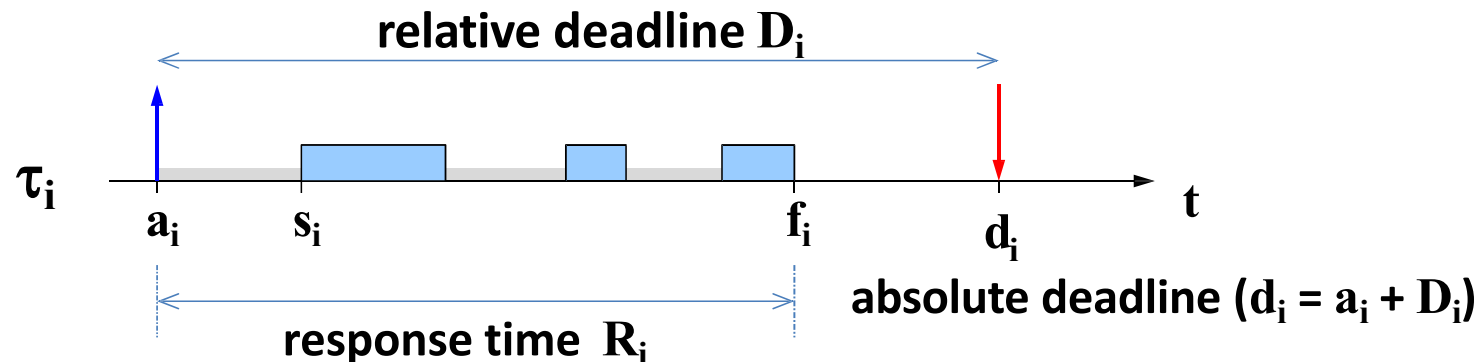
# Definitions

---

- Feasible schedule
  - A schedule  $\sigma$  is said to be feasible if all the tasks can complete according to a set of specified constraints.
- Schedulable set of tasks
  - A set of tasks  $\Gamma$  is said to be schedulable if there exists at least one algorithm that can produce a feasible schedule for it.
- Hard real-time task: missing deadline may have catastrophic consequences, so deadline violations are not permitted. A system able to handle hard real-time tasks is a hard real-time system
  - sensory acquisition
  - low-level control
  - sensory-motor planning
- Soft real-time task: missing deadlines causes Quality-of-Service(QoS)/performance degradation, so deadline violations are expected and permitted
  - reading data from the keyboard—user command interpretation
  - message displaying
  - graphical activities

# Real-Time Task


- A task characterized by a timing constraint on its response time, called deadline:
  - relative deadline  $D_i$ : part of task attribute definition, measured from task arrival time  $a_i$
  - Absolute deadline  $d_i = a_i + D_i$ : measured from some absolute reference time point 0
- Definition: feasible task
  - A real-time task  $\tau_i$  is said to be feasible if it completes within its absolute deadline, that is, if  $f_i \leq d_i$ , or, equivalently, if  $R_i \leq D_i$



# Tasks and Jobs


---

- A task running several times on different input data generates a sequence of instances (jobs)
  - Upwards arrow: task arrival or release times; downwards arrow: task deadlines
- Activation mode:
  - Periodic tasks: the task is activated by the operating system at predefined time intervals
  - Aperiodic tasks: the task is activated at an event arrival



# Estimating WCET is Not Easy


- Each job operates on different data and can take different paths.
- Even for the same data, computation time depends on processor state (cache state, number of preemptions).
- We use  $C_i$  to denote  $C_i^{max}$  Worst-Case Execution Time (WCET) in this lecture, and assume it is given as part of task parameters.





# Predictability/Safety vs. Efficiency


- Tradeoff between safety and efficiency in estimating the WCET  $C_i$ 
  - Setting a large  $C_i$  achieves high predictability and safety, since it is unlikely to be exceeded at runtime; but it hurts efficiency, since the system needs to reserve more CPU time for the task. Suitable for hard real-time tasks.
  - Setting a small  $C_i$  achieves high efficiency, but hurts safety, since the task may execute for more than its  $C_i$  estimate. Suitable for soft real-time tasks.




# Jitter

- It is a measure of the time variation of a periodic event:


## Finish-time Jitter



## Start-time Jitter


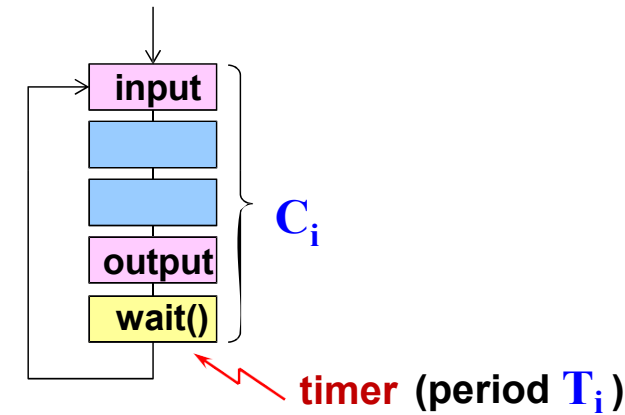


## Completion-time Jitter (I/O Jitter)



# Periodic Task


- A **periodic task**  $\tau_i$  has a tuple of 3 attributes  $(C_i, T_i, D_i)$ :
  - Worst-Case Execution Time (WCET)  $C_i$ ; Period  $T_i$ ; Relative Deadline  $D_i$
  - Implicit deadline if  $D_i = T_i$ ; Constrained deadline if  $D_i \leq T_i$
- It generates an infinite sequence of **jobs** in every period:  $\tau_{i,1}, \tau_{i,1}, \dots, \tau_{i,k}, \dots$



$r_i$  release time (arrival time  $a_i$ )  
 $s_i$  start time  
 $C_i$  worst-case execution time (wcet)  
 $d_i$  absolute deadline  
 $D_i$  relative deadline  
 $f_i$  finishing time

A job of task  $\tau_i$

$\tau_i (C_i, T_i, D_i)$



task phase or  
Release offset

$$\begin{aligned}
 a_{i,k} &= \Phi_i + (k-1) T_i \\
 d_{i,k} &= a_{i,k} + D_i
 \end{aligned}$$

often  
 $D_i = T_i$


Multiple jobs released by task  $\tau_i$

# Aperiodic & Sporadic Task

- Aperiodic task: jobs may arrive at arbitrary time instants
- Sporadic task: arrival times with a minimum interarrival time constraint

- **Aperiodic:**  $a_{i,k+1} > a_{i,k}$
- **Sporadic:**  $a_{i,k+1} \geq a_{i,k} + T_i$

minimum  
interarrival time



# Types of Constraints


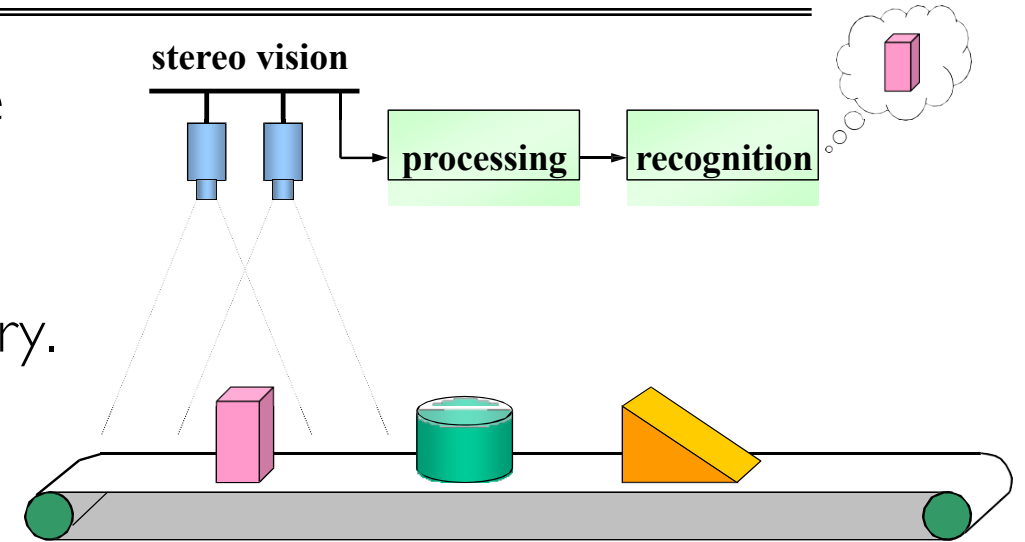
---

- **Timing constraints**
  - Deadline, jitter
- **Precedence constraints**
  - Relative ordering among task executions
- **Resource constraints**
  - Synchronization when accessing mutually-exclusive resources (shared data)

# Precedence Constraints


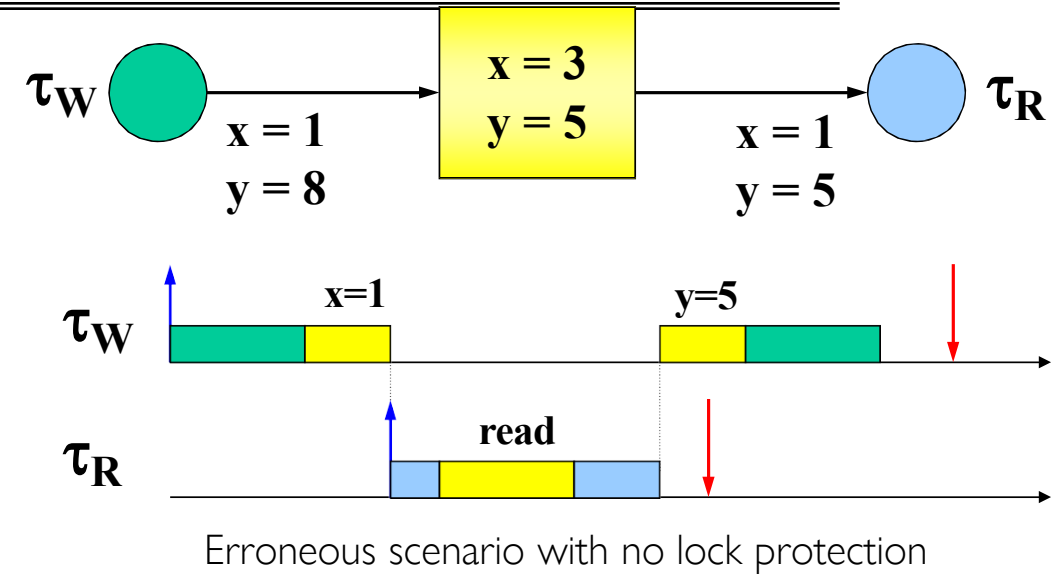
- Tasks must be executed with specific precedence relations, specified by a Directed Acyclic Graph (Precedence Graph)
- Example application of parts inspection in a factory.  
Tasks:

- Image acquisition (acq1, acq2)
- Edge detection (edge1, edge2)
- Shape detection (shape), pixel disparities (disp)
- Height determination (height), recognition (rec)




# Resource Constraints

- To ensure data consistency, shared data must be accessed in mutual exclusion
- Example: the writer task  $\tau_W$  writes to variables  $x$  and  $y$ ; the reader task  $\tau_R$  reads  $x$  and  $y$ . The pair of variables  $(x, y)$  should be updated atomically, i.e.,  $\tau_R$  should read either  $(x, y) = (1, 8)$  or  $(x, y) = (3, 5)$ .
- Left upper: an erroneous scenario when  $\tau_R$  reads a set of inconsistent values  $(x, y) = (3, 5)$ .
- Left lower: protecting the critical section (yellow parts) with a mutex lock ensures atomicity.



# Scheduling Metrics

- Lateness  $L_i = f_i - d_i$  represents the delay of a task completion with respect to its deadline; if a task completes before the deadline, its lateness is negative.
- Tardiness or exceeding time  $E_i = \max(0, L_i)$  is the time a task stays active after its deadline; if a task completes before the deadline, its tardiness is 0.





$r_i$  release time (arrival time  $a_i$ )  
 $s_i$  start time  
 $C_i$  worst-case execution time (wcet)  
 $d_i$  absolute deadline  
 $D_i$  relative deadline  
 $f_i$  finishing time



## Example: Lateness

- Which schedule is better depends on application requirements:
- In (a), the maximum lateness is minimized with  $L_{max} = 3$ , but all jobs  $J_1$  to  $J_5$  miss their deadlines.
- In (b), the maximal lateness is larger with  $L_{max} = 23$ , but only one job  $J_1$  misses its deadline.




# Scheduling Algorithms

---

- Static cyclic scheduling (offline)
  - All task invocation times are computed offline and stored in a table; Runtime dispatch is a simple table lookup
- Online scheduling;
  - Fixed priority scheduling
    - » Each task is assigned a fixed priority; Runtime dispatch is priority-based
  - Dynamic priority scheduling
    - » Task priorities are assigned dynamically at runtime, e.g., Earliest Deadline First (EDF), Least-Laxity First (LLF)
  - Non-real-time scheduling, e.g., round-robin, multi-level queue...

# Static Cyclic Scheduling

- The same schedule is executed once during each hyper-period (least common multiple of all task periods in a taskset).
  - The hyper-period is partitioned into frames of length  $f$ .
    - » If a task's WCET exceeds  $f$ , then programmer needs to cut it to fit within a frame, and save/restore program state manually
  - The schedule is computed offline and stored in a table. Runtime task dispatch is a simple table lookup.
- Pros:
  - Deals with precedence, exclusion, and distance constraints
  - Efficient, low-overhead for runtime task dispatch
  - Lock-free at runtime
- Cons:
  - Task table can get very large if task periods are relatively prime
  - Maintenance nightmare: complete redesign when new tasks are added, or old tasks are deleted
- Not widely used
  - Except in certain safety-critical systems such as avionic systems



---

# Fixed-Priority Scheduling

# Fixed Priority Scheduling

---


- Each task is assigned a fixed priority for all its invocations
- Pros:
  - Predictability
  - Low runtime overhead
  - Temporal isolation during overload
- Cons:
  - Cannot achieve 100% utilization in general, except when task periods are harmonic
- Widely used in most commercial RTOSes and CAN bus

# Rate Monotonic & Deadline Monotonic Scheduling

- Rate Monotonic Scheduling (RMS)
  - Assign higher priority to task with smaller period
  - When  $D = T$ , RMS is the optimal priority assignment, i.e., if a taskset is not schedulable with RMS priority assignment, then it is not schedulable with any other fixed priority assignment

- Deadline Monotonic Scheduling (DMS)
  - Assign higher priority to task with smaller relative deadline (indicated by red downward arrow for  $\tau_2$ )
  - When  $D < T$ , DMS is the optimal priority assignment

- Why do we want deadline  $<$  period?
  - Some events happen infrequently, but need to be handled urgently



# Two Schedulability Analysis Approaches


---

- Utilization bound test
  - Calculate total CPU utilization and compare it to a known bound
  - Polynomial time complexity
  - Pessimistic: sufficient but not necessary condition for schedulability
- Response Time Analysis (RTA)
  - Calculate Worst-Case Response Time  $R_i$  for each task  $\tau_i$  and compare it to its deadline  $D_i$
  - Pseudo-polynomial time complexity
    - » An algorithm runs in pseudo-polynomial time if its running time is polynomial in the numeric value of the input (which is exponential in the length of the input – its number of digits).
  - Accurate: necessary and sufficient condition for schedulability

# Utilization Bound Test

# Tasks	RM Util Bound
1	1.00
2	0.828
3	0.780
4	0.757
5	0.743
10	0.718
inf	0.693

- A taskset is schedulable under RM scheduling if its system utilization  $U = \sum_{i=1}^N \frac{C_i}{T_i} \leq N(2^{1/N} - 1)$ 
  - $U \rightarrow 0.69$  as  $N \rightarrow \infty$
  - Assumptions: task period equal to deadline ( $P_i = D_i$ ); task with smaller period  $P_i$  is assigned higher priority (RM priority assignment); tasks are independent (no resource sharing)
- Sufficient but not necessary condition
  - Guaranteed to be schedulable if test succeeds
  - May still be schedulable even if test fails
  - If periods are harmonic (larger period divisible by smaller periods), then utilization bound is 1






# Utilization Bound Test Examples

We use the notation  $\tau_i(C_i, T_i, D_i)$  to denote task  $\tau_i$  with WCET  $C_i$  Period  $T_i$ , Deadline  $D_i$

Taskset  $\tau_1(3, 6, 6), \tau_2(4, 9, 9)$

non-schedulable


$$U = \frac{3}{6} + \frac{4}{9} = 0.944 > 0.828$$



Taskset  $\tau_1(3, 6, 6), \tau_2(3, 9, 9)$

schedulable (UB test is sufficient but not necessary condition)


$$U = \frac{3}{6} + \frac{3}{9} = 0.833 > 0.828$$



Taskset  $\tau_1(2, 4, 4), \tau_2(4, 8, 8)$


schedulable (periods are harmonic)

$$U = \frac{2}{4} + \frac{4}{8} = 1.0 > 0.828$$



# Response Time Analysis (RTA)

- Assumptions:
  - Consider the synchronous taskset: all tasks are initially released at time 0 simultaneously, called the critical instant. This is the worst-case when each task experiences maximum amount of interference from higher priority tasks: if the taskset is schedulable with this assumption, then it will be schedulable for any other release offset.
  - No resource sharing (no critical sections)
  - Figure shows task  $\tau_2$  has longer response time  $R_2$  when released at time 0, the critical instant, simultaneously with higher priority task  $\tau_1$



## Response Time Analysis (RTA)

---

- For each task  $\tau_i$ , compute its Worst-Case Response Time (WCRT)  $R_i$  and compare to its deadline  $D_i$ .  $\tau_i$  is schedulable if  $R_i \leq D_i$
- The taskset is schedulable if all tasks are schedulable.
- WCRT  $R_i$  is computed by solving the following recursive equation to find the minimum fixed-point solution:

- $$R_i = C_i + \sum_{\forall j \in hp(i)} \left\lceil \frac{R_i}{T_j} \right\rceil C_j$$

- where  $hp(i)$  is the set of tasks with higher priority than task  $\tau_i$ .
- $\lceil \cdot \rceil$  is the ceiling operator, e.g.,  $\lceil 1.1 \rceil = 2$ ,  $\lceil 1.0 \rceil = 1$
- $\left\lceil \frac{R_i}{T_j} \right\rceil$  denotes the number of times HP task  $\tau_j$  pre-empts  $\tau_i$  during its one job execution;  $\left\lceil \frac{R_i}{T_j} \right\rceil C_j$  denotes the total preemption delay caused by HP task  $\tau_j$  to  $\tau_i$  during its one job execution

## An Example Taskset

---


- Consider a taskset of 3 task with  $(C_i, T_i, D_i)$  of  $(10, 30, 30), (10, 40, 40), (12, 52, 52)$ . Under RM, task priorities are assigned to be High for T1, Medium for T2, and Low for T3
- System Utilization  $U = \sum_{i=1}^3 \frac{C_i}{T_i} = \frac{10}{30} + \frac{10}{40} + \frac{12}{52} = 0.81 > 0.78$ 
  - Utilization Bound  $(N = 3) = 3 * (2^{1/3} - 1) = 0.78$
- Utilization bound test fails, but taskset is actually schedulable

Task	T=D	C	Prio
T1	30	10	H
T2	40	10	M
T3	52	12	L

## Task T1

Task	T=D	C	Prio
T1	30	10	H
T2	40	10	M
T3	52	12	L


- T1 is the highest priority task, with no interference from other tasks  $hp(1) = \emptyset$
- $R_1 = C_1 + 0 = 10$
- $R_1 < D_1$ , so T1 is schedulable



## Task T2

Task	T=D	C	Prio
T1	30	10	H
T2	40	10	M
T3	52	12	L


- T2 is the medium priority task, with interference from higher priority Task 1  $hp(2) = 1$
- $R_2 = C_2 + \lceil \frac{R_2}{T_1} \rceil * C_1 = 10 + \lceil \frac{R_2}{30} \rceil * 10$
- Solve for  $R_2$  iteratively, starting with initial value  $R_2 = C_2 = 10$ :
  - Iteration 1:  $R_2 = 10 + \lceil \frac{10}{30} \rceil * 10 = 10 + 1 * 10 = 20$
  - Iteration 2:  $R_2 = 10 + \lceil \frac{20}{30} \rceil * 10 = 10 + 1 * 10 = 20$
- Hence  $R_2 = 20 < D_2 = 40$ , so T2 is schedulable



## Task T3

Task	T=D	C	Prio
T1	30	10	H
T2	40	10	M
T3	52	12	L


- T3 is the lowest priority task, with interference from higher priority tasks  $hp(3) = \{1,2\}$
- $R_3 = C_3 + \lceil \frac{R_3}{T_1} \rceil * C_1 + \lceil \frac{R_3}{T_2} \rceil * C_2 = 12 + \lceil \frac{R_3}{30} \rceil * 10 + \lceil \frac{R_3}{40} \rceil * 10$
- Solve for  $R_3$  iteratively, starting with initial value  $R_3 = C_3 = 12$ :
  - Iteration 1:  $R_3 = 12 + \lceil 12/30 \rceil * 10 + \lceil 12/40 \rceil * 10 = 32$
  - Iteration 2:  $R_3 = 12 + \lceil 32/30 \rceil * 10 + \lceil 32/40 \rceil * 10 = 42$
  - Iteration 3:  $R_3 = 12 + \lceil 42/30 \rceil * 10 + \lceil 42/40 \rceil * 10 = 52$
  - Iteration 4:  $R_3 = 12 + \lceil 52/30 \rceil * 10 + \lceil 52/40 \rceil * 10 = 52$
- Hence  $R_3 = 52 \leq D_3 = 52$ , so T3 is schedulable



## RTA for T3: Initial Condition

Task	T=D	C	Prio
T1	30	10	H
T2	40	10	M
T3	52	12	L

- Initially  $R_3 = 12$
- We have not taken into account any preemption delays from higher priority tasks T1 and T2 yet






## RTA for Task 3: Iteration 1

Task	T=D	C	Prio
T1	30	10	H
T2	40	10	M
T3	52	12	L


- $R_3 = 12 + \left\lceil \frac{12}{30} \right\rceil * 10 + \left\lceil \frac{12}{40} \right\rceil * 10$
- $= 12 + 1 * 10 + 1 * 10 = 32$
- T1 preempts T3 once, and T2 preempts T3 once
  - since all 3 tasks are released at time 0 (synchronous release time assumption), and T1 and T2 have higher priority than T3



## RTA for Task 3: Iteration 2

Task	T=D	C	Prio
T1	30	10	H
T2	40	10	M
T3	52	12	L


- $R_3 = 12 + \left\lceil \frac{32}{30} \right\rceil * 10 + \left\lceil \frac{32}{40} \right\rceil * 10$
- $= 12 + 2 * 10 + 1 * 10 = 42$
- T1 preempts T3 twice, and T2 preempts T3 once
  - Since T3 has not finished execution at time 30, and another job of higher priority task T1 is released at time 30 and preempts T3



## RTA for Task 3: Iteration 3

Task	T=D	C	Prio
T1	30	10	H
T2	40	10	M
T3	52	12	L


- $R_3 = 12 + \left\lceil \frac{42}{30} \right\rceil * 10 + \left\lceil \frac{42}{40} \right\rceil * 10$
- $= 12 + 2 * 10 + 2 * 10 = 52$
- T1 preempts T3 twice, and T2 preempts T3 twice
  - Since T3 has not finished execution at time 40, and another job of higher priority task T2 is released at time 40 and preempts T3



## RTA for Task 3: Iteration 4

Task	T=D	C	Prio
T1	30	10	H
T2	40	10	M
T3	52	12	L

- $R_3 = 12 + \left\lceil \frac{52}{30} \right\rceil * 10 + \left\lceil \frac{52}{40} \right\rceil * 10 = 12 + 2 * 10 + 2 * 10 = 52$
- T1 preempts T3 twice, and T2 preempts T3 twice
  - Since T3 has finished execution at time 52, and the next arrivals of T1 and T2 are at time 60 and 80, respectively, so T3 will not experience additional preemptions from T1 and T2
- Now the recursive equation has converged, and we have obtained the WCRT of T3  $R_3 = 52 \leq D_3 = 52$



## When T3 is Unschedulable


- The recursive equation may not converge, i.e., a task's WCRT may be infinity, e.g., suppose we change T1's WCET to be 20, then:
- $$R_3 = C_3 + \lceil \frac{R_3}{T_1} \rceil * C_1 + \lceil \frac{R_3}{T_2} \rceil * C_2 = 12 + \lceil \frac{R_3}{30} \rceil * 10 + \lceil \frac{R_3}{40} \rceil * 20$$
- Solve for  $R_3$  iteratively, starting with initial value  $R_3 = C_3 = 12$ :
  - Iteration 1:  $R_3 = 12 + \lceil 12/30 \rceil * 10 + \lceil 12/40 \rceil * 20 = 42$
  - Iteration 2:  $R_3 = 12 + \lceil 42/30 \rceil * 10 + \lceil 42/40 \rceil * 20 = 72$
  - Iteration 3:  $R_3 = 12 + \lceil 72/30 \rceil * 10 + \lceil 72/40 \rceil * 20 = 82$
  - Iteration 4:  $R_3 = 12 + \lceil 82/30 \rceil * 10 + \lceil 82/40 \rceil * 20 = 102$
  - ...
- Hence  $R_3 \rightarrow \infty$ . This means that T3's first job never finishes execution due to interferences by higher priority tasks, hence T3 is unschedulable
- It is also possible for T3 to be unschedulable if  $R_3$  converges but it exceeds its deadline  $D_3$ , e.g., if we set  $D_3 = 50$ , then  $R_3 = 52 > D_3 = 50$  (another job of T3 is released at time 50, but RTA for the current job is not affected by the newly-released job.)

Task	T=D	C	Prio
T1	30	10	H
T2	40	20	M
T3	52	12	L

Task	T=D	C	Prio
T1	30	10	H
T2	40	10	M
T3	50	12	L


## RM with $D < T$

- Deadline monotonic (Fixed Priority):
  - A task with smaller **relative** deadline gets higher priority  $P_i \propto 1/D_i$
  - No Utilization Bound test for RM or DM, for taskset with  $D_i < T_i$ ; need to use Response Time Analysis (RTA)
  - Consider a taskset with two tasks both with  $(C_i, T_i, D_i) = (1, 2, 1)$ . Using RTA, assuming  $\tau_1$  has higher priority (since task periods are equal, we can assign either task higher priority), we can determine  $R_1 = 1 \leq D_2 = 1, R_2 = 2 > D_2 = 1$ , hence it is unschedulable




# RM vs. DM Example

- Three tasks:  $\tau_1 = (0.5, 3, 3)$ ,  $\tau_2 = (1, 4, 4)$ ,  $\tau_3 = (2, 6, 6)$
- Under RM (or DM), priority ordering  $\tau_1 > \tau_2 > \tau_3$



- Three tasks with  $\tau_2$  assigned a smaller deadline of  $D_2 = 2$ :  $\tau_1 = (0.5, 3, 3)$ ,  $\tau_2 = (1, 4, \mathbf{2})$ ,  $\tau_3 = (2, 6, 6)$
- Under DM, priority ordering  $\tau_2 > \tau_1 > \tau_3$




---

# Earliest Deadline First (EDF) Scheduling




# Earliest Deadline First (EDF)

- As each task enters the system, it is assigned a deadline, and its priority is determined by its absolute deadline  $d_i$ 
  - The task with the earlier deadline is assigned the higher priority
  - Upwards arrows indicate arrival time; Downwards arrows indicate deadline
- Pros:
  - Optimal: can achieve 100% CPU utilization
- Cons:
  - High runtime overhead
  - Poor temporal isolation during overload




# EDF Scheduling Example

- Say you have two tasks, both released at time 0
  - T1 has WCET 5 ms, with deadline of 20 ms
  - T2 has WCET 10 ms, with deadline of 12 ms
- Non-EDF scheduling: T1 before T2, T2 misses its deadline at 12
- EDF scheduling: T2 before T1, both tasks meet their deadlines



# Schedulable Utilization Bound: RM vs. EDF

- The schedulable utilization bound for EDF Scheduling is 1:
  - A taskset is schedulable under EDF scheduling if its system utilization  $U = \sum_{i=1}^N \frac{C_i}{T_i} \leq 1$
  - Assumptions: task period equal to deadline ( $P_i = D_i$ ); tasks are independent (no resource sharing)
- Recall the schedulable utilization bound for Fixed-Priority scheduling:
- A taskset is schedulable under RM scheduling if its system utilization  $U = \sum_{i=1}^N \frac{C_i}{T_i} \leq N(2^{1/N} - 1)$ 
  - $U \rightarrow 0.69$  as  $N \rightarrow \infty$




# RM vs. EDF Example

Task	T=D	C
$\tau_1$	6	3
$\tau_2$	9	4


Under RM (Fixed-Priority scheduling), all jobs of  $\tau_1$  have higher priority than all jobs of  $\tau_2$ . Taskset non-schedulable with RM

$$U = \frac{3}{6} + \frac{4}{9} = 0.944 > 0.828$$



Under EDF (Dynamic Priority scheduling), different jobs of  $\tau_1$  and  $\tau_2$  may have different priorities, depending on their absolute deadlines  $d_i$ . Taskset schedulable with EDF

$$U = \frac{3}{6} + \frac{4}{9} = 0.944 < 1.0$$




When two jobs have equal priority, the newly arrived job does not preempt the running job

# RM vs. EDF: Robustness under Overload

- Under permanent overload, with CPU utilization  $U > 1$ 
  - Under EDF, all tasks execute at a slower rate with “period rescaling”, i. e., all tasks are delayed evenly
  - Under RM, higher priority tasks are protected while lower priority tasks are delayed or complete blocked
  - Recall [Slide 25 Example Lateless](#)
- Under transient overload, when some job overruns (executes longer than expected temporarily)
  - Under EDF, task overruns can cause deadline miss of arbitrary task
  - Under RM: task overruns only affect lower priority tasks
- Conclusion: RM offers better temporal isolation for higher priority tasks

## RM under permanent overload


$$U = \frac{4}{8} + \frac{6}{12} + \frac{5}{20} = 1.25$$



- High priority tasks execute at the proper rate
- Low priority tasks are completely blocked

## EDF under permanent overload


$$U = \frac{4}{8} + \frac{6}{12} + \frac{5}{20} = 1.25$$



- All tasks execute at a slower rate
- No task is blocked

# EDF Period Rescaling

- Theorem on Period Rescaling [Cervin et al. 2003]:
  - If system utilization  $U > 1$ , tasks are executed with an average period  $T'_i = T_i U$  under EDF scheduling




	$T_i$	$T'_i$
$\tau_1$	8	10
$\tau_2$	12	15
$\tau_3$	20	25

12

# EDF with $D < T$

- Earliest Deadline First (Dynamic-Priority):
  - A task with smaller **absolute** deadline gets higher priority  $P_i \propto 1/d_i$
  - Instead of Utilization Bound of 1, we use Density Bound of 1
  - Density of task  $\tau_i$  is defined as  $\delta_i = \frac{c_i}{\min(D_i, T_i)}$ . Taskset is schedulable if system density does not exceed 1:  $\delta = \sum_i \delta_i \leq 1$  (sufficient but not necessary condition)
  - Process Demand Analysis can be used as necessary and sufficient condition (discussion omitted)
  - Consider a taskset with two tasks both with  $(C_i, T_i, D_i) = (1, 2, 1)$ . It is obviously unschedulable under EDF. System utilization is  $U = \frac{1}{2} + \frac{1}{2} = 1$ ; System density  $\Delta = \frac{1}{1} + \frac{1}{1} = 2$ . However we cannot determine schedulability based on  $\Delta > 1$ .
  - Consider a taskset with two tasks  $\tau_1 = (0.6, 2, 1)$ ,  $\tau_2 = (2.3, 5, 5)$ .  $\Delta = \frac{0.6}{1} + \frac{2.3}{5} = 1.06$ . Yet the taskset is schedulable under EDF:





---

# Least Laxity First (LLF) Scheduling



# Least Laxity First (LLF)


- LLF assigns priority to jobs dynamically based on their laxity (slack)
  - With absolute deadline  $d_i$  and remaining execution time  $e_i$ , laxity of  $\tau_i$ 's job at time  $t$  is  $l_i = d_i - t - e_i$ . Job with the smallest laxity has the highest priority
  - If an active job runs in the previous time slot, then its laxity remains the same, as  $t$  is incremented by 1, and  $e_i$  is decremented by 1
  - If an active job does not run in the previous time slot, then its laxity is decremented by 1, as  $t$  is incremented by 1, and  $e_i$  remains the same
  - While an active job waits and does not run, its laxity decreases and its priority increases



- EDF and LLF are both optimal scheduling algorithms, i.e., they both have schedulable utilization bound of 1
  - LLF incurs frequent context switches, hence is less practical than EDF


# RM, EDF, LLF Example

Task	T=D	C
T1	5	2
T2	6	3



EDF and RM have the same schedule

Time	$\tau_1$ Laxity	$\tau_2$ Laxity	Running Task
t=0	$5-0-2=3$	$6-0-3=3$	$\tau_1$ (tie)
t=1	$5-1-1=3$	$6-1-3=2$	$\tau_2$
t=2	$5-2-1=2$	$6-2-2=2$	$\tau_2$ (tie)
t=3	$5-3-1=1$	$6-3-1=2$	$\tau_1$
t=4	$\tau_1$ done	$6-4-1=1$	$\tau_2$
t=5	$10-5-2=3$	$\tau_2$ done	$\tau_1$
t=6	$10-6-1=3$	$12-6-3=3$	$\tau_1$ (tie)
t=7	$\tau_1$ done	$12-7-3=2$	$\tau_2$



LLF has more frequent context switches

---

# Preemptive vs. Non-Preemptive Scheduling

# Preemptive vs. Non-Preemptive Scheduling

- Non-preemptive scheduling pros:

- It reduces runtime overhead
  - Less context switches
  - No mutex locks needed for critical sections
- It preserves program locality, improving the effectiveness of CPU cache
  - As a result, task WCET becomes smaller and execution time distribution becomes more predictable (shown on right)
- Sometimes NP scheduling can improve schedulability

- Cons:


- Reduced schedulability
- Scheduling anomalies

- Preemptive scheduling pros:

- Better schedulability (higher CPU utilization)

- Cons:

- Runtime overhead due to frequent context-switches
- Destroys program locality so task WCET becomes larger




# Sometimes NP Scheduling Improves Schedulability


- An example where NP scheduling improves schedulability (for fixed-priority sched'g)

$$U = \frac{2}{5} + \frac{4}{7} \cong 0.97$$

**RM**




**NP-RM**



## Disadvantage of NP Scheduling: Reduced Schedulability

- In general, NP scheduling reduces schedulability. The utilization bound under NP scheduling drops to zero due to blocking time
- An example with two tasks T1 and T2, CPU utilization of nearly 0, yet non-schedulable.
  - If  $C_2$  (WCET of T2)  $\geq T_1$  (period of T1), then the taskset is non-schedulable with arbitrarily small system CPU utilization  $\frac{C_1}{T_1} + \frac{C_2}{T_2} \rightarrow \frac{0}{T_1} + \frac{C_2}{\infty}$  (when  $C_1$  goes to 0 and  $T_2$  goes to infinity)
  - This example is valid whether  $\tau_1$  or  $\tau_2$  has higher priority: even if  $\tau_1$  has higher priority, it may be released very shortly after  $\tau_2$  is released at time 0, and it has to wait for  $\tau_2$  to finish due to NP scheduling



# Disadvantage of NP Scheduling: Scheduling Anomalies

- Scheduling anomaly: doubling the processor speed (reducing task execution times by half) makes task  $\tau_1$  miss its deadline





---

# Resource Synchronization Protocols (for Fixed-Priority Scheduling)




# Resource Sharing

- Some shared resources do not allow simultaneous accesses but require mutual exclusion. A piece of code executed under mutual exclusion constraints is called a critical section.
- When two tasks access shared resource, mutexes or binary semaphores are used to protect critical section
- Each shared resource  $R_i$  must be protected by a semaphore  $S_i$ , and each critical section (CS) using resource  $R_i$  must begin with  $\text{wait}(S_i)$  and end with  $\text{signal}(S_i)$
- A task waiting for an exclusive resource is said to be blocked on that resource. Otherwise, it proceeds by entering the critical section and holds the resource. When a task leaves a critical section, the associated resource becomes free.
- Tasks blocked on the same resource are kept in a queue. When a running task invokes  $\text{wait}(S_i)$  when  $S_i$  is locked, it enters a waiting state, until another task invokes  $\text{signal}(S_i)$  to unlock  $S_i$



# Blocking Time

- Lower-priority tasks can cause delay to higher-priority tasks due to blocking time
- If higher priority task  $\tau_1$  tries to lock a semaphore that is locked by lower priority task  $\tau_2$ ,  $\tau_1$  blocks until  $\tau_2$  unlocks the semaphore, and  $\tau_1$  experiences a blocking delay  $\Delta$ .
  - Since typical Critical Sections are very short, it seems this blocking time delay  $\Delta$  is bounded by the longest critical section in lower-priority tasks, hence acceptable?
- No, blocking delay may be unbounded!



## Example Taskset

---

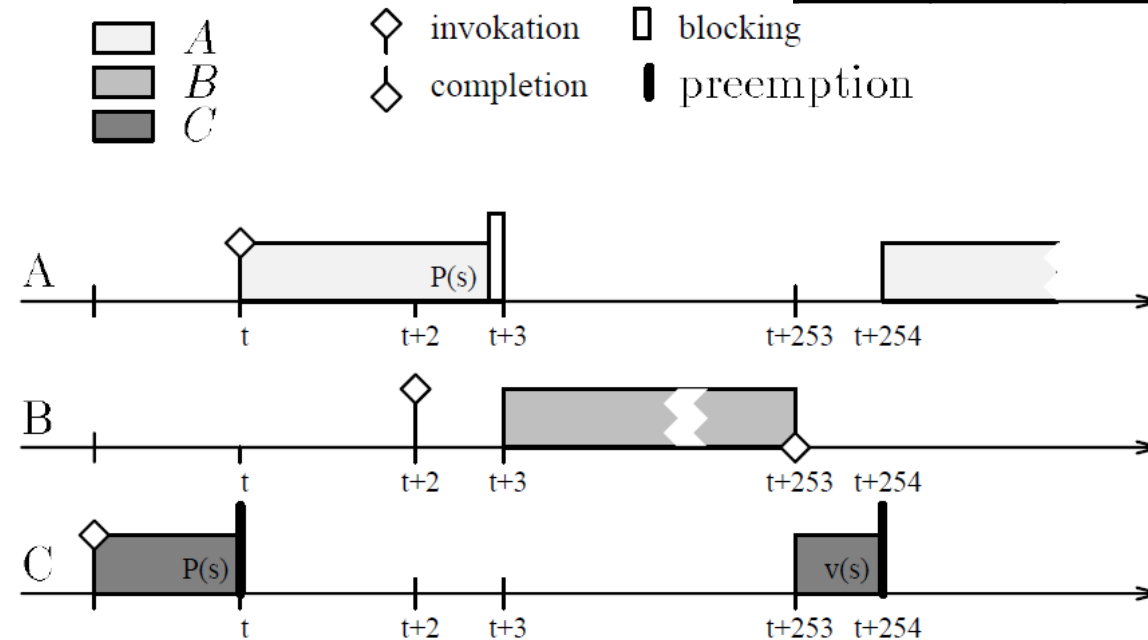
- system utilization  $U = \frac{5}{50} + \frac{250}{500} + \frac{1000}{3000} = 0.93 > 0.780$ 
  - Since utilization exceeds the Utilization Bound of 0.780 of 3 tasks under RM scheduling, we cannot determine schedulability by the Utilization Bound test
- RTA shows that the taskset is schedulable by computing WCRT of each task (without shared resources):
  - $R_A = C_A + 0 = 5 + 0 = 5$
  - $R_B = C_B + \left\lceil \frac{R_B}{T_A} \right\rceil \cdot C_A = 250 + \left\lceil \frac{R_B}{50} \right\rceil \cdot 5 = 280$
  - $R_C = C_C + \left\lceil \frac{R_C}{T_A} \right\rceil \cdot C_A + \left\lceil \frac{R_C}{T_B} \right\rceil \cdot C_B = 1000 + \left\lceil \frac{R_C}{50} \right\rceil \cdot 5 + \left\lceil \frac{R_C}{500} \right\rceil \cdot 250 = 2500$

Task	T	D	C	Prio
A	50	10	5	H
B	500	500	250	M
C	3000	3000	1000	L

# Priority Inversion

Task	T	D	C	Prio
A	50	10	5	H
B	500	500	250	M
C	3000	3000	1000	L

- HP: High-Priority; MP: Medium-Priority; LP: Low-Priority
- $t$ : LP C locks  $s$
- $t+\Delta$ : HP task A preempts LP task C
- $t+2$ : MP task B is invoked, but cannot start running due to HP task A running
- $t+3$ : HP task A tries to lock  $s$ , blocks since LP task C is holding  $s$ ; MP task B starts running
- $t+253$ : MP task B finishes; LP task C starts running
- $t+254$ : LP task C unlocks  $s$ ; HP task A starts running, but it already missed its deadline long ago!



# Priority Inversion and Priority Inheritance


---

- Priority inversion happened
  - High priority task (A) is blocked by low-priority task (B) for an **unbounded interval of time**.
    - » More than the longest critical section of B
- In 1997, this bug caused the Mars pathfinder to freeze up occasionally without explanation, and then starts working again
- Fixed by uploading a software patch enabling Priority-Inheritance Protocol (PIP)
  - When a task locks a semaphore, it inherits the highest priority of all tasks blocked waiting for the semaphore
- A task in a CS increases its priority if it blocks other higher priority tasks, by inheriting the highest priority among those tasks it blocks.
  - $P_{CS} = \max\{P_k | \tau_k \text{ blocked on CS}\}$

# With PIP


Task	T	D	C	Prio
A	50	10	5	H
B	500	500	250	M
C	3000	3000	1000	L

- $t$ : LP task C locks  $s$
- $t+\Delta$ : HP task A preempts LP task C
- $t+2$ : MP task B is invoked, but cannot start running due to HP task A running
- $t+3$ : HP task A tries to lock  $s$ , blocks since LP task C is holding  $s$ ; **C inherits A's priority and starts running**
  - MP task B cannot start running, hence cannot cause unbounded blocking to HP task A
- $t+4$ : LP task C unlocks  $s$ , and returns to its regular Low priority; HP task A locks  $s$  and starts running
- $t+5$ : HP task A finishes and meets its deadline.




# PIP Properties

- PIP does not prevent deadlocks. Classic deadlock scenario:
  - LP task  $\tau_2$  locks  $s_1$  at time  $t$
  - HP task  $\tau_1$  starts running after  $t$  and locks  $s_2$ , and tries to lock  $s_1$ , blocked by  $\tau_2$
  - $\tau_2$  inherits  $\tau_1$ 's priority and starts running
  - $\tau_2$  tries to lock  $s_2$ , but  $\tau_1$  holds  $s_2$ . Deadlocked!



# PIP Properties

- Chained blocking:
  - Theorem: task  $\tau_i$  can be blocked at most once by each lower priority task
  - If  $n$  is the number of tasks with priority less than  $\tau_i$ , and  $m$  is the number of semaphores on which  $\tau_i$  can be blocked, then
    - **Theorem:**  $\tau_i$  can be blocked at most for the duration of  $\min(n, m)$  critical sections
- Priority Ceiling Protocol is a more advanced protocol, which prevents deadlocks and reduces blocking time





# Schedulability Analysis

---

- Schedulable utilization bound for RM scheduling with blocking time:
  - A taskset is schedulable under RM scheduling with blocking time if
  - $\forall i$ , priority level  $i$  utilization  $U_i = \sum_{k=1}^{i-1} \frac{C_k}{T_k} + \frac{C_i + B_i}{T_i} \leq i(2^{1/i} - 1)$
- Response Time Analysis (RTA) for RM scheduling with blocking time:
  - WCRT  $R_i$  is computed by solving the following recursive equation:
  - $R_i = C_i + B_i + \sum_{\forall j \in hp(i)} \left\lceil \frac{R_i}{T_j} \right\rceil C_j$
  - where  $B_i$  is the maximum blocking time experienced by task  $\tau_i$  due to shared resources

---

# Multiprocessor Scheduling

# Multiprocessor models

---

- Identical multiprocessors:
  - each processor has the same computing capacity
- Uniform multiprocessors:
  - different processors have different computing capacities
- Heterogeneous multiprocessors:
  - each (task, processor) pair may have a different computing capacity
- MP scheduling
  - Many NP-hard problems, with few optimal results, mainly heuristic approaches
  - Only sufficient schedulability tests


# Multiprocessor Models

---

Identical multiprocessors: each processor has the same speed

Task T1

Task T2




# Multiprocessor Models

Uniform multiprocessors: different processors have different speeds


Task T1

Task T2




P1

speed = 1



P2

speed = 2




P3

speed = 3

# Multiprocessor Models


**Heterogeneous** multiprocessors: each (task, processor) pair may have a different relative speed, due to specialized processor architectures




# Global vs partitioned scheduling

- Global scheduling
  - All ready jobs are kept in a common (global) queue; when selected for execution, a job can be dispatched to an arbitrary processor, even after being preempted
- Partitioned scheduling
  - Each task may only execute on a specific processor

**Global scheduling:  
Single system-wide queue**



**Partitioned scheduling:  
per-processor queues**



# Global Scheduling vs. Partitioned Scheduling

---

- Global Scheduling

- Advantages:

- Runtime load-balancing across cores
  - » More effective utilization of processors and overload management
- Supported by most multiprocessor operating systems
  - » Windows, Linux, MacOS...

- Disadvantages:

- Low schedulable utilization
- Weak theoretical framework

- Partitioned Scheduling

- Advantages:

- Mature scheduling framework
- Uniprocessor scheduling theory scheduling are applicable on each core; uniprocessor resource access protocols (PIP, PCP...) can be used
- Partitioning of tasks can be done by efficient bin-packing algorithms

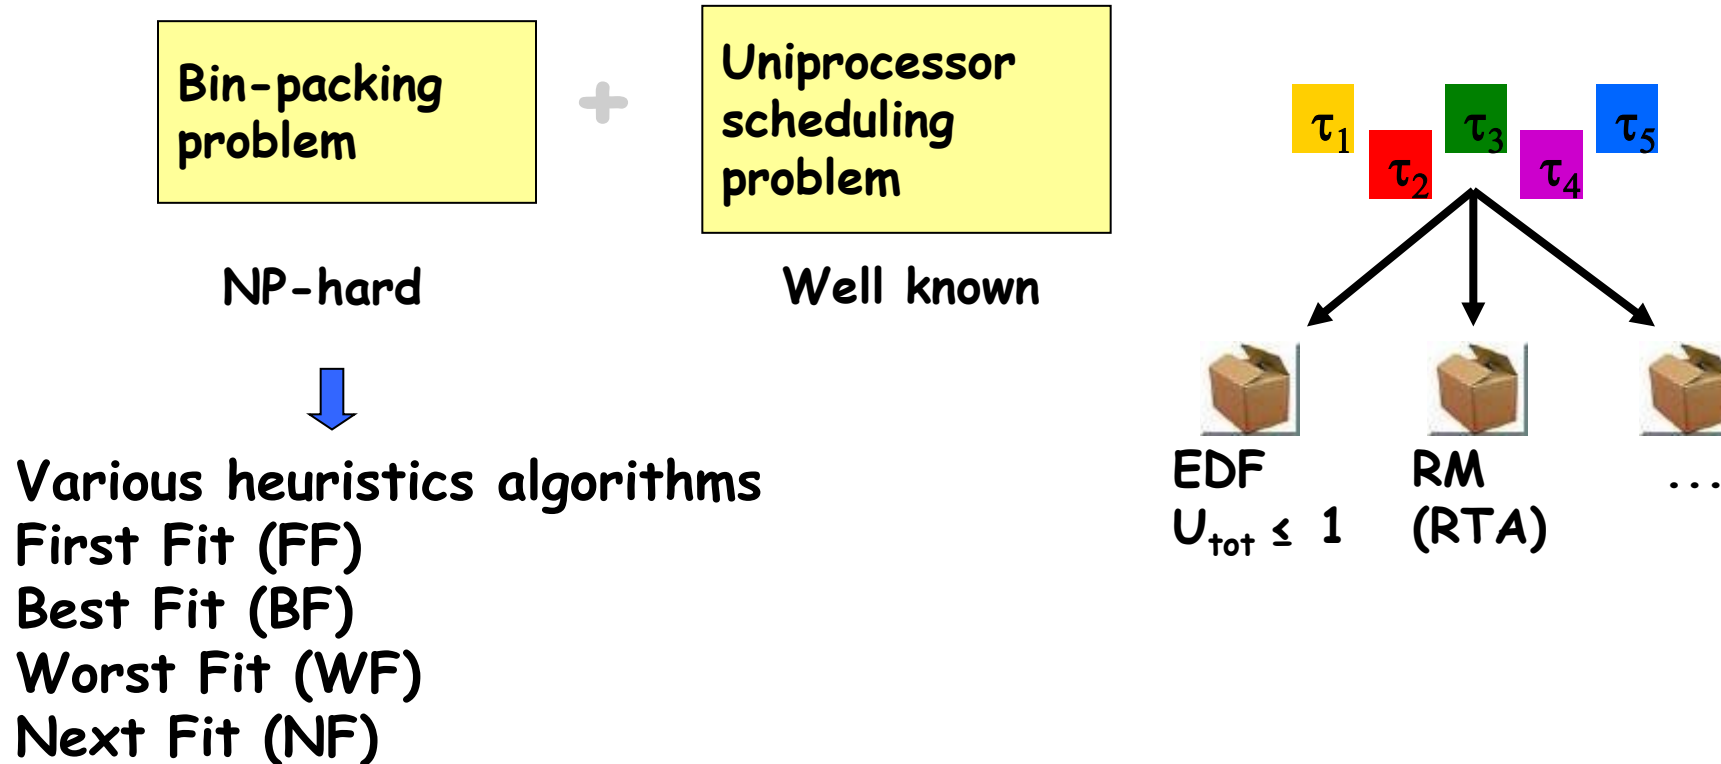
- Disadvantages:

- No runtime load-balancing; surplus CPU time cannot be shared among processors



# Partitioned Scheduling

- Scheduling problem reduces to:



# Partitioned Scheduling

---

- Bin-packing algorithms:
  - The problem concerns packing objects of varying sizes in boxes ("bins") with some optimization objective, e.g., minimizing number of used boxes (best-fit), or minimizing the maximum workload for each box (worst-fit)
- Application to multiprocessor scheduling:
  - Bins are represented by processors and objects by tasks
  - The decision whether a processor is "full" or not is derived from a utilization-based feasibility test.
- Since optimal bin-packing is a NP-complete problem, partitioned scheduling is also NP-complete
- Example: Rate-Monotonic-First-Fit (RMFF): (Dhall and Liu, 1978)
  - Let the processors be indexed as 1, 2, ...
  - Assign the tasks to processor in the order of increasing periods (that is, RM order)
  - For each task  $\tau_i$ , choose the lowest previously-used processor  $j$  such that  $\tau_i$ , together with all tasks that have already been assigned to processor  $j$ , can be feasibly scheduled according to the utilization-based schedulability test
  - Additional processors are added if needed

# Assumptions for Global Scheduling


---

- Identical multiprocessors
- Work-conserving:
  - At each instant, the highest-priority jobs that are eligible to execute are selected for execution upon the available processors
  - No processor is ever idle when the ready queue is non-empty
- Preemption and Migration support
  - For global schedulers, a preempted task can resume its execution on a different processor with 0 overhead
  - Cost of preemption/migration integrated into task WCET
- No job-level parallelism
  - the same job cannot be *simultaneously* executed on more than one processor

# Source of Difficulty

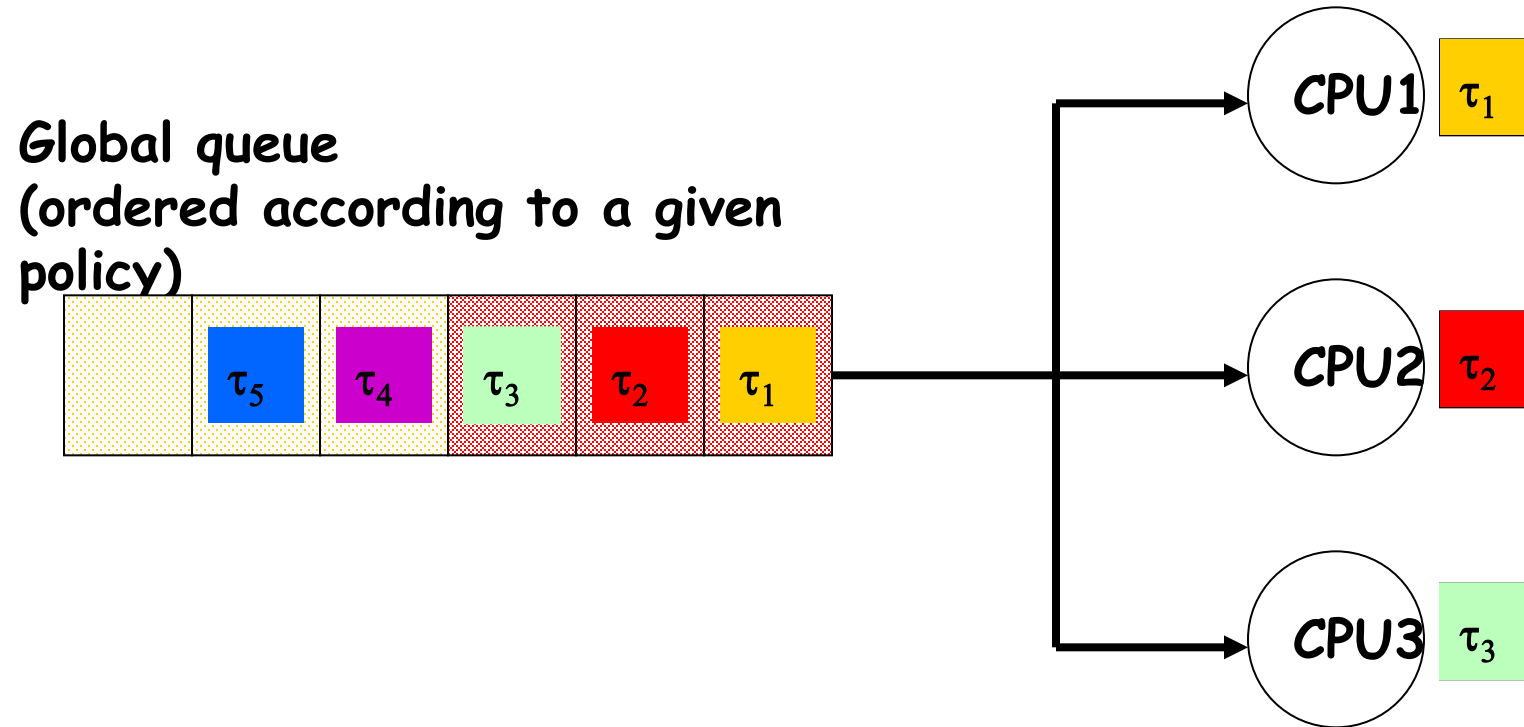
---

- The “no job-level parallelism” assumption leads to difficult scheduling problems
- “The simple fact that a task can use only one processor even when several processors are free at the same time adds a surprising amount of difficulty to the scheduling of multiple processors” [Liu’69]




# Global scheduling example

---



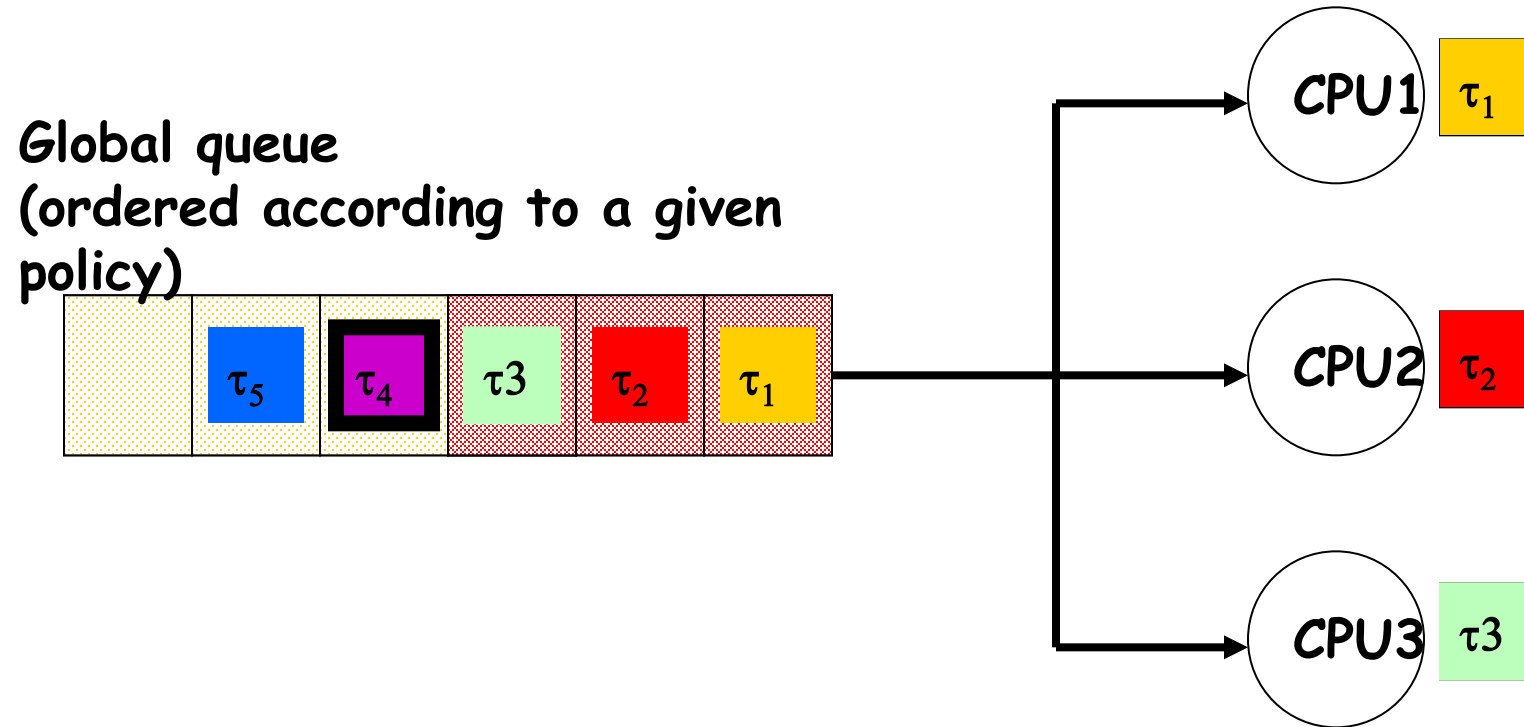
The first  $m$  jobs in the queue are scheduled upon the  $m$  CPUs

# Global scheduling example




When a job  $\tau_3$  finishes its execution, the next job in the queue  $\tau_4$  is scheduled on the available CPU

# Global scheduling example



When a new higher priority job  $\tau_3$  arrives, it preempts the job with lowest priority  $\tau_4$  among the executing ones

# Global scheduling example



When another job  $\tau_1$  finishes its execution, the preempted job  $\tau_4$  can resume its execution. Net effect:  $\tau_4$  "migrated" from CPU3 to CPU1



## Global vs. Partitioned


---

- Global (work-conserving) and partitioned scheduling algorithms are incomparable:
  - There are tasksets that are schedulable with a global scheduler, but not with a partitioned scheduler, and vice versa.

# Global vs Partitioned (FP) Scheduling

Task	T=D	C	Prio
T1	2	1	H
T2	3	2	M
T3	3	2	L

- A taskset schedulable with global scheduling, but not partitioned scheduling. System utilization  $U = \frac{1}{2} + \frac{2}{3} + \frac{2}{3} = 1.83$
- Global FP scheduling is schedulable with priority assignment  $p_1 > p_2 > p_3$  (or  $p_2 > p_1 > p_3$ )
- Partitioned scheduling is unschedulable, since assigning any two tasks to the same processor will cause that processor's utilization to exceed 1.



A feasible execution trace under global scheduling

# Global vs Partitioned (FP) Scheduling


Task	T=D	C	Prio
T1	6	4	4(H)
T2	12	7	3
T3	12	4	2
T4	24	10	1(L)

- A taskset schedulable with partitioned scheduling, but not global scheduling. System utilization  $U = \frac{4}{6} + \frac{7}{12} + \frac{4}{12} + \frac{10}{24} = 2.0$ , hence the two processors must be fully utilized with no possible idle intervals

- Partitioned FP scheduling with RM priority assignment ( $p_1 > p_2 > p_3 > p_4$ ) is schedulable. T1, T3 assigned to Processor 1; T2, T4 assigned to Processor 2. Both processors have utilization 1.0, and harmonic task periods


- Global FP scheduling with RM priority assignment  $p_1 > p_2 > p_3 > p_4$  is unschedulable. Compared to partitioned scheduling, the difference is at time 7, when T3 (with higher priority than T4) runs on Processor 2. This causes idle intervals on Processor 1 [10,12] and [22,24], since only one task T4 is ready during these time intervals

At time 7, T4 runs on Processor 2



A feasible execution trace under partitioned scheduling

At time 7, T3 runs on Processor 2



An infeasible execution trace under global scheduling


# Difficulties of Global Scheduling

---

- Dhall's effect
  - With RM, DM and EDF, some low-utilization task sets can be unschedulable regardless of how many processors are used.
- Scheduling anomalies
  - Decreasing task execution time or increasing task period may cause deadline misses
- Hard-to-find worst-case
  - The worst-case does not always occur when a task arrives at the same time as all its higher-priority tasks
- Dependence on relative priority ordering (omitted)
  - Changing the relative priority ordering among higher-priority tasks may affect schedulability for a lower-priority task

# Dhall's effect

- Global RM/DM/EDF can fail at very low utilization
- Example:  $m$  processors,  $n=m+1$  tasks. Tasks  $\tau_1, \dots, \tau_m$  are light tasks, with small  $C_i = 1, T_i = D_i = T - 1$ ; Task  $\tau_{m+1}$  is a heavy task, with large  $C_i = T, T_i = D_i = T$ ;
- For global RM/DM/EDF, Task  $\tau_{m+1}$  has lowest priority, so  $\tau_1, \dots, \tau_m$  must run on  $m$  processors starting at time 0, causing  $\tau_{m+1}$  to miss its deadline
- One solution: assign higher priority to tasks with higher utilization



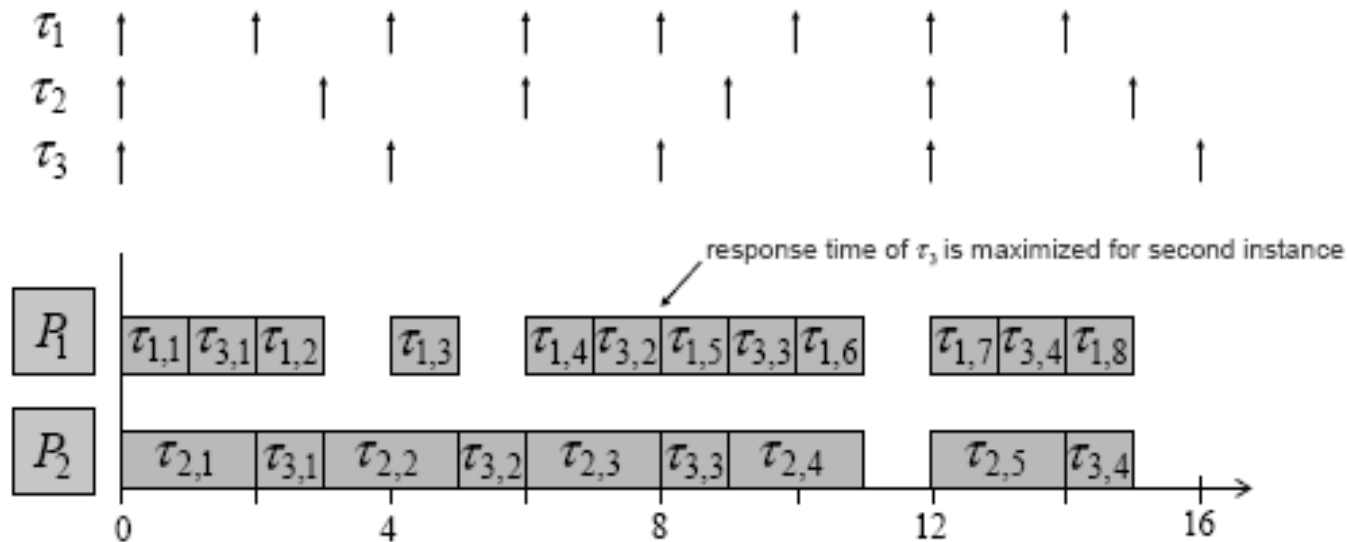
# Hard-to-Find Worst-Case

- For uniprocessor scheduling, the worst case occurs when all tasks are initially released at time 0 simultaneously, called the critical instant (Recall Slide [Response Time Analysis \(RTA\)](#)). This is not true for multiprocessor scheduling.
- The worst-case for a task does not always occur at time 0 when it arrives at the same time as all its higher-priority tasks.
  - Response time for task  $\tau_3$  is maximized for its 2<sup>nd</sup> job  $\tau_{3,2}$  (8-4=4), which does not arrive at the same time as its higher priority tasks; not for its 1<sup>st</sup> job  $\tau_{3,1}$  (3-0=3), which arrives at the same time as its higher priority tasks

Hard-to-find critical instant:

(RM scheduling)

$$\begin{aligned}\tau_1 &= \{C_1 = 1, T_1 = 2\} \\ \tau_2 &= \{C_2 = 2, T_2 = 3\} \\ \tau_3 &= \{C_3 = 2, T_3 = 4\}\end{aligned}$$



# MP Scheduling Anomalies


---

- Decrease in processor demand (decreasing task execution time or increasing task period) may cause deadline misses!
- **Anomaly 1**
  - Decrease in processor demand from higher-priority tasks can *increase the interference on a lower-priority task because of change in the time when the tasks execute*
- **Anomaly 2**
  - Decrease in processor demand of a task *negatively affects the task itself because change in the task arrival times cause it to suffer more interference*


# Scheduling Anomaly Example 1

- Two processors and total taskset utilization  $\sum U_i = 1.83$
- With Task a's period  $T_a = 3$ , WCRT of task c is  $R_c = 12 \leq D_c = 12$ .  $R_c = C_c + I_c = 8 + I_c$ , where  $I_c = 2 + 1 + 1 = 4$  is interference by higher priority tasks a and b. (Task c experiences inference when both processors are busy executing higher priority tasks a and b.) Task c is schedulable but saturated, as any increase in its WCET or interference would make it unschedulable.
- With Task a's period  $T_a = 4$ , WCRT of task c is  $R_c = 14 > D_c = 12$ .  $R_c = 8 + I_c$  where  $I_c = 2 + 2 + 2 = 6$ . This is because execution segments of tasks a and b on two processors are aligned in time, thus causing more interference to task c

Task	T=D	C	Util	Prio
a	3	2	0.67	H
b	4	2	0.5	M
c	12	8	0.67	L



Task	T=D	C	Util	Prio
a	4	2	0.5	H
b	4	2	0.5	M
c	12	8	0.67	L

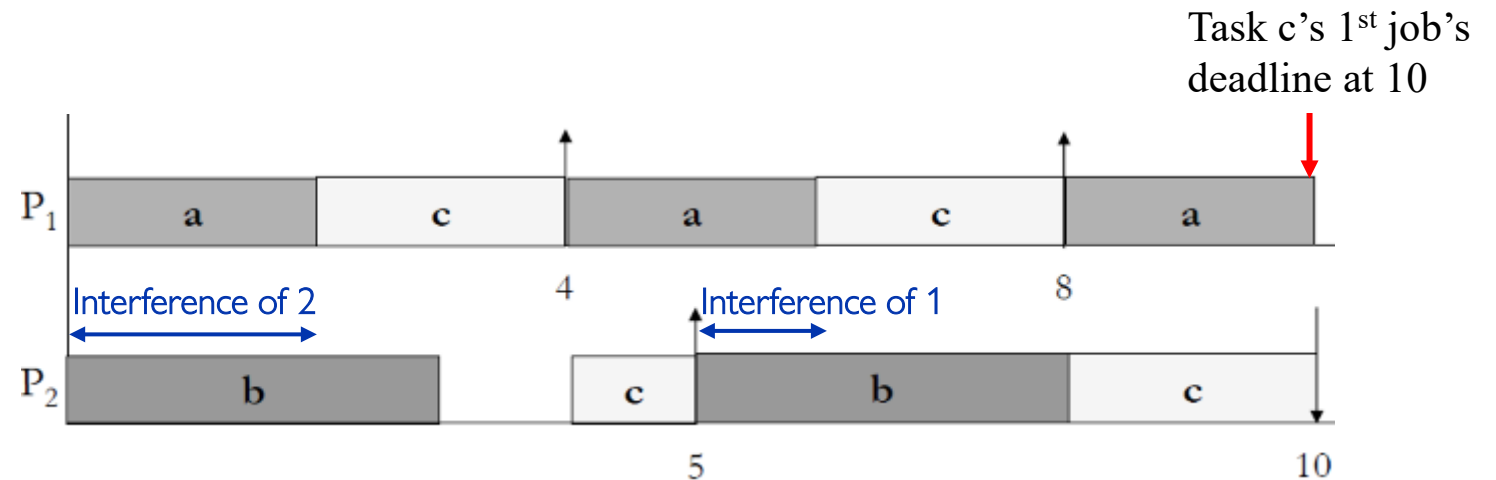




## Scheduling Anomaly Example 2

- Two processors and total task utilization  $\sum U_i = 1.8$
- With Task c's period  $T_c = 10$ , WCRT of task c is  $R_c = 10 \leq D_c = 10$ .  $R_c = C_c + I_c = 7 + 3 = 10$ , where  $I_c = 2 + 1 = 3$  is interference by higher priority tasks a and b. Its 1st job meets its deadline at time 10. This schedule repeats in future periods, hence task c is schedulable but saturated, as any increase in its WCET or interference would make it unschedulable.


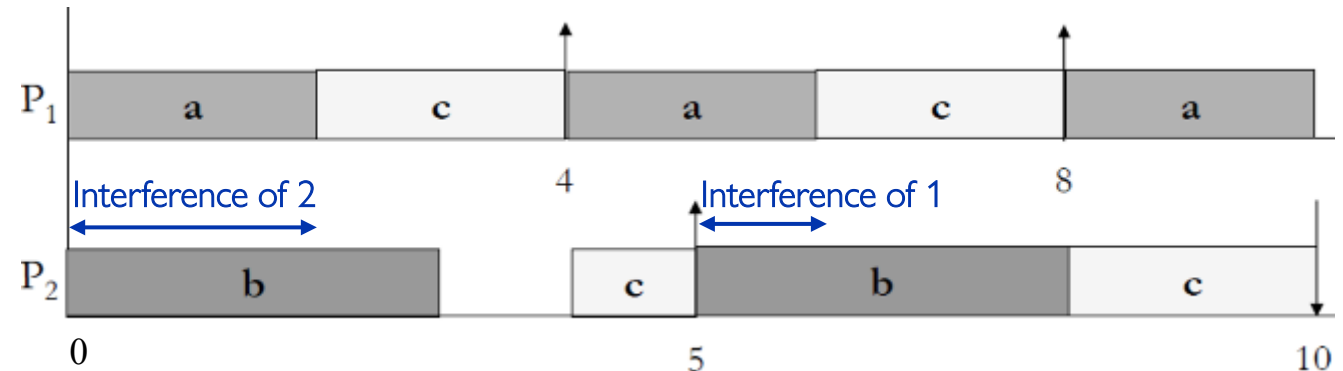
Task	T=D	C	Util	Prio
a	4	2	0.5	H
b	5	3	0.6	M
c	10	7	0.7	L



## Scheduling Anomaly Example 2

- With Task c's period  $T_c = 11$ , WCRT of task c is  $R_c = 12 > D_c = 10$ . Its 1<sup>st</sup> job has response time  $C_c + I_c = 7 + 3 = 10 \leq D_c = 11$ , where  $I_c = 2 + 1 = 3$ , but this is not task c's WCRT.
- Its 2<sup>nd</sup> job has response time  $C_c + I_c = 7 + 5 = 12 > D_c = 11$ , where  $I_c = 1 + 2 + 2 = 5$ . The 2<sup>nd</sup> job finishes at time  $11+12=23$ , and misses its deadline at time 22.
- Another example where the worst-case for task c does not occur at time 0 when it arrives at the same time as all its higher-priority tasks

Task	T=D	C	Util	Prio
a	4	2	0.5	H
b	5	3	0.6	M
c	11	7	0.64	L



## Scheduling Anomaly Example 3

- Doubling processor speed causes T1 to miss its deadline
  - (Yellow part denotes a critical section shared by T1 and T2)

